

# AI Notes

*Huu Duc Nguyen M.Sc.*

29 February 2022

# Contents

<b>Abbreviations</b>	<b>1</b>
<b>1 Introduction</b>	<b>2</b>
<b>2 Overview of Machine Learning</b>	<b>4</b>
2.1 Task $T$ . . . . .	4
2.1.1 Classification . . . . .	4
2.1.2 Regression . . . . .	4
2.1.3 Clustering . . . . .	4
2.1.4 Others . . . . .	5
2.2 Performance $P$ . . . . .	5
2.3 Experience $E$ . . . . .	5
2.3.1 Supervised Learning . . . . .	5
2.3.2 Unsupervised Learning . . . . .	5
2.4 Model Parameters and Loss Function . . . . .	6
<b>3 Probabilities</b>	<b>7</b>
3.1 Definitions . . . . .	7
3.1.1 Basic Definitions . . . . .	7
3.1.2 Independence and Variability . . . . .	7
3.1.3 Bayes Rule . . . . .	8
3.1.4 Expectation . . . . .	8
3.2 Types of Probability Distributions . . . . .	8
3.2.1 Bernoulli Distribution . . . . .	8
3.2.2 Categorical Distribution . . . . .	9
3.2.3 Univariate Normal Distribution . . . . .	9
3.2.4 Multivariate Normal Distribution . . . . .	9
3.2.5 Beta Distribution . . . . .	10
3.3 Parameter Estimation . . . . .	10
3.3.1 Maximum Likelihood Estimation . . . . .	11
3.3.2 Maximum A Posteriori . . . . .	11
3.4 Naive Bayes Classifier . . . . .	11
3.5 Views on the Decision Problem . . . . .	12
3.5.1 Generative Methods . . . . .	12

3.5.2 Discriminative Methods . . . . .	12
3.6 Unknown Notes . . . . .	12
3.7 Probability Density Estimation . . . . .	13
3.7.1 Histogram . . . . .	13
3.7.2 Parametric Probability Density Estimation . . . . .	13
3.7.3 Kernel Methods . . . . .	14
3.7.4 K-Nearest Neighbor . . . . .	14
3.7.5 Mixture of Gaussians . . . . .	15
3.7.6 K-Means Clustering . . . . .	15
3.7.7 EM Clustering . . . . .	16
<b>4 Linear Regression</b>	<b>17</b>
<b>5 Logistics Regression</b>	<b>19</b>
<b>6 Softmax Regression</b>	<b>20</b>
<b>7 Error Functions</b>	<b>21</b>
7.1 Ideal Miss-classification Error . . . . .	21
7.2 Squared Error - $L_2$ Loss . . . . .	21
7.3 Cross Entropy Error . . . . .	22
7.4 Squared Error on Sigmoid / Tanh . . . . .	22
7.5 Hinge Error . . . . .	22
7.6 $L_1, L_0$ Loss . . . . .	22
7.7 Average Loss . . . . .	23
<b>8 Neural Network</b>	<b>24</b>
8.1 Gradient Descent . . . . .	24
8.1.1 Vanilla Gradient Descent . . . . .	24
8.1.2 Momentum . . . . .	24
8.1.3 Nesterov Accelerated Gradient . . . . .	24
8.1.4 Root mean squared prop (RMSprop) . . . . .	25
8.1.5 Adaptive moment estimation (Adam) . . . . .	25
8.2 Tips and Tricks . . . . .	26
<b>Bibliography</b>	<b>I</b>

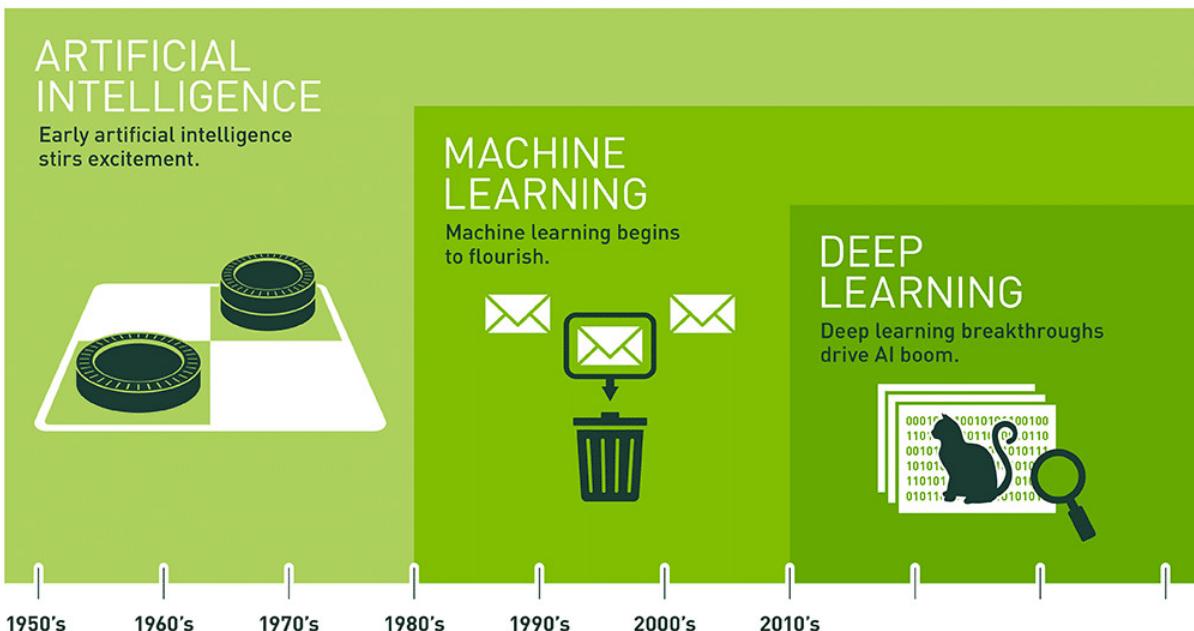


# Abbreviations

<b>AI</b>	Artificial Intelligence
<b>ML</b>	Machine Learning
<b>DL</b>	Deep Learning
<b>CS</b>	Computer Science
<b>CV</b>	Computer Vision
<b>RL</b>	Reinforcement Learning
<b>NLP</b>	Natural Language Processing
<b>prob.</b>	probability
<b>params.</b>	parameters
<b>algor.</b>	algorithms
<b>PDF</b>	Probability Density Function
<b>MLE</b>	Maximum Likelihood Estimation
<b>MAP</b>	Maximum A Posteriori
<b>MoG</b>	Mixture of Gaussians
<b>NAG</b>	Nestorov Accelerated Gradient
<b>RMSprop</b>	Root mean squared prop
<b>Adam</b>	Adaptive moment estimation
<b>PCA</b>	Principal Component Analysis
<b>K-L</b>	Kullback–Leibler
<b>ReLU</b>	Rectified Linear Unit

# 1 Introduction

Artificial Intelligence ([AI](#)) is the study field that leverages the ability of machines to mimic the problem-solving skill of human. It lies in the core of countless novel applications in real life, self-driving cars, virtual assistant, face recognition, etc. Machine Learning ([ML](#)) is a sub-field of Computer Science ([CS](#)) and [AI](#), that “gives computers the ability to learn without being explicitly programmed” ([Wikipedia](#)). As a great amount of collected data and powerful computational hardware arise, Deep Learning ([DL](#)) is then a subset of [ML](#) (Fig. 3.1). Advanced applications which relate to Natural Language Processing ([NLP](#)), Computer Vision ([CV](#)), robotic learning, etc., are with in this [DL](#) subset.



Since an early flush of optimism in the 1950s, smaller subsets of artificial intelligence – first machine learning, then deep learning, a subset of machine learning – have created ever larger disruptions.

**Figure 1.1:** The relation between [AI](#), [ML](#) and [DL](#) ([src](#)).

## [TODO: The structure of the notes]

It's important to understand that there are more to [AI](#) and [ML](#) than just neural networks. In the end, to create a meaningful and working neural network, I believe one should have strong background in the basics of [ML](#) as well. Henceforth, I keep record of what I have learnt in the whole [AI](#) field. The structure of the notes is as follows:

- Chap. 2 introduces common ideas in [ML](#).
- Chap. 3 introduces the mathematics background on probabilities, matrix.

- [TODO: chapter 3] explain basic concepts, the branching of different classes in ML. Later chapters presents each smaller branches.

# 2 Overview of Machine Learning

A machine learning algorithm is an algorithm that has the ability to *learn* from the data. A computer program is said to **learn**, if its performance at tasks in  $T$ , measured by  $P$ , improves with experience  $E$  (in which the experience is equivalent to the data). [GBC16]

## 2.1 Task $T$

A *task* is usually described by how the **ML** model process a single *data point*. This section presents some common **ML** tasks. [Vu18]

### 2.1.1 Classification

The task is to specify a label for the given data point. The labels are usually members of a list.

E.g., in the problem of digit classification, the data point is images of hand-written numbers. The data set comes with their labels as well. The task is then, given a unseen image, the model would be able to tell which number is in that image. In this problem, there are 10 possible labels, i.e.,  $0, 1, \dots, 9$ .

### 2.1.2 Regression

If the desired output is a real value, instead of a label in a list, then it's a regression problem. E.g.:

- with an image as the input data, the model predicts the age of the person
- given a feature vector, the model generates an image

### 2.1.3 Clustering

This is the task of grouping relevant data points based on some relationship between them.

E.g., find the pattern in customer shopping behaviors.

### 2.1.4 Others

Some worth-mentioning tasks:

- Recommendation System
- Machine Translation
- Completion
- Ranking
- Information Retrieval
- Denoising

## 2.2 Performance $P$

Usually, the dataset is divided into *training set* and *test set*. The model uses the training set to tune update the model parameters ([params.](#)) and the test set to examine the performance.

*Online training* is the approach when new data will continuously arise and introduce for the model to learn. E.g. in Reinforcement Learning ([RL](#)). *Offline training* is the opposite, the model learns from the a fixed training set.

## 2.3 Experience $E$

### 2.3.1 Supervised Learning

Supervised Learning is the approach that predict the outputs of new data points based on pairs of known inputs and outputs. This is the most common type of [ML](#) algorithms ([algor.](#)).

### 2.3.2 Unsupervised Learning

On the opposite, with unsupervised Learning, there is no known output, just inputs. Unsupervised Learning [algor.](#) will carry on some tasks based on the characteristics of the dataset, e.g. clustering, dimension reduction.

## **2.4 Model Parameters and Loss Function**

Each **ML** model is described by a set of model **params.**. E.g., in the problem of finding a line passing through points in the 2D plane, the model **params.** are  $a, b$  in the line equation  $y = ax + b$ . The aim of training is to find the model **params.** that leads to the best performance. For classification problems, it means having the least number of incorrect classified data points. For regression problems, it means having the smallest difference with the actual output. It is then equivalent to having an optimization problem, in which we try to minimize a loss/cost function.

# 3 Probabilities

## 3.1 Definitions

### 3.1.1 Basic Definitions

- If  $x$  is discrete:  $\sum_x p(x) = 1$  with  $\forall 0 \leq p(x) \leq 1$
- If  $x$  is continuous:  $\int p(x) dx = 1 \Rightarrow \exists$  a **Probability Density Function (PDF)**  
 $p(x)$  can take any positive value, as long as  $\int p(x) dx = 1$   
[TODO: Add image]  
**NOTE:** : theoretically  $p(x) = 0, \forall x$
- Common types

Joint probability:  $p(x_i, y_i)$  ( $= p(X = x_i, Y = y_i)$ )

Marginal probability:  $p(x_i)$  ( $= p(X = x_i)$ )

Conditional probability:  $p(y_i|x_i)$  ( $= p(Y = y_i|X = x_i)$ )

- Sum rule:  $\sum$  joint probability (prob.) = marginal prob.  
 $\Rightarrow$  Marginalization
  - discrete variable:  $p(x) = \sum_y p(x, y)$
  - continuous variable:  $p(x) = \int p(x, y) dy$
- Product rule: Product of marginal prob. and conditional prob. = joint prob.

### 3.1.2 Independence and Variability

- Independence. E.g.:  $x, y$  are independent, then

$$\begin{cases} p(x|y) = p(x) \\ p(y|x) = p(y) \end{cases} \iff p(x, y) = p(x).p(y)$$

- Variability
  - variance:  
$$var [f] = \mathbb{E} [(f(x) - \mathbb{E} [f(x)])^2] = \mathbb{E} [f(x)^2] - \mathbb{E} [f(x)]^2$$
  - covariance:  
$$cov [x, y] = \mathbb{E}_{x,y} [xy] - \mathbb{E} [x] . \mathbb{E} [y] = \mathbb{E}_{x,y} [xy^T] - \mathbb{E} [x] . \mathbb{E} [y^T]$$

### 3 Probabilities

- covariance matrix

#### 3.1.3 Bayes Rule

$$p(x_i|y_i) \cdot p(y_i) = p(y_i|x_i) \cdot p(x_i) = p(x_i, y_i)$$

$$\Rightarrow p(y_i|x_i) = \frac{p(x_i|y_i) \cdot p(y_i)}{p(x_i)} = \frac{p(x_i|y_i) \cdot p(y_i)}{\sum_y p(x_i|y_i) \cdot p(y_i)}$$

$\Rightarrow$  the ***Bayes equation:***

$$\text{posterior} = \frac{\text{likelihood} \times \text{prior}}{\text{normalization factor}}$$

#### 3.1.4 Expectation

For variable $x$ :	$\mathbb{E}[x] = \sum_x x \cdot p(x)$	$\left( = \int x \cdot p(x) dx \right)$
For function $f(\cdot)$ :	$\mathbb{E}[f(x)] = \sum_x f(x) \cdot p(x)$	$\left( = \int f(x) \cdot p(x) dx \right)$

## 3.2 Types of Probability Distributions

Reference source: [machinelearningcoban.com](http://machinelearningcoban.com).

#### 3.2.1 Bernoulli Distribution

Bernoulli Distribution is a distribution to describe binary discrete variables. It's the case that the variable can only take value in 2 classes  $x \in \{0, 1\}$ . E.g., the probability of throwing a coin. The Bernoulli distribution is defined with parameter  $\lambda \in [0, 1]$ :

$$p(x) = \text{Bern}_x[\lambda] = \begin{cases} p(x=1) = \lambda \\ p(x=0) = 1 - \lambda \end{cases} \quad (3.1)$$

In short form, the above equation can be combined into one:

$$p(x) = \lambda^x (1 - \lambda)^{1-x} \Rightarrow \begin{cases} p(0) = \lambda^0 (1 - \lambda)^1 = 1 - \lambda \\ p(1) = \lambda^1 (1 - \lambda)^0 = \lambda \end{cases} \quad (3.2)$$

### 3.2.2 Categorical Distribution

*Categorical Distribution* is the generalization of *Bernoulli Distribution*, in case there are  $K$  classes for the discrete variable  $x \in \{1, 2, \dots, K\}$ . Accordingly, there will be  $K$  parameters to describe this PDF:  $\lambda = [\lambda_1, \lambda_2, \dots, \lambda_K]$ , with  $\lambda_k \geq 0$  and  $\sum \lambda_k = 1$ . Each  $\lambda_k$  represents the probability to take the output  $k$ :  $p(x = k) = \lambda_k$ . In short:  $p(x) = \text{Cat}_x[\lambda]$ .

Another common way to represent the output is the one-hot vector,  $\mathbf{x} \in \{\mathbf{e}_1, \mathbf{e}_2, \dots, \mathbf{e}_K\}$  with  $\mathbf{e}_k$  is the  $k$ -unit vector, which has all 0-element, except the  $k$ -element equal to 1. E.g., given 3 classes:  $\mathbf{e}_1 = [1, 0, 0]^T, \mathbf{e}_2 = [0, 1, 0]^T, \mathbf{e}_3 = [0, 0, 1]^T$ . We will then have:

$$p(\mathbf{x} = \mathbf{e}_k) = \prod_{j=1}^K \lambda_j^{x_j} = \lambda_k \quad (3.3)$$

because for  $\mathbf{x} = \mathbf{e}_k$ , only  $x_k = 1$ , while  $x_j = 0, \forall j \neq k$ .

### 3.2.3 Univariate Normal Distribution

Univariate Normal Distribution is also known as the Gaussian distribution. For single dimension data (in 1D):  $x \in (-\infty, \infty)$ , the mean  $\mu \in \mathbb{R}$ , and the variance  $\sigma^2$  with  $\sigma \in \mathbb{R}$ .

$$p(x) = \text{Norm}_x [\mu, \sigma^2] = \mathcal{N}(\mu, \sigma^2) = \frac{1}{\sqrt{2\pi\sigma^2}} \cdot \exp\left(-\frac{(x - \mu)^2}{2\sigma^2}\right) \quad (3.4)$$

#### NOTE:

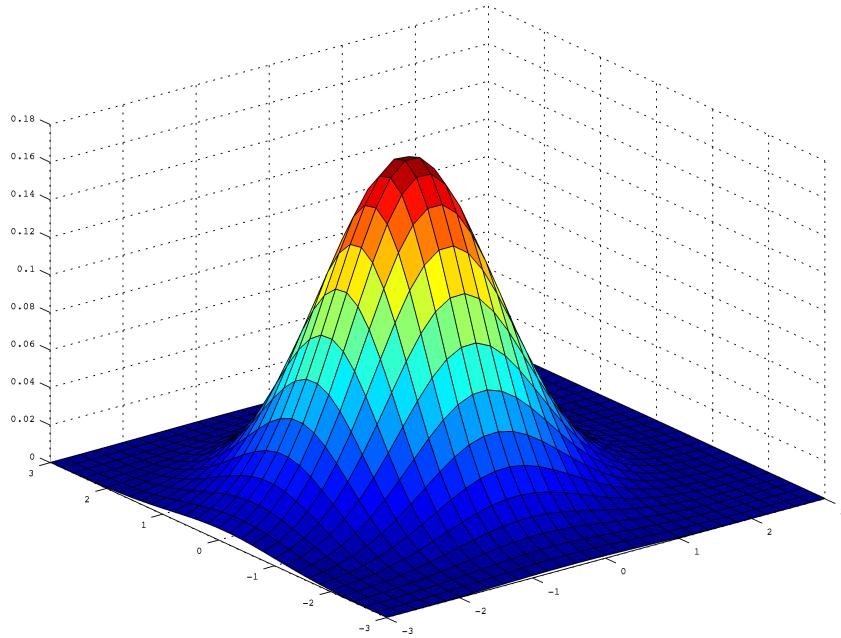
- Marginals prob. of Gaussian are again Gaussian.
- When estimating the params. of a Gaussian, beware the underestimation problem.

$$\begin{aligned} \mathbb{E}[\mu_{ML}] &= \mu \\ \mathbb{E}[\sigma_{ML}^2] &= \left(\frac{N-1}{N}\right)\sigma^2 \\ \Rightarrow \tilde{\sigma}^2 &= \left(\frac{N}{N-1}\right)\sigma_{ML}^2 = \frac{1}{N-1} \sum_{n=1}^N (x_n - \hat{\mu})^2 \end{aligned}$$

### 3.2.4 Multivariate Normal Distribution

*Multivariate Normal Distribution* is the extension of *Univariate Normal Distribution* to multi-dimensional data:  $\mathbf{x}, \boldsymbol{\mu} \in \mathbb{R}^D, \sigma^2 \Rightarrow \Sigma \in \mathbb{S}_{++}^D$  ( $\mathbb{S}_{++}^D$  is the set of positive definite

### 3 Probabilities



**Figure 3.1:** Bivariate Gaussian distribution ([src](#)).

symmetric matrix)

$$p(x) = \text{Norm}_x[\boldsymbol{\mu}, \Sigma] = \mathcal{N}(\boldsymbol{\mu}, \Sigma) = \frac{1}{2\pi^{D/2}|\Sigma|^{\frac{1}{2}}}\exp\left(-\frac{1}{2}(\mathbf{x} - \boldsymbol{\mu})^T\Sigma^{-1}(\mathbf{x} - \boldsymbol{\mu})\right) \quad (3.5)$$

#### 3.2.5 Beta Distribution

This distribution describes the parameter for another distributions. E.g., Dirichlet [PDF](#) describes Categorical Distribution (Subsec. [3.2.2](#))

### 3.3 Parameter Estimation

Many of [ML](#) problems are boiled down to finding *statistical models*. Those models could predict the [prob.](#) for the classification problem, [prob.](#) of events that will happen, etc. It all end up with finding the suitable set of [params.](#) for these *statistical models*.

### 3.3.1 Maximum Likelihood Estimation

Maximum Likelihood Estimation (**MLE**) finds the parameters that maximize the **prob.** of the existing data.

$$\theta = \operatorname{argmax}_{\theta} p(x_1, x_2, \dots, x_N | \theta) \quad (3.6)$$

Assuming independent variables:

$$\theta = \operatorname{argmax}_{\theta} \prod_{n=1}^N p(x_n | \theta) \quad (3.7)$$

Maximum log-likelihood:

$$\theta = \operatorname{argmax}_{\theta} \sum_{n=1}^N [\log p(x_n | \theta)] \quad (3.8)$$

Minimum negative log-likelihood:

$$\theta = \operatorname{argmin}_{\theta} \sum_{n=1}^N [-\log p(x_n | \theta)] \quad (3.9)$$

### 3.3.2 Maximum A Posteriori

Sometimes, we have prior knowledge of the **PDF**. E.g., we know that the **prob.** of getting head when flipping a coin is around 50%. Maximum A Posteriori (**MAP**) takes advantage of the prior knowledge  $p(\theta)$  on the parameters  $\theta$  by applying Bayes rule (Subsec. 3.1.3)

$$\theta = \operatorname{argmax}_{\theta} \prod_{n=1}^N p(x_n | \theta) p(\theta) \quad (3.10)$$

**MLE suffers when there is not enough data  $\Rightarrow$  use MAP**

## 3.4 Naive Bayes Classifier

Naive implies having the independence assumption on the variables.

$$\begin{aligned} c &= \operatorname{argmax}_{c \in \mathcal{C}} p(c|x) \\ &= \operatorname{argmax}_{c \in \mathcal{C}} p(x|c) p(c) \end{aligned}$$

If  $x$  is:

- continuous variable  $\Rightarrow$  Gaussian Naive Bayes
- feature vector  $\Rightarrow$  Multinomial Naive Bayes
- binary vector  $\Rightarrow$  Bernoulli Naive Bayes

### 3 Probabilities

Minimize the expected loss:  $\mathbb{E}[L] = \sum_k \sum_j \int_{R_j} L_{kj} p(x, C_k) dx$  by choosing region  $R_j$  such that  $\mathbb{E}[L] = \sum_k L_{kj} p(C_k | x)$

## 3.5 Views on the Decision Problem

### 3.5.1 Generative Methods

First determine the class-conditional densities and separately infer the prior class prob.  
⇒ Bayes theorem ⇒ class membership

$$p(x|C_k) p(C_k) \Rightarrow y_k(x)$$

E.g., Mixture of Gaussians

### 3.5.2 Discriminative Methods

First solve the inference problem of determined the posterior class prob.

## 3.6 Unknown Notes

E.g., 2 class  $C_1, C_2$ , 2 decisions  $\alpha_1, \alpha_2$ .

The loss:  $L(\alpha_j | C_k) = L_{kj}$ .

The expected loss is equal to the  $Risk(R)$ .

$$\begin{aligned}\mathbb{E}_{\alpha_1}[L] &= R(\alpha_1 | x) = L_{11} p(C_1 | x) + L_{21} p(C_2 | x) \\ \mathbb{E}_{\alpha_2}[L] &= R(\alpha_2 | x) = L_{12} p(C_1 | x) + L_{22} p(C_2 | x)\end{aligned}$$

Choose  $\alpha_1$  if  $R(\alpha_1 | x) < R(\alpha_2 | x)$

## 3.7 Probability Density Estimation

### 3.7.1 Histogram

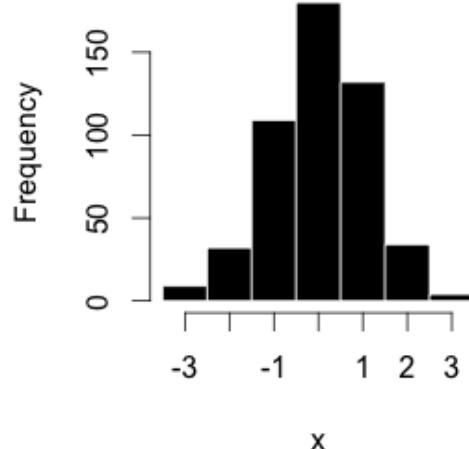
This is **non-parametric** prob. density estimation. All other approaches are **parametric**.  
The prob. of a bin:

$$p_i = \frac{n_i}{N \cdot \Delta_i} \quad (3.11)$$

in which  $n_i$  is the number of data points in that bin,  $N$  is the total number of data point,  $\Delta_i$  is the width of the bin, often  $\Delta_i = \Delta$ .

#### Notes:

- $\Delta$  serves as the **smoothing factor**
- With  $D$  as the dimensions of the data points. The number of bins grow exponentially with  $\mathcal{O}(k^D)$



**Figure 3.2:** Example of a histogram,  $\Delta = 1$  ([src](#)).

### 3.7.2 Parametric Probability Density Estimation

In the other hands, one could find the prob. from  $p = \int_{\mathcal{R}} p(y)dy \approx p(x)V$  where the region  $\mathcal{R}$  is sufficiently small.

$\Rightarrow p(x) \approx \frac{K}{N \cdot V}$ , where  $K$  is the number of data points in the region,  $V$  is the volume of the region.

### 3 Probabilities

#### 3.7.3 Kernel Methods

The kernel methods fix  $V$  and determine  $K$ . The volume  $V$  is the space restricted within a parzen window  $k(u)$  that satisfies  $k(u) \geq 0$ .

A hyper-space cube: 
$$k(u) = \begin{cases} 1 & \text{if } |u_i| \leq \frac{1}{2}h, \quad i = 1, 2, \dots, D \\ 0 & \text{else} \end{cases} \quad (3.12)$$

The number of points inside: 
$$K = \sum_{n=1}^N k(x - x_n) \quad (3.13)$$

The region *volume*: 
$$V = \int k(u)du = h^D \quad (3.14)$$

The probability: 
$$\Rightarrow p(x) \approx \frac{K}{N.V} = \frac{1}{N.h^D} \sum_{n=1}^N k(x - x_n) \quad (3.15)$$

The **symmetric Gaussian kernel is a better substitution** for the asymmetric parzen window.

A Gaussian kernel: 
$$k(u) = \frac{1}{\sqrt{2\pi h^2}} \exp\left(\frac{-u^2}{2h^2}\right) \quad (3.16)$$

The region *volume*: 
$$V = \int k(u)du = 1 \quad (3.17)$$

The probability: 
$$\Rightarrow p(x) \approx \frac{1}{N} \sum_{n=1}^N \frac{1}{(2\pi)^{D/2} h} \exp\left(\frac{-||x - x_n||^2}{2h^2}\right) \quad (3.18)$$

For Kernel methods,  $h$  is the **smoothing factor**.

Generalization:  $k(u) \geq 0, \int k(u)du = 1$ .

**Size of the hypersphere is proportional to  $h^2$ .**

#### 3.7.4 K-Nearest Neighbor

When you fix  $K$  and determine  $V$ , it leads to K-Nearest Neighbor.

[TODO: Add image]

$$p(x) \approx \frac{K}{NV} \quad (3.19)$$

Here,  $K$  is the **smoothing factor**.

- Too much bias  $\Rightarrow$  too smooth
- Too much variance  $\Rightarrow$  NOT smooth enough

$\Rightarrow$  combine parametric methods to a mixture model

Mixture distribution = multi parametric model

### 3.7.5 Mixture of Gaussians

Mixture of Gaussians (MoG), as **Generative Model**, is defined from the prob. sum of elemental Gaussians:  $p(x|\theta) = \sum_{j=1}^M p(x|\theta_j)p(j)$ , where  $p(x|\theta_j)$  is a **mixture component**,  $p(j) = \pi_j$  is the **weight of the component**

$$p(x|\theta_j) = \frac{1}{\sqrt{2\pi}\sigma_j} \exp\left[\frac{-(x - \mu_j)^2}{2\sigma_j^2}\right], \quad p(j) = \pi_j, \quad \sum \pi_j = 1 \quad (3.20)$$

$$p(x|\theta_j) = \frac{1}{(2\pi)^{\frac{D}{2}} |\Sigma_j|^{\frac{1}{2}}} \exp\left[-\frac{1}{2}(x - \mu_j)^T \Sigma_j^{-1} (x - \mu_j)\right] \quad (3.21)$$

### 3.7.6 K-Means Clustering

There are 3 steps:

- Pick  $K$  centroids
- Assign sample to the centroid
- Adjust centroids

Step 2 and 3 are repeated until there is no change.

This leads to a local optimum, depends on initialization. It's sensitive to outliers, detects **spherical clusters only**.

[TODO: Add images]

Application: e.g., image compression.

### 3 Probabilities

#### 3.7.7 EM Clustering

It's short for Expectation-Maximization. Assuming  $N$  data points and  $K$  Gaussians.

- **E-Step:** Fix the Gaussians, find  $\gamma_j(x)$ , which represent the **responsibility of component  $j$  for  $x$** .

$$\gamma_j(x_n) = \frac{\pi_j \mathcal{N}(x_n | \mu_j, \Sigma_j)}{\sum_{k=1}^K \pi_k \mathcal{N}(x_n | \mu_k, \Sigma_k)} \quad \forall j = 1, 2, \dots, K, \quad n = 1, 2, \dots, N \quad (3.22)$$

- **M-Step:** Fix  $\gamma_j(x)$ , update the Gaussians.

$$\hat{N}_j = \sum_{n=1}^N \gamma_j(x_n) \quad (3.23)$$

$$\hat{\mu}_j = \frac{1}{\hat{N}_j} \sum_{n=1}^N \gamma_j(x_n) x_n \quad (3.24)$$

$$\hat{\pi}_j = \frac{\hat{N}_j}{N} \quad (3.25)$$

$$\hat{\Sigma}_j = \frac{1}{\hat{N}_j} \sum_{n=1}^N \gamma_j(x_n) (x_n - \hat{\mu}_j) (x_n - \hat{\mu}_j)^T \quad (3.26)$$

#### Notes:

- Regularization with  $\Sigma + \sigma_{min} I$
- Initialization  $\mu_j$  with K-Means
- Hard-assignment: each data point to 1 class  $\Rightarrow$  K-Means
- Soft-assignment: each data point  $\Rightarrow$  prob. to fall into many classes  $\Rightarrow$  EM Clustering

EM **needs more iteration**, because there are **more params..**

# 4 Linear Regression

This is the simplest regression problem in ML.

**Problem statement:** Given data points  $\mathbf{x}_i \in \mathbb{R}^D$  and their labels  $y_i \in \mathbb{R}$ , find the "line" that fit these data points. The line is represented via parameters  $\mathbf{w}$ . For each data point  $\mathbf{x}$  and its label  $y$

$$\begin{aligned}\mathbf{w} &= [w_0, w_1, \dots, w_n]^T \\ \bar{\mathbf{x}} &= [1, x_0, \dots, x_n] \quad (\text{x bar}) \\ y &\approx \hat{y} = \bar{\mathbf{x}} \cdot \mathbf{w} \quad (\text{y hat}) \\ \Rightarrow \frac{1}{2} e^2 &= \frac{1}{2} (y - \bar{\mathbf{x}} \cdot \mathbf{w})^2\end{aligned}$$

The loss function for all points

$$\mathcal{L}(\mathbf{w}) = \frac{1}{2} \sum_{i=1}^N (y_i - \bar{\mathbf{x}}_i \cdot \mathbf{w})^2 \tag{4.1}$$

The average loss:

$$\mathcal{L}(\mathbf{w}) = \frac{1}{2N} \sum_{i=1}^N (y_i - \bar{\mathbf{x}}_i \cdot \mathbf{w})^2 \tag{4.2}$$

We need to find the weights that minimize the loss function

$$\mathbf{w}^* = \underset{\mathbf{w}}{\operatorname{argmin}} \mathcal{L}(\mathbf{w}) \tag{4.3}$$

We now can write the loss function using matrix form:

$$\mathcal{L}(\mathbf{w}) = \frac{1}{2} \|\mathbf{y} - \bar{\mathbf{X}} \cdot \mathbf{w}\|_2^2 \tag{4.4}$$

$$\text{with } \bar{\mathbf{X}} = \begin{bmatrix} \bar{\mathbf{x}}_1 \\ \bar{\mathbf{x}}_2 \\ \vdots \\ \bar{\mathbf{x}}_n \end{bmatrix} \text{ and } \mathbf{y} = \begin{bmatrix} y_1 \\ y_2 \\ \vdots \\ y_n \end{bmatrix}$$

## 4 Linear Regression

**Solution:**

$$\frac{\partial \mathcal{L}(\mathbf{w})}{\partial \mathbf{w}} = \bar{\mathbf{X}}^T (\bar{\mathbf{X}}\mathbf{w} - \mathbf{y}) = 0 \quad (4.5)$$

$$\iff \bar{\mathbf{X}}^T \bar{\mathbf{X}}\mathbf{w} = \bar{\mathbf{X}}^T \mathbf{y} \quad (4.6)$$

$$\iff \mathbf{w} = (\bar{\mathbf{X}}^T \bar{\mathbf{X}})^{\dagger} \bar{\mathbf{X}}^T \mathbf{y} \quad (4.7)$$

in which,  $A^{\dagger}$  (A dagger) is the pseudo inverse of a matrix, because it's might not inverse-able.  $\mathbf{A}^{\dagger} = (\mathbf{A}^T \mathbf{A})^{-1} \mathbf{A}^T$

**NOTE:** Sensitive to outliers  $\Rightarrow$  pre-processing

Multi-variable:

$$\begin{aligned} \mathcal{L}(\mathbf{w}) &= \frac{1}{2N} \|\mathbf{y} - \bar{\mathbf{X}}\mathbf{w}\|_2^2 \\ \Rightarrow \nabla_{\mathbf{w}} \mathcal{L}(\mathbf{w}) &= \frac{1}{N} \bar{\mathbf{X}}^T (\bar{\mathbf{X}}\mathbf{w} - \mathbf{y}) \end{aligned}$$

Linear Discriminant Functions:

$$\begin{aligned} y(\mathbf{x}) &= \widetilde{\mathbf{W}}^T \tilde{\mathbf{x}} \\ E(\mathbf{w}) &= \frac{1}{2} \sum_{n=1}^N \sum_{k=1}^K \left( \mathbf{w}_k^T \mathbf{x}_n - t_{kn} \right)^2 \\ \widetilde{\mathbf{W}} &= \left( \widetilde{\mathbf{X}}^T \widetilde{\mathbf{X}} \right)^{-1} \widetilde{\mathbf{X}}^T T \\ &= \widetilde{\mathbf{X}}^{\dagger} T \end{aligned}$$

Generalized Discriminant:

$$y_k(x) = \sum_{j=1}^M w_{kj} \phi_j(x) + w_{k0} = \sum_{j=0}^M w_{kj} \phi_j(x), \quad \phi_0(x) = 1 \quad (4.8)$$

# 5 Logistics Regression

Though named as "regression", logistics regression is used for classification problem. For the classification problem with 2 classes, instead of a hard threshold (from linear classification), we could have a soft one, and find the **prob.** of the input belongs to either class.

$$f(x) = \theta(\omega^T x) \quad (5.1)$$

The sigmoid function:

$$f(s) = \frac{1}{1 + e^{-s}} \triangleq \sigma(s) \quad (5.2)$$

$$s = \ln\left(\frac{\sigma}{1 - \sigma}\right) \quad (5.3)$$

$$\sigma'(s) = \sigma(s)(1 - \sigma(s)) \quad (5.4)$$

The cross entropy error function:

$$J(w, x, y) = -(y_i \log z_i + (1 - y_i) \log(1 - z_i)) \quad (5.5)$$

with  $z_i = f(w^T x_i)$

$$\Rightarrow \frac{\partial J}{\partial w} = (z_i - y_i)x_i \quad (5.6)$$

$$\Rightarrow w = w + \eta(y_i - z_i)x_i \quad (5.7)$$

**NOTE:** Require less **params.**, only  $D$  dimensions, compared to Gaussians with

$$\frac{M(M+5)}{2} + 1 \text{ params.}$$

# 6 Softmax Regression

Multinomial Logistics Regression, Maximum Entropy Classifier

$$a_i = \frac{\exp(z_i)}{\sum_{j=1}^C \exp(z_j)} \quad (6.1)$$

$$\text{so that } \begin{cases} a_i > 0 \\ \sum a_i = 1 \\ z_m > z_n \iff a_m > a_n \text{ (order)} \end{cases}$$

When  $z_i$  is too big

$$\frac{\exp(z_i)}{\sum_{j=1}^C \exp(z_j)} = \frac{\exp(z_i - c)}{\sum_{j=1}^C \exp(z_j - c)} \quad (6.2)$$

with  $c = \max_i z_i$

$$J(w, x, y) = - \sum_{i=1}^N \sum_{j=1}^C y_{ij} \log(a_{ij}) \quad (6.3)$$

$$J(w, x, y) = - \sum_{n=1}^N \sum_{k=0}^1 [\mathbb{I}(t_n = k) \ln p(y_n = k | x_n; w)] \quad (6.4)$$

$$\Rightarrow E(w) = - \sum_{n=1}^N \sum_{k=1}^K \left[ \mathbb{I}(t_n = k) \ln \frac{\exp(w_k^T x)}{\sum_{j=1}^K \exp(w_j^T x)} \right] \quad (6.5)$$

$$\frac{\partial J_i(w)}{\partial W} = x_i e_i^T = x_i (a_{ij} - y_{ij})^T \quad (6.6)$$

$$W = W + \eta x_i (y_i - a_i)^T \quad (6.7)$$

# 7 Error Functions

As mention in Sec. 2.4, the loss function is the metric to assess the performance of the ML model with a certain task. Different loss functions fit with different tasks. Sometimes, the design of the loss function is what makes the differences.

[TODO: Add graph and explanation]

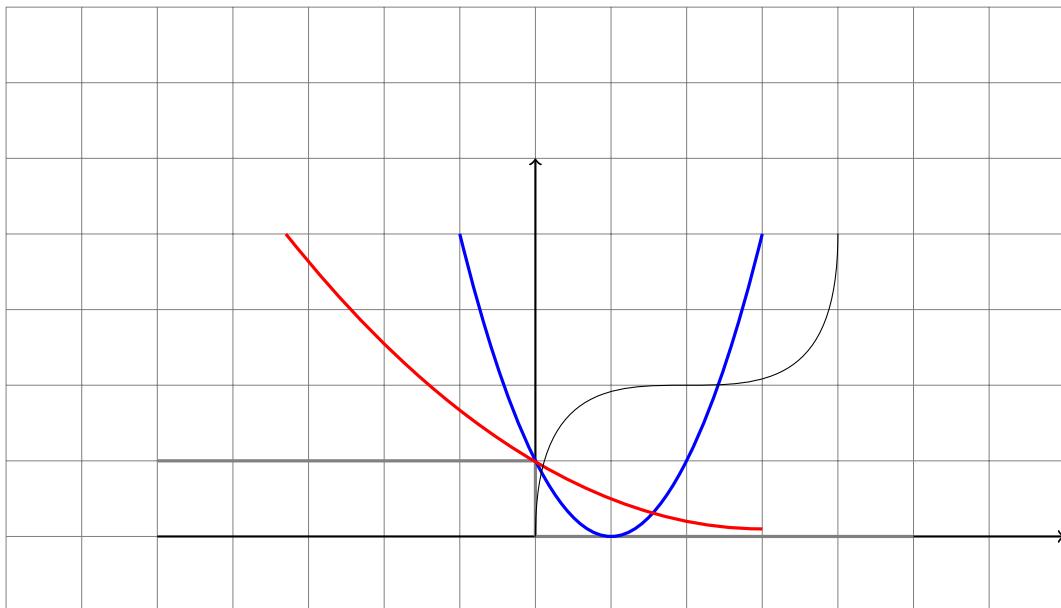


Figure 7.1: Different error functions: Ideal Miss-classification Error (gray line)

## 7.1 Ideal Miss-classification Error

Gradient = 0  $\Rightarrow$  can't use gradient descent.

It simply counts incorrectly classified points.

## 7.2 Squared Error - $L_2$ Loss

- Leads to closed form solutions
- Sensitive to outliers
- Penalize "too correct" data points

### 7.3 Cross Entropy Error

- Concave function  $\Rightarrow$  unique minimum exists
- Robust to outliers, error increases only roughly linear
- No closed-form solution, requires iterative method

### 7.4 Squared Error on Sigmoid / Tanh

- No penalty for "too correct" points
- Zero gradient for confidently incorrect classifications

$\Rightarrow$  **Do NOT** use  $L_2$  loss with sigmoid outputs, instead, use cross-entropy.

### 7.5 Hinge Error

- Robust to outliers
- Zero error for points outside margin  $\Rightarrow$  sparsity
- Not differentiable around  $z_n = 1$

**NOTE:** Want the correct class to have a score that is higher than incorrect class by a fixed margin  $\Delta$ .

$$L_i = \sum_{j \neq y_i} \max(0, s_j - s_{y_i} + \Delta) \quad (7.1)$$

in which,  $s_j$  is other classes score,  $s_{y_i}$  is real class score.

### 7.6 $L_1, L_0$ Loss

Median, no wrong points

$$L_1 = \sum |t - y| \quad (7.2)$$

$$L_2 = \sum (t - y)^2 \quad (7.3)$$

## **7.7 Average Loss**

Mathematically, dividing the loss by the amount of data  $N$  (in each batch or epoch) doesn't have any effect on the result. However, it's usually advisable to take the average to have more meaningful judgment and avoiding overflow when there are numerous data points.

# 8 Neural Network

## 8.1 Gradient Descent

NOTE: Just use ADAM??

### 8.1.1 Vanilla Gradient Descent

$$\theta_{t+1} = \theta_t - \eta \nabla_{\theta} f(\theta_t) \quad (8.1)$$

Check derivative:

$$f'(x) \approx \frac{f(x + \varepsilon) - f(x - \varepsilon)}{2\varepsilon} \quad (\text{numerical gradient}) \quad (8.2)$$

### 8.1.2 Momentum

- Init:  $v_{dW_0} = 0, v_{db_0} = 0$
- Calculate  $dW, db$
- Update  $W, b$

$$\Rightarrow \begin{cases} v_{dW} &= \beta v_{dW} + (1 - \beta)dW \\ v_{db} &= \beta v_{db} + (1 - \beta)db \end{cases} \Rightarrow \begin{cases} W &= W - \alpha v_{dW} \\ b &= b - \alpha v_{db} \end{cases} \quad (8.3)$$

The above formulas are to calculate the moving average of  $v_{dW}$  and  $v_{db}$ .

- Tips: Choose  $\beta_1 = 0.9$ , implying taking average of the last 10 steps.
- Reference source: [DeepLearning.AI](#).

### 8.1.3 Nesterov Accelerated Gradient

Nestorov Accelerated Gradient ([NAG](#)):

$$v_t = \gamma v_{t-1} + \eta \nabla_t J(\theta - \gamma v_{t-1}) \quad (8.4)$$

### 8.1.4 RMSprop

- Init  $s_{dW_0} = 0, s_{db_0} = 0$
- Calculate  $dW, db$
- Update  $W, b$

$$\begin{cases} s_{dW} &= \beta s_{dW} + (1 - \beta) dW^2 \\ s_{db} &= \beta s_{db} + (1 - \beta) db^2 \end{cases} \Rightarrow \begin{cases} W &= W - \alpha \frac{dW}{\sqrt{s_{dW}} + \varepsilon} \\ b &= b - \alpha \frac{db}{\sqrt{s_{db}} + \varepsilon} \end{cases} \quad (8.5)$$

- Tips: choose  $\beta_2 = 0.999, \varepsilon = 10^{-7}$
- Reference source: [DeepLearning.AI](#).

### 8.1.5 Adam

[Adam](#) is basically the combination of Momentum and RMSprop.

- Init  $v_{dW_0}, s_{dW_0}, v_{db_0}, s_{db_0} = 0$
- Calculate  $dW, db$
- Update  $W, b$

$$\begin{cases} v_{dW} &= \beta_1 v_{dW} + (1 - \beta_1) dW \\ v_{db} &= \beta_1 v_{db} + (1 - \beta_1) db \\ s_{dW} &= \beta_2 s_{dW} + (1 - \beta_2) dW^2 \\ s_{db} &= \beta_2 s_{db} + (1 - \beta_2) db^2 \end{cases} \Rightarrow \begin{cases} v_{dW}^{cor.} &= \frac{v_{dW}}{1 - \beta_1^t} \\ v_{db}^{cor.} &= \frac{v_{db}}{1 - \beta_1^t} \\ s_{dW}^{cor.} &= \frac{s_{dW}}{1 - \beta_2^t} \\ s_{db}^{cor.} &= \frac{s_{db}}{1 - \beta_2^t} \end{cases} \Rightarrow \begin{cases} W &= W - \alpha \frac{v_{dW}^{cor.}}{\sqrt{s_{dW}^{cor.}} + \varepsilon} \\ b &= b - \alpha \frac{v_{db}^{cor.}}{\sqrt{s_{db}^{cor.}} + \varepsilon} \end{cases} \quad (8.6)$$

- Tips: choose  $\beta_1 = 0.9, \beta_2 = 0.999, \varepsilon = 10^{-7}$
- Reference source: [DeepLearning.AI](#).

## 8.2 Tips and Tricks

- Shuffling
- Data Augmentation: reshape, rescale, crops, zooming, change color (color Principal Component Analysis ([PCA](#)))
- Normalizing the inputs

Convergence is the fastest if

- The mean of each input variable = 0
- Scale  $\Rightarrow$  same covariance

Mean cancellation  $\Rightarrow$  Kullback–Leibler ([K-L](#)) expansion  $\Rightarrow$  covariance equalization (if possible)

- Leaky Rectified Linear Unit ([ReLU](#)) is better a bit than [ReLU](#), ELU
- Weights initialization: Xavier-Glorot:

$$W \sim U \left( 0, \sqrt{\frac{6}{n_{in} + n_{out}}} \right)$$

- Batch Norm(alization): Normalize after each layer  
 $\Rightarrow$  learn the moving average
- Drop out

**NOTE:** When in inferencing (after training), must multiply the activation output with the [prob.](#) that the weights are set to 0

# Bibliography

- [GBC16] I. Goodfellow, Y. Bengio, and A. Courville. *Deep learning*. MIT press, 2016.
- [Vu18] H. T. Vu. *Deep learning*. 2018.