

AI Notes

Huu Duc Nguyen M.Sc.

29 February 2022

Contents

Abbreviations	1
1 Introduction	2
2 Overview of Machine Learning	4
2.1 Task T	4
2.1.1 Classification	4
2.1.2 Regression	4
2.2 Experience E	4
2.3 Performance P	4
3 Probabilities	5
3.1 Definitions	5
3.1.1 Basic Definitions	5
3.1.2 Independence and Variability	5
3.1.3 Bayes Rule	6
3.1.4 Expectation	6
3.2 Types of Probability Distributions	6
3.2.1 Bernoulli Distribution	6
3.2.2 Categorical Distribution	7
3.2.3 Univariate Normal Distribution	7
3.2.4 Multivariate Normal Distribution	7
3.2.5 Beta Distribution	8
3.3 Parameter Estimation	8
3.3.1 Maximum Likelihood Estimation	9
3.3.2 Maximum A Posteriori	9
3.4 Naive Bayes Classifier	9
3.5 Views on the Decision Problem	10
3.5.1 Generative Methods	10
3.5.2 Discriminative Methods	10
3.6 Unknown Notes	10
3.7 Probability Density Estimation	11
3.7.1 Histogram	11
3.7.2 Parametric Probability Density Estimation	11

3.7.3	Kernel Methods	12
3.7.4	K-Nearest Neighbor	12
3.7.5	Mixture of Gaussians	13
3.7.6	K-Means Clustering	13
3.7.7	EM Clustering	14
4	Linear Regression	15
5	Logistics Regression	17
6	Softmax Regression	18
7	Error Functions	19
7.1	Ideal Miss-classification Error	19
7.2	Squared Error - L_2 Loss	19
7.3	Cross Entropy Error	19
7.4	Squared Error on Sigmoid / Tanh	19
7.5	Hinge Error	19
7.6	L_1, L_0 Loss	20
8	Neural Network	21
8.1	Gradient Descent	21
8.1.1	Vanilla Gradient Descent	21
8.1.2	Momentum	21
8.1.3	Nesterov Accelerated Gradient	21
8.1.4	Root mean squared prop (RMSprop)	22
8.1.5	Adaptive moment estimation (Adam)	22
8.2	Tips and Tricks	23
Bibliography		I

Abbreviations

General Terms

AI	Artificial Intelligence
ML	Machine Learning
DL	Deep Learning
CS	Computer Science
CV	Computer Vision
NLP	Natural Language Processing

Convenient Abbreviations

prob.	probability
params.	parameters

Probabilities

PDF	Probability Density Function
MLE	Maximum Likelihood Estimation
MAP	Maximum A Posteriori

Machine Learning related terms

MoG	Mixture of Gaussians
NAG	Nestorov Accelerated Gradient
RMSprop	Root mean squared prop
Adam	Adaptive moment estimation
PCA	Principal Component Analysis
K-L	Kullback–Leibler
ReLU	Rectified Linear Unit

1 Introduction

Artificial Intelligence ([AI](#)) is the study field that leverages the ability of machines to mimic the problem-solving skill of human. It lies in the core of countless novel applications in real life, self-driving cars, virtual assistant, face recognition, etc. Machine Learning ([ML](#)) is a sub-field of Computer Science ([CS](#)) and [AI](#), that “gives computers the ability to learn without being explicitly programmed” ([Wikipedia](#)). As a great amount of collected data and powerful computational hardware arise, Deep Learning ([DL](#)) is then a subset of [ML](#) (Fig. 3.1). Advanced applications which relate to Natural Language Processing ([NLP](#)), Computer Vision ([CV](#)), robotic learning, etc., are with in this [DL](#) subset.




Figure 1.1: The relation between [AI](#), [ML](#) and [DL](#) ([src](#)).

[TODO: The structure of the notes]

It's important to understand that there are more to [AI](#) and [ML](#) than just neural networks. In the end, to create a meaningful and working neural network, I believe one should have strong background in the basics of [ML](#) as well. Henceforth, I keep record of what I have learnt in the whole [AI](#) field. The structure of the notes is as follows:

- Chap. 3 introduces the mathematics background on probabilities, matrix.

- [TODO: chapter 3] explain basic concepts, the branching of different classes in ML. Later chapters presents each smaller branches.

2 Overview of Machine Learning

A machine learning algorithm is an algorithm that has the ability to *learn* from the data. A computer program is said to **learn**, if its performance at tasks in T , measured by P , improves with experience E (in which the experience is equivalent to the data). [GBC16]

2.1 Task T

A *task* is usually described by how the **ML** model process a single *data point*. This section presents some common **ML** tasks. [Vu18]

2.1.1 Classification

The task is to specify a label for the given data point. The labels are usually members of a list.

E.g., in the problem of digit classification, the data point is images of hand-written numbers. The data set comes with their labels as well. The task is then, given a unseen image, the model would be able to tell which number is in that image. In this problem, there are 10 possible labels, i.e., $0, 1, \dots, 9$.

2.1.2 Regression

2.2 Experience E

2.3 Performance P

3 Probabilities

3.1 Definitions

3.1.1 Basic Definitions

- If x is discrete: $\sum_x p(x) = 1$ with $\forall 0 \leq p(x) \leq 1$
- If x is continuous: $\int p(x) dx = 1 \Rightarrow \exists$ a **Probability Density Function (PDF)**
 $p(x)$ can take any positive value, as long as $\int p(x) dx = 1$
[TODO: Add image]
NOTE: : theoretically $p(x) = 0, \forall x$
- Common types

Joint probability: $p(x_i, y_i)$ ($= p(X = x_i, Y = y_i)$)

Marginal probability: $p(x_i)$ ($= p(X = x_i)$)

Conditional probability: $p(y_i|x_i)$ ($= p(Y = y_i|X = x_i)$)

- Sum rule: \sum joint probability (prob.) = marginal prob.
 \Rightarrow Marginalization
 - discrete variable: $p(x) = \sum_y p(x, y)$
 - continuous variable: $p(x) = \int p(x, y) dy$
- Product rule: Product of marginal prob. and conditional prob. = joint prob.

3.1.2 Independence and Variability

- Independence. E.g.: x, y are independent, then

$$\begin{cases} p(x|y) = p(x) \\ p(y|x) = p(y) \end{cases} \iff p(x, y) = p(x).p(y)$$

- Variability
 - variance:
$$var [f] = \mathbb{E} [(f(x) - \mathbb{E} [f(x)])^2] = \mathbb{E} [f(x)^2] - \mathbb{E} [f(x)]^2$$
 - covariance:
$$cov [x, y] = \mathbb{E}_{x,y} [xy] - \mathbb{E} [x] . \mathbb{E} [y] = \mathbb{E}_{x,y} [xy^T] - \mathbb{E} [x] . \mathbb{E} [y^T]$$

3 Probabilities

- covariance matrix

3.1.3 Bayes Rule

$$p(x_i|y_i) \cdot p(y_i) = p(y_i|x_i) \cdot p(x_i) = p(x_i, y_i)$$

$$\Rightarrow p(y_i|x_i) = \frac{p(x_i|y_i) \cdot p(y_i)}{p(x_i)} = \frac{p(x_i|y_i) \cdot p(y_i)}{\sum_y p(x_i|y_i) \cdot p(y_i)}$$

\Rightarrow the ***Bayes equation:***

$$\text{posterior} = \frac{\text{likelihood} \times \text{prior}}{\text{normalization factor}}$$

3.1.4 Expectation

For variable x :	$\mathbb{E}[x] = \sum_x x \cdot p(x)$	$\left(= \int x \cdot p(x) dx \right)$
For function $f(\cdot)$:	$\mathbb{E}[f(x)] = \sum_x f(x) \cdot p(x)$	$\left(= \int f(x) \cdot p(x) dx \right)$

3.2 Types of Probability Distributions

Reference source: machinelearningcoban.com.

3.2.1 Bernoulli Distribution

Bernoulli Distribution is a distribution to describe binary discrete variables. It's the case that the variable can only take value in 2 classes $x \in \{0, 1\}$. E.g., the probability of throwing a coin. The Bernoulli distribution is defined with parameter $\lambda \in [0, 1]$:

$$p(x) = \text{Bern}_x[\lambda] = \begin{cases} p(x=1) = \lambda \\ p(x=0) = 1 - \lambda \end{cases} \quad (3.1)$$

In short form, the above equation can be combined into one:

$$p(x) = \lambda^x (1 - \lambda)^{1-x} \Rightarrow \begin{cases} p(0) = \lambda^0 (1 - \lambda)^1 = 1 - \lambda \\ p(1) = \lambda^1 (1 - \lambda)^0 = \lambda \end{cases} \quad (3.2)$$

3.2.2 Categorical Distribution

Categorical Distribution is the generalization of *Bernoulli Distribution*, in case there are K classes for the discrete variable $x \in \{1, 2, \dots, K\}$. Accordingly, there will be K parameters to describe this **PDF**: $\lambda = [\lambda_1, \lambda_2, \dots, \lambda_K]$, with $\lambda_k \geq 0$ and $\sum \lambda_k = 1$. Each λ_k represents the probability to take the output k : $p(x = k) = \lambda_k$. In short: $p(x) = \text{Cat}_x[\lambda]$.

Another common way to represent the output is the one-hot vector, $\mathbf{x} \in \{\mathbf{e}_1, \mathbf{e}_2, \dots, \mathbf{e}_K\}$ with \mathbf{e}_k is the k -unit vector, which has all 0-element, except the k -element equal to 1. E.g., given 3 classes: $\mathbf{e}_1 = [1, 0, 0]^T, \mathbf{e}_2 = [0, 1, 0]^T, \mathbf{e}_3 = [0, 0, 1]^T$. We will then have:

$$p(\mathbf{x} = \mathbf{e}_k) = \prod_{j=1}^K \lambda_j^{x_j} = \lambda_k \quad (3.3)$$

because for $\mathbf{x} = \mathbf{e}_k$, only $x_k = 1$, while $x_j = 0, \forall j \neq k$.

3.2.3 Univariate Normal Distribution

Univariate Normal Distribution is also known as the Gaussian distribution. For single dimension data (in 1D): $x \in (-\infty, \infty)$, the mean $\mu \in \mathbb{R}$, and the variance σ^2 with $\sigma \in \mathbb{R}$.

$$p(x) = \text{Norm}_x [\mu, \sigma^2] = \mathcal{N}(\mu, \sigma^2) = \frac{1}{\sqrt{2\pi\sigma^2}} \cdot \exp\left(-\frac{(x - \mu)^2}{2\sigma^2}\right) \quad (3.4)$$

NOTE:

- Marginals prob. of Gaussian are again Gaussian.
- When estimating the parameters ([params.](#)) of a Gaussian, beware the underestimation problem.

$$\begin{aligned} \mathbb{E} [\mu_{ML}] &= \mu \\ \mathbb{E} [\sigma_{ML}^2] &= \left(\frac{N-1}{N}\right) \sigma^2 \\ \Rightarrow \tilde{\sigma}^2 &= \left(\frac{N}{N-1}\right) \sigma_{ML}^2 = \frac{1}{N-1} \sum_{n=1}^N (x_n - \hat{\mu})^2 \end{aligned}$$

3.2.4 Multivariate Normal Distribution

Multivariate Normal Distribution is the extension of *Univariate Normal Distribution* to multi-dimensional data: $\mathbf{x}, \boldsymbol{\mu} \in \mathbb{R}^D, \sigma^2 \Rightarrow \Sigma \in \mathbb{S}_{++}^D$ (\mathbb{S}_{++}^D is the set of positive definite

3 Probabilities




Figure 3.1: Bivariate Gaussian distribution ([src](#)).

symmetric matrix)

$$p(x) = \text{Norm}_x[\boldsymbol{\mu}, \Sigma] = \mathcal{N}(\boldsymbol{\mu}, \Sigma) = \frac{1}{2\pi^{D/2}|\Sigma|^{\frac{1}{2}}}\exp\left(-\frac{1}{2}(\mathbf{x} - \boldsymbol{\mu})^T\Sigma^{-1}(\mathbf{x} - \boldsymbol{\mu})\right) \quad (3.5)$$

3.2.5 Beta Distribution

This distribution describes the parameter for another distributions. E.g., Dirichlet [PDF](#) describes Categorical Distribution (Subsec. [3.2.2](#))

3.3 Parameter Estimation

Many of [ML](#) problems are boiled down to finding *statistical models*. Those models could predict the [prob.](#) for the classification problem, [prob.](#) of events that will happen, etc. It all end up with finding the suitable set of [params.](#) for these *statistical models*.

3.3.1 Maximum Likelihood Estimation

Maximum Likelihood Estimation ([MLE](#)) finds the parameters that maximize the [prob.](#) of the existing data.

$$\theta = \operatorname{argmax}_{\theta} p(x_1, x_2, \dots, x_N | \theta) \quad (3.6)$$

Assuming independent variables:

$$\theta = \operatorname{argmax}_{\theta} \prod_{n=1}^N p(x_n | \theta) \quad (3.7)$$

Maximum log-likelihood:

$$\theta = \operatorname{argmax}_{\theta} \sum_{n=1}^N [\log p(x_n | \theta)] \quad (3.8)$$

Minimum negative log-likelihood:

$$\theta = \operatorname{argmin}_{\theta} \sum_{n=1}^N [-\log p(x_n | \theta)] \quad (3.9)$$

3.3.2 Maximum A Posteriori

Sometimes, we have prior knowledge of the [PDF](#). E.g., we know that the [prob.](#) of getting head when flipping a coin is around 50%. Maximum A Posteriori ([MAP](#)) takes advantage of the prior knowledge $p(\theta)$ on the parameters θ by applying Bayes rule (Subsec. [3.1.3](#))

$$\theta = \operatorname{argmax}_{\theta} \prod_{n=1}^N p(x_n | \theta) p(\theta) \quad (3.10)$$

MLE suffers when there is not enough data \Rightarrow use MAP

3.4 Naive Bayes Classifier

Naive implies having the independence assumption on the variables.

$$\begin{aligned} c &= \operatorname{argmax}_{c \in \mathcal{C}} p(c|x) \\ &= \operatorname{argmax}_{c \in \mathcal{C}} p(x|c) p(c) \end{aligned}$$

If x is:

- continuous variable \Rightarrow Gaussian Naive Bayes
- feature vector \Rightarrow Multinomial Naive Bayes
- binary vector \Rightarrow Bernoulli Naive Bayes

3 Probabilities

Minimize the expected loss: $\mathbb{E}[L] = \sum_k \sum_j \int_{R_j} L_{kj} p(x, C_k) dx$ by choosing region R_j such that $\mathbb{E}[L] = \sum_k L_{kj} p(C_k | x)$

3.5 Views on the Decision Problem

3.5.1 Generative Methods

First determine the class-conditional densities and separately infer the prior class prob.
⇒ Bayes theorem ⇒ class membership

$$p(x|C_k) p(C_k) \Rightarrow y_k(x)$$

E.g., Mixture of Gaussians

3.5.2 Discriminative Methods

First solve the inference problem of determined the posterior class prob.

3.6 Unknown Notes

E.g., 2 class C_1, C_2 , 2 decisions α_1, α_2 .

The loss: $L(\alpha_j | C_k) = L_{kj}$.

The expected loss is equal to the $Risk(R)$.

$$\begin{aligned}\mathbb{E}_{\alpha_1}[L] &= R(\alpha_1 | x) = L_{11} p(C_1 | x) + L_{21} p(C_2 | x) \\ \mathbb{E}_{\alpha_2}[L] &= R(\alpha_2 | x) = L_{12} p(C_1 | x) + L_{22} p(C_2 | x)\end{aligned}$$

Choose α_1 if $R(\alpha_1 | x) < R(\alpha_2 | x)$

3.7 Probability Density Estimation

3.7.1 Histogram

This is **non-parametric** prob. density estimation. All other approaches are **parametric**.
The prob. of a bin:

$$p_i = \frac{n_i}{N \cdot \Delta_i} \quad (3.11)$$

in which n_i is the number of data points in that bin, N is the total number of data point, Δ_i is the width of the bin, often $\Delta_i = \Delta$.

Notes:

- Δ serves as the **smoothing factor**
- With D as the dimensions of the data points. The number of bins grow exponentially with $\mathcal{O}(k^D)$




Figure 3.2: Example of a histogram, $\Delta = 1$ ([src](#)).

3.7.2 Parametric Probability Density Estimation

In the other hands, one could find the prob. from $p = \int_{\mathcal{R}} p(y)dy \approx p(x)V$ where the region \mathcal{R} is sufficiently small.

$\Rightarrow p(x) \approx \frac{K}{N \cdot V}$, where K is the number of data points in the region, V is the volume of the region.

3 Probabilities

3.7.3 Kernel Methods

The kernel methods fix V and determine K . The volume V is the space restricted within a parzen window $k(u)$ that satisfies $k(u) \geq 0$.

A hyper-space cube:
$$k(u) = \begin{cases} 1 & \text{if } |u_i| \leq \frac{1}{2}h, \quad i = 1, 2, \dots, D \\ 0 & \text{else} \end{cases} \quad (3.12)$$

The number of points inside:
$$K = \sum_{n=1}^N k(x - x_n) \quad (3.13)$$

The region *volume*:
$$V = \int k(u)du = h^D \quad (3.14)$$

The probability:
$$\Rightarrow p(x) \approx \frac{K}{N.V} = \frac{1}{N.h^D} \sum_{n=1}^N k(x - x_n) \quad (3.15)$$

The **symmetric Gaussian kernel is a better substitution** for the asymmetric parzen window.

A Gaussian kernel:
$$k(u) = \frac{1}{\sqrt{2\pi h^2}} \exp\left(\frac{-u^2}{2h^2}\right) \quad (3.16)$$

The region *volume*:
$$V = \int k(u)du = 1 \quad (3.17)$$

The probability:
$$\Rightarrow p(x) \approx \frac{1}{N} \sum_{n=1}^N \frac{1}{(2\pi)^{D/2} h} \exp\left(\frac{-||x - x_n||^2}{2h^2}\right) \quad (3.18)$$

For Kernel methods, h is the **smoothing factor**.

Generalization: $k(u) \geq 0, \int k(u)du = 1$.

Size of the hypersphere is proportional to h^2 .

3.7.4 K-Nearest Neighbor

When you fix K and determine V , it leads to K-Nearest Neighbor.

[TODO: Add image]

$$p(x) \approx \frac{K}{NV} \quad (3.19)$$

Here, K is the **smoothing factor**.

- Too much bias \Rightarrow too smooth
- Too much variance \Rightarrow NOT smooth enough

\Rightarrow combine parametric methods to a mixture model

Mixture distribution = multi parametric model

3.7.5 Mixture of Gaussians

Mixture of Gaussians (MoG), as **Generative Model**, is defined from the prob. sum of elemental Gaussians: $p(x|\theta) = \sum_{j=1}^M p(x|\theta_j)p(j)$, where $p(x|\theta_j)$ is a **mixture component**, $p(j) = \pi_j$ is the **weight of the component**

$$p(x|\theta_j) = \frac{1}{\sqrt{2\pi}\sigma_j} \exp\left[\frac{-(x - \mu_j)^2}{2\sigma_j^2}\right], \quad p(j) = \pi_j, \quad \sum \pi_j = 1 \quad (3.20)$$

$$p(x|\theta_j) = \frac{1}{(2\pi)^{\frac{D}{2}} |\Sigma_j|^{\frac{1}{2}}} \exp\left[-\frac{1}{2}(x - \mu_j)^T \Sigma_j^{-1} (x - \mu_j)\right] \quad (3.21)$$

3.7.6 K-Means Clustering

There are 3 steps:

- Pick K centroids
- Assign sample to the centroid
- Adjust centroids

Step 2 and 3 are repeated until there is no change.

This leads to a local optimum, depends on initialization. It's sensitive to outliers, detects **spherical clusters only**.

[TODO: Add images]

Application: e.g., image compression.

3 Probabilities

3.7.7 EM Clustering

It's short for Expectation-Maximization. Assuming N data points and K Gaussians.

- **E-Step:** Fix the Gaussians, find $\gamma_j(x)$, which represent the **responsibility of component j for x** .

$$\gamma_j(x_n) = \frac{\pi_j \mathcal{N}(x_n | \mu_j, \Sigma_j)}{\sum_{k=1}^K \pi_k \mathcal{N}(x_n | \mu_k, \Sigma_k)} \quad \forall j = 1, 2, \dots, K, \quad n = 1, 2, \dots, N \quad (3.22)$$

- **M-Step:** Fix $\gamma_j(x)$, update the Gaussians.

$$\hat{N}_j = \sum_{n=1}^N \gamma_j(x_n) \quad (3.23)$$

$$\hat{\mu}_j = \frac{1}{\hat{N}_j} \sum_{n=1}^N \gamma_j(x_n) x_n \quad (3.24)$$

$$\hat{\pi}_j = \frac{\hat{N}_j}{N} \quad (3.25)$$

$$\hat{\Sigma}_j = \frac{1}{\hat{N}_j} \sum_{n=1}^N \gamma_j(x_n) (x_n - \hat{\mu}_j) (x_n - \hat{\mu}_j)^T \quad (3.26)$$

Notes:

- Regularization with $\Sigma + \sigma_{min} I$
- Initialization μ_j with K-Means
- Hard-assignment: each data point to 1 class \Rightarrow K-Means
- Soft-assignment: each data point \Rightarrow prob. to fall into many classes \Rightarrow EM Clustering

EM **needs more iteration**, because there are **more params..**

4 Linear Regression

Problem statement: Given data points $\mathbf{x}_i \in \mathbb{R}^D$ and their labels $y_i \in \mathbb{R}$, find the "line" that fit these data points. The line is represented via parameters \mathbf{w} . For each data point \mathbf{x} and its label y

$$\begin{aligned}\mathbf{w} &= [w_0, w_1, \dots, w_n]^T \\ \bar{\mathbf{x}} &= [1, x_0, \dots, x_n] \quad (\text{x bar}) \\ y &\approx \hat{y} = \bar{\mathbf{x}} \cdot \mathbf{w} \quad (\text{y hat}) \\ \Rightarrow \frac{1}{2} e^2 &= \frac{1}{2} (y - \bar{\mathbf{x}} \cdot \mathbf{w})^2\end{aligned}$$

The loss function for all points

$$\mathcal{L}(\mathbf{w}) = \frac{1}{2} \sum_{i=1}^N (y_i - \bar{\mathbf{x}}_i \cdot \mathbf{w})^2 \quad (4.1)$$

We need to find the weights that minimize the loss function

$$\mathbf{w}^* = \underset{\mathbf{w}}{\operatorname{argmin}} \mathcal{L}(\mathbf{w}) \quad (4.2)$$

We now can write the loss function using matrix form:

$$\mathcal{L}(\mathbf{w}) = \frac{1}{2} \|\mathbf{y} - \bar{\mathbf{X}} \cdot \mathbf{w}\|_2^2 \quad (4.3)$$

$$\text{with } \bar{\mathbf{X}} = \begin{bmatrix} \bar{\mathbf{x}}_1 \\ \bar{\mathbf{x}}_2 \\ \vdots \\ \bar{\mathbf{x}}_n \end{bmatrix} \text{ and } \mathbf{y} = \begin{bmatrix} y_1 \\ y_2 \\ \vdots \\ y_n \end{bmatrix}$$

Solution:

$$\frac{\partial \mathcal{L}(\mathbf{w})}{\partial \mathbf{w}} = \bar{\mathbf{X}}^T (\bar{\mathbf{X}} \mathbf{w} - \mathbf{y}) = 0 \quad (4.4)$$

$$\iff \bar{\mathbf{X}}^T \bar{\mathbf{X}} \mathbf{w} = \bar{\mathbf{X}}^T \mathbf{y} \quad (4.5)$$

$$\iff \mathbf{w} = (\bar{\mathbf{X}}^T \bar{\mathbf{X}})^\dagger \bar{\mathbf{X}}^T \mathbf{y} \quad (4.6)$$

4 Linear Regression

in which, A^\dagger (A dagger) is the pseudo inverse of a matrix, because it's might not inverseable. $\mathbf{A}^\dagger = (\mathbf{A}^T \mathbf{A})^{-1} \mathbf{A}^T$

NOTE: Sensitive to outliers \Rightarrow pre-processing

Multi-variable:

$$\begin{aligned}\mathcal{L}(\mathbf{w}) &= \frac{1}{2N} \|\mathbf{y} - \bar{\mathbf{X}} \cdot \mathbf{w}\|_2^2 \\ \Rightarrow \nabla_{\mathbf{w}} \mathcal{L}(\mathbf{w}) &= \frac{1}{N} \bar{\mathbf{X}}^T (\bar{\mathbf{X}} \mathbf{w} - \mathbf{y})\end{aligned}$$

Linear Discriminant Functions:

$$\begin{aligned}y(\mathbf{x}) &= \tilde{\mathbf{W}}^T \tilde{\mathbf{x}} \\ E(\mathbf{w}) &= \frac{1}{2} \sum_{n=1}^N \sum_{k=1}^K (\mathbf{w}_k^T \mathbf{x}_n - t_{kn})^2 \\ \tilde{\mathbf{W}} &= (\tilde{\mathbf{X}}^T \tilde{\mathbf{X}})^{-1} \tilde{\mathbf{X}}^T T \\ &= \tilde{\mathbf{X}}^\dagger T\end{aligned}$$

Generalized Discriminant:

$$y_k(x) = \sum_{j=1}^M w_{kj} \phi_j(x) + w_{k0} = \sum_{j=0}^M w_{kj} \phi_j(x), \quad \phi_0(x) = 1 \quad (4.7)$$

5 Logistics Regression

Though named as "regression", logistics regression is used for classification problem. For the classification problem with 2 classes, instead of a hard threshold (from linear classification), we could have a soft one, and find the **prob.** of the input belongs to either class.

$$f(x) = \theta(\omega^T x) \quad (5.1)$$

The sigmoid function:

$$f(s) = \frac{1}{1 + e^{-s}} \triangleq \sigma(s) \quad (5.2)$$

$$s = \ln\left(\frac{\sigma}{1 - \sigma}\right) \quad (5.3)$$

$$\sigma'(s) = \sigma(s)(1 - \sigma(s)) \quad (5.4)$$

The cross entropy error function:

$$J(w, x, y) = -(y_i \log z_i + (1 - y_i) \log(1 - z_i)) \quad (5.5)$$

with $z_i = f(w^T x_i)$

$$\Rightarrow \frac{\partial J}{\partial w} = (z_i - y_i)x_i \quad (5.6)$$

$$\Rightarrow w = w + \eta(y_i - z_i)x_i \quad (5.7)$$

NOTE: Require less **params.**, only D dimensions, compared to Gaussians with

$$\frac{M(M+5)}{2} + 1 \text{ params.}$$

6 Softmax Regression

Multinomial Logistics Regression, Maximum Entropy Classifier

$$a_i = \frac{\exp(z_i)}{\sum_{j=1}^C \exp(z_j)} \quad (6.1)$$

$$\text{so that } \begin{cases} a_i > 0 \\ \sum a_i = 1 \\ z_m > z_n \iff a_m > a_n \text{ (order)} \end{cases}$$

When z_i is too big

$$\frac{\exp(z_i)}{\sum_{j=1}^C \exp(z_j)} = \frac{\exp(z_i - c)}{\sum_{j=1}^C \exp(z_j - c)} \quad (6.2)$$

with $c = \max_i z_i$

$$J(w, x, y) = - \sum_{i=1}^N \sum_{j=1}^C y_{ij} \log(a_{ij}) \quad (6.3)$$

$$J(w, x, y) = - \sum_{n=1}^N \sum_{k=0}^1 [\mathbb{I}(t_n = k) \ln p(y_n = k | x_n; w)] \quad (6.4)$$

$$\Rightarrow E(w) = - \sum_{n=1}^N \sum_{k=1}^K \left[\mathbb{I}(t_n = k) \ln \frac{\exp(w_k^T x)}{\sum_{j=1}^K \exp(w_j^T x)} \right] \quad (6.5)$$

$$\frac{\partial J_i(w)}{\partial W} = x_i e_i^T = x_i (a_{ij} - y_{ij})^T \quad (6.6)$$

$$W = W + \eta x_i (y_i - a_i)^T \quad (6.7)$$

7 Error Functions

[TODO: Add graph and explanation]

7.1 Ideal Miss-classification Error

Gradient = 0 \Rightarrow can't use gradient descent.
It simply counts incorrectly classified points.

7.2 Squared Error - L_2 Loss

- Leads to closed form solutions
- Sensitive to outliers
- Penalize "too correct" data points

7.3 Cross Entropy Error

- Concave function \Rightarrow unique minimum exists
- Robust to outliers, error increases only roughly linear
- No closed-form solution, requires iterative method

7.4 Squared Error on Sigmoid / Tanh

- No penalty for "too correct" points
- Zero gradient for confidently incorrect classifications

\Rightarrow **Do NOT** use L_2 loss with sigmoid outputs, instead, use cross-entropy.

7.5 Hinge Error

- Robust to outliers
- Zero error for points outside margin \Rightarrow sparsity
- Not differentiable around $z_n = 1$

7 Error Functions

NOTE: Want the correct class to have a score that is higher than incorrect class by a fixed margin Δ .

$$L_i = \sum_{j \neq y_i} \max(0, s_j - s_{y_i} + \Delta) \quad (7.1)$$

in which, s_j is other classes score, s_{y_i} is real class score.

7.6 L_1, L_0 Loss

Median, no wrong points

$$L_1 = \sum |t - y| \quad (7.2)$$

$$L_2 = \sum (t - y)^2 \quad (7.3)$$

8 Neural Network

8.1 Gradient Descent

NOTE: Just use ADAM??

8.1.1 Vanilla Gradient Descent

$$\theta_{t+1} = \theta_t - \eta \nabla_{\theta} f(\theta_t) \quad (8.1)$$

Check derivative:

$$f'(x) \approx \frac{f(x + \varepsilon) - f(x - \varepsilon)}{2\varepsilon} \quad (\text{numerical gradient}) \quad (8.2)$$

8.1.2 Momentum

- Init: $v_{dW_0} = 0, v_{db_0} = 0$
- Calculate dW, db
- Update W, b

$$\Rightarrow \begin{cases} v_{dW} &= \beta v_{dW} + (1 - \beta)dW \\ v_{db} &= \beta v_{db} + (1 - \beta)db \end{cases} \Rightarrow \begin{cases} W &= W - \alpha v_{dW} \\ b &= b - \alpha v_{db} \end{cases} \quad (8.3)$$

The above formulas are to calculate the moving average of v_{dW} and v_{db} .

- Tips: Choose $\beta_1 = 0.9$, implying taking average of the last 10 steps.
- Reference source: [DeepLearning.AI](#).

8.1.3 Nesterov Accelerated Gradient

Nestorov Accelerated Gradient ([NAG](#)):

$$v_t = \gamma v_{t-1} + \eta \nabla_t J(\theta - \gamma v_{t-1}) \quad (8.4)$$

8.1.4 RMSprop

- Init $s_{dW_0} = 0, s_{db_0} = 0$
- Calculate dW, db
- Update W, b

$$\begin{cases} s_{dW} &= \beta s_{dW} + (1 - \beta) dW^2 \\ s_{db} &= \beta s_{db} + (1 - \beta) db^2 \end{cases} \Rightarrow \begin{cases} W &= W - \alpha \frac{dW}{\sqrt{s_{dW}} + \varepsilon} \\ b &= b - \alpha \frac{db}{\sqrt{s_{db}} + \varepsilon} \end{cases} \quad (8.5)$$

- Tips: choose $\beta_2 = 0.999, \varepsilon = 10^{-7}$
- Reference source: [DeepLearning.AI](#).

8.1.5 Adam

[Adam](#) is basically the combination of Momentum and RMSprop.

- Init $v_{dW_0}, s_{dW_0}, v_{db_0}, s_{db_0} = 0$
- Calculate dW, db
- Update W, b

$$\begin{cases} v_{dW} &= \beta_1 v_{dW} + (1 - \beta_1) dW \\ v_{db} &= \beta_1 v_{db} + (1 - \beta_1) db \\ s_{dW} &= \beta_2 s_{dW} + (1 - \beta_2) dW^2 \\ s_{db} &= \beta_2 s_{db} + (1 - \beta_2) db^2 \end{cases} \Rightarrow \begin{cases} v_{dW}^{cor.} &= \frac{v_{dW}}{1 - \beta_1^t} \\ v_{db}^{cor.} &= \frac{v_{db}}{1 - \beta_1^t} \\ s_{dW}^{cor.} &= \frac{s_{dW}}{1 - \beta_2^t} \\ s_{db}^{cor.} &= \frac{s_{db}}{1 - \beta_2^t} \end{cases} \Rightarrow \begin{cases} W &= W - \alpha \frac{v_{dW}^{cor.}}{\sqrt{s_{dW}^{cor.}} + \varepsilon} \\ b &= b - \alpha \frac{v_{db}^{cor.}}{\sqrt{s_{db}^{cor.}} + \varepsilon} \end{cases} \quad (8.6)$$

- Tips: choose $\beta_1 = 0.9, \beta_2 = 0.999, \varepsilon = 10^{-7}$
- Reference source: [DeepLearning.AI](#).

8.2 Tips and Tricks

- Shuffling
- Data Augmentation: reshape, rescale, crops, zooming, change color (color Principal Component Analysis ([PCA](#)))
- Normalizing the inputs

Convergence is the fastest if

- The mean of each input variable = 0
- Scale \Rightarrow same covariance

Mean cancellation \Rightarrow Kullback–Leibler ([K-L](#)) expansion \Rightarrow covariance equalization (if possible)

- Leaky Rectified Linear Unit ([ReLU](#)) is better a bit than [ReLU](#), ELU
- Weights initialization: Xavier-Glorot:

$$W \sim U \left(0, \sqrt{\frac{6}{n_{in} + n_{out}}} \right)$$

- Batch Norm(alization): Normalize after each layer
 \Rightarrow learn the moving average

- Drop out

NOTE: When in inferencing (after training), must multiply the activation output with the [prob.](#) that the weights are set to 0

Bibliography

- [GBC16] I. Goodfellow, Y. Bengio, and A. Courville. *Deep learning*. MIT press, 2016.
- [Vu18] H. T. Vu. *Deep learning*. 2018.