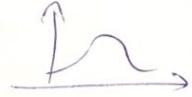
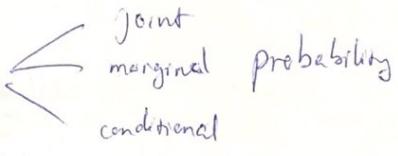
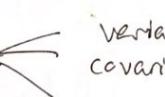


Probabilities

- + if x is discrete $\sum_x p(x) = 1 \rightarrow \forall p(x) \leq 1$
- + if x is continuous : theoretically $p(x)=0 \Rightarrow \int p(x) dx = 1$ 
- \Rightarrow There must be x such that ~~probability density function~~ $(\underline{\text{pdf}})$
- that the total sum $\int p(x) dx = 1$
- + Definition 
- | | |
|----------------|--|
| $p(x_i, y_i)$ | $\left\{ \begin{array}{l} \text{full: } p(X=x_i, Y=y_i) \\ p(X=x_i) \\ p(Y=y_i X=x_i) \end{array} \right.$ |
| $p(x_i)$ | |
| $p(y_i x_i)$ | |
- + Bayes rule: $p(y_i) \cdot p(x_i | y_i) = p(x_i, y_i) = p(x_i) \cdot p(y_i | x_i) \Rightarrow p(y_i | x_i) = \frac{p(x_i, y_i)}{p(x_i) p(y_i)}$
- posterior = $\frac{\text{likelihood} \times \text{prior}}{\text{normalization factor}} = \frac{p(x_i | y_i) p(y_i)}{\sum_y p(x_i | y_i) p(y_i)}$
- + Marginalization
- discrete: $p(x) = \sum_y p(x, y) = \sum_y p(x_i, y_i) / \sum_y p(x_i, y_i) p(y_i)$
 - continuous: $p(x) = \int p(x, y) dy$
- Nếu 2 biến độc lập \Rightarrow Independence $\Rightarrow \begin{cases} p(x, y) = p(x) \\ p(y | x) = p(y) \end{cases} \Leftrightarrow p(x, y) = p(x)p(y)$
- ④ Expectation - $E[x] = \sum_x x p(x) \quad (= \int x p(x) dx)$
- ④ $\forall f(\cdot) \Rightarrow E[f(x)] = \sum_x f(x) p(x) \quad (= \int f(x) p(x) dx)$
- ④ Measure variability \Rightarrow 
- | | |
|-------------------|--|
| variance | $\text{var}[f] = E[(f(x) - E[f(x)])^2] = E[f(x)^2] - E[f]^2$ |
| covariance matrix | $\text{cov}[x, y] = E_{x,y}[xy] - E[x]E[y]$ |
| | $= E_{x,y}[x^T y] - E[x]E[y^T]$ |
- ④ Sum Rule : Sum of joint prob. = marginal prob.
- Product Rule : Product of marginal prob & conditional prob = joint prob

Probability distribution:

$$x \in \{0, 1\} \quad p(x) = \text{Bern}_x[\lambda]$$

1) Bernoulli distribution:

$$\begin{aligned} x &= [0, 1] \Rightarrow \begin{cases} p(x=1) = \lambda \\ p(x=0) = 1-\lambda \end{cases} \quad \text{hay } p(x) = \lambda^x(1-\lambda)^{1-x} \\ \text{Tổng} \downarrow & \quad \downarrow \end{aligned}$$

Example = đồng xu

2) Categorical: $x \in \{e_1, e_2, \dots, e_k\} \approx \text{one-hot vector}$

$$p(x=k) = \lambda_k \quad p(x) = \text{Cat}_x[\lambda]$$

$$\Rightarrow p(x=e_k) = \prod_{j=1}^k \lambda_j^{x_j} = \lambda_k$$

$$\text{với } x = [x_1, x_2, x_3, \dots, x_k] \Rightarrow (= \lambda_1^{\circ} \lambda_2^{\circ} \dots \lambda_k^{\circ} \dots \lambda_k^{\circ})$$

3) Univariate Normal Distribution: Gaussian

$$\begin{aligned} p(x) &= f(\mu, \sigma^2) = \frac{1}{\sqrt{2\pi\sigma^2}} \exp\left(-\frac{(x-\mu)^2}{2\sigma^2}\right) \\ &= \text{Norm}_x[\mu, \sigma^2] \end{aligned}$$

4) Multivariate Normal Distribution: $\mu \in \mathbb{R}^p, \sigma^2 \Sigma \in \mathbb{S}_{++}^p$

$$\begin{aligned} p(x) &= \frac{1}{(2\pi)^{p/2} |\Sigma|^{1/2}} \exp\left(-\frac{1}{2} (\mathbf{x}-\mu)^T \Sigma^{-1} (\mathbf{x}-\mu)\right) \\ &= \text{Norm}_x[\mu, \Sigma] \end{aligned}$$

5) Beta distribution \Rightarrow mô tả thêm số che 1 distribution khác (VD: Bernoulli)

↳ Dirichlet Distribution (VD mô tả Categorical Distribution)

MAP MLE

Maximum

Maximum Likelihood Estimation

$$\theta = \underset{\theta}{\operatorname{argmax}} \ p(x_1, \dots, x_N | \theta)$$

$$\text{Independence assumption } \theta = \underset{\theta}{\operatorname{argmax}} \ \prod_{n=1}^N p(x_n | \theta)$$

$$\text{Maximum log-likelihood } \theta = \underset{\theta}{\operatorname{argmax}} \sum_{n=1}^N \log(p(x_n | \theta))$$

Minimum negative log-likelihood

MAP

Maximum a Posteriori
(Bayes rule)

$$\theta = \underset{\theta}{\operatorname{argmax}} \left[\prod_{i=1}^N p(x_i | \theta) p(\theta) \right]$$

Marginals of Gaussian are again Gaussian

Underestimation problem

$$\mathbb{E}[\mu_{ML}] = \mu$$

$$\mathbb{E}[\sigma^2_{ML}] = \left(\frac{N-1}{N}\right)\sigma^2$$

$$\Rightarrow \hat{\sigma}^2 = \frac{N}{N-1} \sigma^2_{ML}$$
$$= \frac{1}{N-1} \sum_{n=1}^N (x_n - \hat{\mu})^2$$

Thanh toán viên

Kiểm soát

Giám đốc

MLE suffers when there is not much data

⇒ Solve by MAP

Form NHBL 19 - 06/99, Li	
có 3 (Credit 3):	VAT amount:
Số tiền bằng chữ: Hai Muoi Hai Nghin Dong	Amount in words: NGUYEN HUU DUC; nop tien tk to chuc,
Remarks: Nội dung: NGUYEN HUU DUC; nop tien tk to chuc,	Net Income
PHÒNG KIỂM SÁT PHÒNG KHẢO DỤC HỌA	ĐÃ THU

Naive Bayes Classifier

Independence assumption \Rightarrow Naive: $p(c|x) = \text{const.} \cdot p(c)p(x|c)$

$$c = \arg \max_{c \in C} p(c|x)$$

$$= \arg \max_c p(c)p(x|c)$$

brain tumor \Rightarrow

Gaussian Naive Bayes

x feature vector \Rightarrow Multinomial
binary \Rightarrow Bernoulli

④ Minimize the expected loss: $E[L] = \sum_k \sum_j L_{kj} \int_{R_j} p(x, c_k) dx$
by choosing regimen R_j such that: $E[L] = \sum_k L_{kj} p(c_k | x)$

Views on the Decision Problem

+ Generative Methods: First determine the class-conditional densities & separately infer the prior class prob
 \Rightarrow Bayes Theorem \Rightarrow class membership
 $p(x|c_k) p(c_k) \Rightarrow y_k(x)$ Example: Mixture of Gaussians

+ Discriminative Methods: First solve the inference problem of determining the posterior class probabilities

ĐỀ NGHỊ GHI CÓ TÀI KHOẢN (Please Credit account):

SỐ TK (A/C No): 0011344844717	SỐ TIỀN
TÊN TK (A/C Name): Tín hanel	Bảng số (In which):
Địa chỉ (Address): Số 9, Trúc Bạch, Ba Đình, Hà Nội	Bảng chữ
NGÂN HÀNG (With Bank): Vietcombank	Địa chỉ

Form NHBL 01 - 06/99, Liên 1: Lưu, Liên 2: KH

Example:

2 Class $C_1 \quad C_2$

2 Decision $\alpha_1 \quad \alpha_2$

$$\text{Loss: } L(\alpha_j | C_k) = L_{kj}$$

$$\text{Expected loss} = \text{Risk}(R)$$

$$E_{\alpha_1}[L] = R(\alpha_1)$$

$$= L_{11} p(C_1|x) + L_{21} p(C_2|x)$$

$$E_{\alpha_2}[L] = R(\alpha_2|x)$$

$$= L_{12} p(C_1|x) + L_{22} p(C_2|x)$$

Choose α_1 if

$$R(\alpha_1|x) < R(\alpha_2|x)$$

002 HAN0704 21Aug1815:46:13 1100 CSD CA CSH DEP
7E49.0072 0-001-1-37-184471-7/04261
SỐ TIỀN MẤT HỌP: USD ****250,00 23.210,00/23.200,00
TIỀN GHI LỜI: (USD) *****50,00 23.300,00

DÀNH CHO NGÂN HÀNG (For Bank's Use only) MÃ VAT:

TRONG ĐÓ
In which:
KÝ TÊN (Signature):

SỐ TIỀN

non parametric

Probability Density Estimation

- Histograms: $p_i = \frac{n_i}{N\Delta_i}$  often $\Delta_i = \Delta$; D -dimension
 Δ is the smoothing factor
no of bins: $\in O(k^D)$

- $P = \int_R p(y) dy \approx p(x) V$ when R is sufficiently small
 $\Rightarrow p(x) \approx \frac{K}{NV} \quad (P = \frac{K}{N})$

determine K , fixed $V \leftarrow$

Kernel Methods

Parzen Window

$$k(u) = \begin{cases} 1 & |u_i| \leq \frac{1}{2}h, i=1, \dots, D \\ 0 & \text{else} \end{cases}$$

$$K = \sum_{n=1}^N k(x-x_n); \quad V = \int k(u) du = h^D$$

$$\Rightarrow p(x) \approx \frac{K}{NV} = \frac{1}{Nh^D} \sum_{n=1}^N k(x-x_n)$$

Gaussian kernel

$$k(u) = \frac{1}{\sqrt{2\pi h^2}} \exp\left\{-\frac{u^2}{2h^2}\right\}$$

$$K = \sum_{n=1}^N k(x-x_n); \quad V = \int k(u) du = 1$$

$$\Rightarrow p(x) \approx \frac{1}{N} \sum_{n=1}^N \frac{1}{(2\pi)^{D/2} h} \exp\left\{-\frac{\|x-x_n\|^2}{2h^2}\right\}$$

h is smoothing factor

General: $k(u) \geq 0 \quad \int k(u) du = 1$

Size of the hypersphere is proportional to h^2

$p(x|t_j)$: mixture component
 $p(t_j) = \pi_j$: weight of ↑

often $\Delta_i = \Delta$; D -dimension
no of bins: $\in O(k^D)$

→ fixed K , determine V

K-Nearest Neighbor

$K=3$

$$p(x) \approx \frac{K}{NV}$$


K is smoothing factor

Too much bias variance too smooth
too not smooth enough

⇒ including parametric methods
mixture models

Mixture distributions = multi parametric models

↓
Mixture of Gaussians (MoG)

$$p(x|\theta) = \sum_{j=1}^M p(x|\theta_j) p(j)$$

Generative Model
 $p(x|\theta_j) = \frac{1}{\sqrt{2\pi\sigma_j^2}} \exp\left\{-\frac{(x-\mu_j)^2}{2\sigma_j^2}\right\}; \quad p(j) = \pi_j$
 $\sum \pi_j = 1$

$$p(x|\theta_j) = \frac{1}{(2\pi)^{D/2} |\Sigma_j|^{1/2}} \exp\left\{-\frac{1}{2}(x-\mu_j)^T \Sigma_j^{-1} (x-\mu_j)\right\}$$

- K-means clustering: 3 steps
 - pick K centroids + assign sample
 - adjust centroids till no change

local optimum, depends on initialization

sensitive to outliers, detects spherical clusters only

⇒ Image compression
- EM Clustering (Expectation - Maximization)

$$E\text{-Step: } \gamma_j(x_n) = \frac{\pi_j N(x_n | \mu_j, \Sigma_j)}{\sum_{k=1}^K \pi_k N(x_n | \mu_k, \Sigma_k)} \Rightarrow \forall j=1\dots K, n=1\dots N$$

responsibility of component j for x

N data
K Gaussians

$$M\text{-Step: } \hat{N}_j = \sum_{n=1}^N \gamma_j(x_n) ; \quad \hat{\mu}_j = \frac{1}{\hat{N}_j} \sum_{n=1}^N \gamma_j(x_n) x_n$$

$$\hat{\pi}_j = \frac{\hat{N}_j}{N}$$

Regularization: $\Sigma + \sigma_{min} I$

Init μ_j with KMeans

$$\hat{\Sigma}_j = \frac{1}{\hat{N}_j} \sum_{n=1}^N \gamma_j(x_n)(x_n - \hat{\mu}_j)(x_n - \hat{\mu}_j)^T$$

Hard-assignments: 1 point to 1 class \Rightarrow Kmeans

Soft: Each point \Rightarrow prob to fall into many classes \Rightarrow EM

EM needs more iteration cause there are more parameters

Machine Learning

Linear Regression

$$w = [w_0, w_1, \dots, w_n]^T$$

$$\bar{x} = [1, x_0, \dots, x_n] \quad (\text{x bar})$$

$$y \approx \hat{y} = \bar{x} \cdot w \quad (\text{y hat})$$

$$\Rightarrow \frac{1}{2} e^2 = \frac{1}{2} (y - \bar{x} \cdot w)^2$$

$$\Rightarrow L(w) = \frac{1}{2} \sum_{i=1}^N (y_i - \bar{x}_i \cdot w)^2 \quad (\text{loss function})$$

$$\Rightarrow w^* = \arg \min_w L(w)$$

$$L(w) = \frac{1}{2} \|y - \bar{x} \cdot w\|_2^2$$

$$\bar{x} = \begin{bmatrix} \bar{x}_1 \\ \bar{x}_2 \\ \vdots \\ \bar{x}_n \end{bmatrix} \quad y = \begin{bmatrix} y_1 \\ y_2 \\ \vdots \\ y_n \end{bmatrix}$$

Solution: $\frac{\partial L(w)}{\partial w} = \bar{x}^T (\bar{x} \cdot w - y) = 0$

$$\Leftrightarrow \bar{x}^T \bar{x} \cdot w = \bar{x}^T y$$

$$\Leftrightarrow w = (\bar{x}^T \bar{x})^{-1} \bar{x}^T y \quad (\text{dagger})$$

vì chưa chắc
khả nghịch
nên dùng giao nghịch đảo
pseudo inverse

$$A^+ = (A^T A)^{-1} A^T$$

(?)

Sensitive to outliers \Rightarrow pre-processing

Multivariable: $\theta_{t+1} = \theta_t - \eta \nabla_{\theta} f(\theta_t)$ or $\theta_{t+1} = \theta_t - \eta \nabla_{\theta} f(\theta)$

$$L(w) = \frac{1}{2N} \|y - \bar{x} \cdot w\|_2^2 \Rightarrow \nabla_w L(w) = \frac{1}{N} \bar{x}^T (\bar{x} \cdot w - y)$$

Check derivative: $f'(x) \approx \frac{f(x+\epsilon) - f(x-\epsilon)}{2\epsilon}$ (numerical gradient)

- Linear Discriminant Functions:

$$y(x) = \tilde{w}^\top \tilde{x}$$

$$E(w) = \frac{1}{2} \sum_{n=1}^N \sum_{k=1}^K (w_k^\top x_n - b_{kn})^2$$

$$\begin{aligned}\tilde{w} &= (\tilde{x}^\top \tilde{x})^{-1} \tilde{x}^\top \cdot T \\ &= \tilde{x}^+ \cdot T\end{aligned}$$

- Generalized Discriminants

$$y_k(x) = \sum_{j=1}^M w_{kj} \phi_j(x) + w_{k0} = \sum_{j=0}^M w_{kj} \phi_j(x), \quad \phi_0(x) = 1$$

Logistic Regression

$$f(x) = \Theta(\omega^T x)$$

$$f(s) = \frac{1}{1+e^{-s}} \triangleq g(s) \quad \text{sigmoid func}$$

$s = \ln\left(\frac{g}{1-g}\right)$

$$g'(s) = g(s)(1-g(s))$$

$$J(w, x, y) = -\left(y_i \log z_i + (1-y_i) \log(1-z_i)\right)$$

with $z_i = f(w^T x_i)$ cross entropy error func

$$\Rightarrow \frac{\partial J}{\partial w} = (z_i - y_i) x_i$$

$$\Rightarrow w = w + \eta (y_i - z_i) x_i$$

② Require less parameters (only M: dimensions)
 compared to ... Gaussians .. $\frac{M(M+5)}{2} + 1$

$$J = -\sum_{n=1}^N \sum_{k=0}^1 \left\{ \mathbb{I}(t_n=k) \ln P(y_n=k | x_n; w) \right\} \Rightarrow E(w) = -\sum_{n=1}^N \sum_{k=1}^K \left\{ \mathbb{I}(t_n=k) \ln \frac{\exp(w_k^T x)}{\sum_{j=1}^K \exp(w_j^T x)} \right\}$$

Multinomial Logistic Regression, Maximum Entropy Classifier

Softmax Regression

$$a_i = \frac{\exp(z_i)}{\sum_{j=1}^c \exp(z_j)}$$

so that

$$\begin{cases} a_i > 0 \\ \sum a_i = 1 \\ z_m > z_n \Leftrightarrow a_m > a_n \text{ (order)} \end{cases}$$

→ when z_i too big

$$\frac{\exp(z_i)}{\sum_{j=1}^c \exp(z_j)} = \frac{\exp(z_i - c)}{\sum_{j=1}^c \exp(z_j - c)}$$

$$c = \max_i(z_i)$$

$$J(w, x, y) = -\sum_{j=1}^N \sum_{i=1}^c y_{ji} \log(a_{ji})$$

$$\frac{\partial J_i(w)}{\partial w} = x_i \varrho_i^T = x_i (a_{ji} - y_{ji})^T$$

$$w = w + \eta x_i (y_{ji} - a_{ji})^T$$

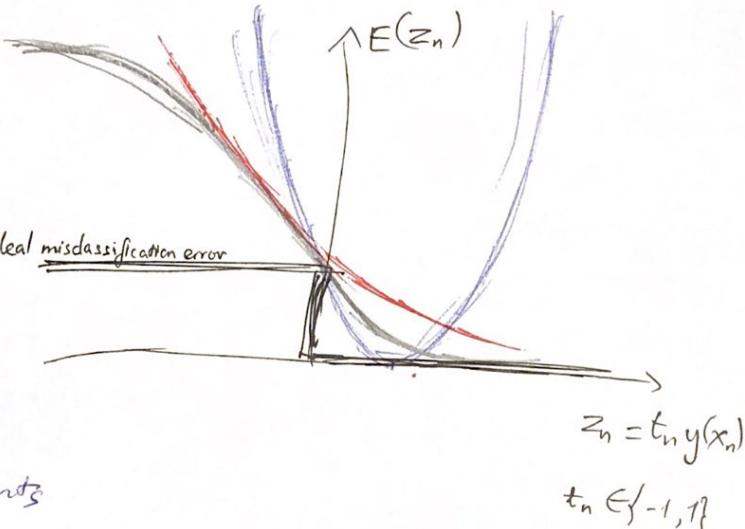
Error functions

→ Ideal misclassification error:

gradient = 0 \Rightarrow cannot use gradient descent

\Rightarrow Count wrong-classified points

→ Squared error - L₂ loss $\xrightarrow{\text{ideal misclassification error}}$



Leads to closed form solutions

sensitive to outliers

Penalize "too correct" data points

+ Cross entropy error

Concave function \Rightarrow unique minimum exists

Robust to outliers, error increases only roughly linear

No closed-form solution, requires iterative method

+ Squared error on sigmoid/tanh

No penalty for "too-correct" points

\Rightarrow Do NOT use L₂ loss with sigmoid outputs
(instead cross-entropy)

Zero gradient for confidently incorrect classifications

+ Hinge error $\sum [1 - t_n y(x_n)]$

robust to outliers

zero error for points outside margin \Rightarrow sparsity
not differentiable around $z_n = 1$

want the correct class to have a score that is higher than incorrect classes by a fixed margin Δ

+ L₁, L₀ loss

median No wrong points

$$L_1 = \sum |t - y|$$

$$L_2 = \sum (t - y)^2$$

$$L_i = \sum_{j \neq i} \max(0, s_j - s_i + \Delta)$$

other classes real class score

Support Vector Machine

ML

$$\text{margin} = \min_n \frac{y_n(w^T x_n + b)}{\|w\|_2}$$

$$\Rightarrow (w, b) = \arg \max_{w, b} \left\{ \min_n \frac{y_n(w^T x_n + b)}{\|w\|_2} \right\} = \arg \max_{w, b} \left\{ \frac{1}{\|w\|_2} \min_n y_n(w^T x_n + b) \right\}$$

$$\Rightarrow \text{Problem: } (w, b) = \arg \min_{w, b} \frac{1}{2} \|w\|_2^2$$

$$\text{subject to } 1 - y_n(w^T x_n + b) \leq 0, \forall n = 1, 2, \dots, N$$

Primal formulation of SUM:

$$L(w, b, \lambda) = \frac{1}{2} \|w\|_2^2 + \sum_{n=1}^N \lambda_n (1 - y_n(w^T x_n + b)) \quad \lambda_n \geq 0 \text{ thr}$$

$$\frac{\partial L(w, b, \lambda)}{\partial w} = 0 \Rightarrow w = \sum_{n=1}^N \lambda_n y_n x_n$$

$$\frac{\partial L(w, b, \lambda)}{\partial b} = 0 \Rightarrow \sum_{n=1}^N \lambda_n y_n = 0$$

$$\Rightarrow g(\lambda) = \sum_{n=1}^N \lambda_n - \frac{1}{2} \sum_{n=1}^N \sum_{m=1}^N \lambda_n \lambda_m y_n y_m x_n^T x_m \quad (\text{Dual formulation of SUM})$$

$$\text{Set } V = [y_1 x_1, y_2 x_2, \dots, y_N x_N]$$

$$\Rightarrow g(\lambda) = -\frac{1}{2} \lambda^T V^T V \lambda + 1^T \lambda \quad \text{is concave with } \lambda \geq 0$$

$$\Rightarrow \text{Find } \lambda \text{ by solving Quadratic Programming} \quad \sum_{n=1}^N \lambda_n y_n = 0$$

- Then use KKT to find w, b

$$S = \{n; \lambda_n \neq 0\} \Rightarrow \begin{cases} b = \frac{1}{N_S} \sum_{n \in S} \left(y_n - \sum_{m \in S} \lambda_m y_m x_m^T x_n \right) \\ w = \sum_{m \in S} \lambda_m y_m x_m \end{cases}$$

KKT conditions: $\begin{cases} \lambda_i \geq 0 \\ \lambda_i f(x_i) = 0 \\ f(x_i) \leq 0 \end{cases}$ if we have $L(\dots) = \dots + \lambda_i f(x_i)$

Hard Margin SVM: $(w, b) = \arg \min_{w, b} \frac{1}{2} \|w\|_2^2$

subject to: $y_n(w^\top x_n + b) \geq 1 \quad \forall n$

slack variable

Soft Margin SVM: $\frac{1}{2} \|w\|_2^2 + C \sum_{n=1}^N \xi_n$

"Soft" constraints: $y_n(w^\top x_n + b) \geq 1 - \xi_n \Leftrightarrow 1 - \xi_n - y_n(w^\top x_n + b) \leq 0 \quad \forall n$
(and $\xi_n \geq 0 \quad \forall n$)

C and margins are proportional $\begin{cases} C \uparrow \Rightarrow \text{margin} \downarrow \\ C \downarrow \Rightarrow \text{margin} \uparrow \end{cases}$

$$\Rightarrow (w, b, \xi) = \arg \min_{w, b, \xi} \frac{1}{2} \|w\|_2^2 + C \sum_{n=1}^N \xi_n$$

subject to $1 - \xi_n - y_n(w^\top x_n + b) \leq 0 \quad \forall n$
 $-\xi_n \leq 0 \quad \forall n$

$$\Rightarrow \mathcal{L}(w, b, \xi, \lambda, \mu) = \frac{1}{2} \|w\|_2^2 + C \sum_{n=1}^N \xi_n + \sum_{n=1}^N \lambda_n (1 - \xi_n - y_n(w^\top x_n + b)) - \sum_{n=1}^N \mu_n \xi_n$$

where $\lambda \geq 0$ and $\mu \geq 0$

$$g(\lambda, \mu) = \inf_{w, b, \xi} \mathcal{L}(w, b, \xi, \lambda, \mu)$$

$$\frac{\partial \mathcal{L}}{\partial w} = 0 \Leftrightarrow w = \sum_{n=1}^N \lambda_n y_n x_n; \frac{\partial \mathcal{L}}{\partial b} = 0 \Leftrightarrow \sum_{n=1}^N \lambda_n y_n = 0; \frac{\partial \mathcal{L}}{\partial \xi_n} = 0 \Leftrightarrow \lambda_n = C - \mu_n$$

$$\Rightarrow g(\lambda, \mu) = \sum_{n=1}^N \lambda_n - \frac{1}{2} \sum_{n=1}^N \sum_{m=1}^N \lambda_n \lambda_m y_n y_m x_n^\top x_m$$

$$\Rightarrow \lambda = \arg \max_{\lambda} g(\lambda) \text{ subject to } \sum_{n=1}^N \lambda_n y_n = 0$$

$$0 \leq \lambda_n \leq C \quad \forall n$$

$$\mathcal{N} = \{n: 0 < \lambda_n < C\} \Rightarrow b = \frac{1}{N_N} \sum_{n \in \mathcal{N}} \left(y_n - \sum_{m \in S} \lambda_m y_m x_m^\top x_n \right)$$

$$\mathcal{S} = \{m: 0 < \lambda_m \leq C\} \Rightarrow w = \sum_{m \in S} \lambda_m y_m x_m$$

$$\left\{ \begin{array}{l} \lambda_n = 0 \Rightarrow \text{safe points} \\ \xi_n = 0 \end{array} \right. \parallel \left\{ \begin{array}{l} 0 < \lambda_n < C \\ \xi_n = 0 \end{array} \right. \Rightarrow \text{marginal points} \parallel \left\{ \begin{array}{l} \lambda_n = C \\ \xi_n \leq 1 \end{array} \right. \begin{array}{l} \text{still correct} \\ \xi_n > 1: \text{wrong} \end{array}$$

Kernel Support Vector Machine

1/ Mercer's condition

Mercer's condition: kernel func. must satisfy (theoretically)

$$\sum_{n=1}^N \sum_{m=1}^N k(x_m, x_n) c_n c_m \geq 0 \quad \forall c_i \in \mathbb{R}, i = 1, 2, \dots, N$$

Thực hiện, khiêm kernel ko rõ ràng, nhưng vẫn chấp nhận được

+ $k(x, z) = x^T z$ (Linear) + $k(x, z) = (r + \gamma x^T z)^d$ polynomial

+ $k(x, z) = \exp(-\gamma \|x - z\|_2^2), \gamma > 0$ + $k(x, z) = \tanh(\gamma x^T z + r)$ sigmoid
radial basic function

The Lagrange dual function

Primal prob

$$x^* = \arg \min_x f_0(x) \text{ subject to } f_i(x) \leq 0, i=1, \dots, m \\ h_j(x) = 0, j=1, \dots, p$$

$$\Rightarrow L(x, \lambda, \nu) = f_0(x) + \sum_{i=1}^m \lambda_i f_i(x) + \sum_{j=1}^p \nu_j h_j(x)$$

dual variables

x^* là mâu
 $\nabla_{x,\lambda} L(x, \lambda) = 0$

Dual function: $g(\lambda, \nu) = \inf_{x \in D} L(x, \lambda, \nu)$ infimum

With $\lambda_i \geq 0, \forall i, \nu_j$: $g(\lambda, \nu) \leq p^*$ (optimal value of f_0)

\Rightarrow The Lagrange Dual problem:

$$\lambda^*, \nu^* = \arg \max_{\lambda, \nu} g(\lambda, \nu) \text{ subject to } \lambda \geq 0$$

$\Rightarrow g(\lambda, \nu)$ is concave

\Rightarrow Lagrange Dual Prob is a convex prob

$\Rightarrow \lambda^*$ and ν^* are called dual optimal / optimal Lagrange multipliers

② Weak duality: $\max g(\lambda, \nu) = d^* \leq p^* = \min f_0(x)$

$p^* - d^*$: optimal duality gap

③ Strong duality & Slater's constraint qualification

if $p^* = d^* \Rightarrow$ strong duality

If

$$x = \arg \min_x f_0(x) \\ \text{subject to } f_i(x) \leq 0 \quad i=1, \dots, n \\ Ax = b$$

f_0, f_1, \dots, f_n are convex \Rightarrow usually (not always) have strong duality

④ Slater's condition:

+ A feasible point is strictly feasible if: $\begin{cases} f_i(x) < 0, i=1, \dots, n \\ Ax = b \end{cases}$

\Rightarrow Slater's condition / theorem: if a strictly feasible point exists and primal problem is convex then strong duality holds

when strong duality happens: $f_0(x^*) = g(x^*, \lambda^*)$

$$\Rightarrow \sum_{i=1}^m \lambda_i^* f_i(x^*) = 0 \Rightarrow \lambda_i^* f_i(x^*) = 0 \quad i=1, \dots, m$$

Complementary slackness

④ Karush - Kuhn - Tucker conditions

f_0 is not necessarily convex, strong duality holds, then x^* , λ^* , ν^* must hold:

$$f_i(x^*) \leq 0, \quad i=1, \dots, m$$

$$h_j(x^*) = 0, \quad j=1, \dots, p$$

$$\lambda_i^* \geq 0, \quad i=1, \dots, m$$

$$\lambda_i^* f_i(x^*) = 0, \quad i=1, \dots, m$$

$$\nabla f_0(x^*) + \sum_{i=1}^m \lambda_i^* \nabla f_i(x^*) + \sum_{j=1}^p \nu_j^* \nabla h_j(x^*) = 0 \iff \textcircled{4}$$

These are necessary conditions

If f_0 convex, then KKT conditions are necessary and sufficient conditions for solutions

Convexity

1) Convex sets

→ Def 1: line connect 2 points of the set lies within the set

→ Def 2: C is convex iff for $x_1, x_2 \in C$:

$$x_\theta = \theta x_1 + (1-\theta)x_2 \in C \quad \forall 0 \leq \theta \leq 1$$

→ Hyperplanes: $a_1 x_1 + a_2 x_2 + \dots + a_n x_n = a^T x = b$
 $b, a_i \ (i \in \mathbb{N}) \in \mathbb{R}$

Hyperplane is also a convex set

→ Halfspace: $a_1 x_1 + a_2 x_2 + \dots + a_n x_n = a^T x \leq b$

Halfspace is a convex set

→ Positive definite: $x^T A^{-1} x \geq 0 \quad \forall x \in \mathbb{R}^n$ $A \succ 0$
 (if $\exists A^{-1} \Rightarrow \forall \lambda \neq 0 \text{ then } \lambda \neq 0$)

Positive semi-definite: $x^T B x \geq 0 \quad \forall x \in \mathbb{R}^n$ $B \succeq 0$

→ Intersection of convex sets is a convex set

⇒ Intersection of halfspaces & hyperplanes is polyhedra (is convex)

→ Convex combination of x_1, x_2, \dots, x_k : x if
 $x = \theta_1 x_1 + \theta_2 x_2 + \dots + \theta_k x_k$, with $\theta_1 + \theta_2 + \dots + \theta_k = 1$

Convex hull of a set (x_1, \dots, x_k) is ^{a set of} \forall all possible convex combination of that set

Convex hull of a set is the smallest convex set that contains that set

\Rightarrow 2 convex sets C and D are disjoint then exist a, b :

$$\begin{cases} a^T x \leq b & \forall x \in C \\ a^T x \geq b & \forall x \in D \end{cases}$$

Set of all x that $a^T x - b = 0$ is a hyperplane that separate C and D \Rightarrow separating hyperplane

Convex function

Definition $f: \mathbb{R}^n \rightarrow \mathbb{R}$ is a convex func if $\text{dom } f$ is a convex set and

$$f(\theta x + (1-\theta)y) \leq \theta f(x) + (1-\theta)f(y)$$
$$\forall x, y \in \text{dom } f, 0 \leq \theta \leq 1$$

Tập xác định của $f(\cdot)$: $\text{dom } f$

\Rightarrow If $f(\theta x + (1-\theta)y) < \theta f(x) + (1-\theta)f(y) \Rightarrow$ strictly convex

Nếu có các trị số đó là duy nhất và cũng là global minimum

Affine function : $f(x) = a^T x + b$ is both $\begin{cases} \text{convex} \\ \text{concave} \end{cases}$

if variable is a matrix $X \Rightarrow f(X) = \text{trace}(A^T X) + b$

Quadratic form : $f(x) = x^T A x + b^T x + c$ $\begin{array}{l} \text{convex if } A \succeq 0 \\ \text{concave if } -A \succeq 0 \end{array}$

Hàm sốs mìn SĐK cùs Norm \Rightarrow convex function

α -sublevel set of $f: \mathbb{R}^n \rightarrow \mathbb{R}$: $C_\alpha = \{x \in \text{dom } f \mid f(x) \leq \alpha\}$

Check if f is convex:

1) First-order condition:

iff differentiable with convex domain
 $f(\mathbf{x}) \geq f(\mathbf{x}_0) + \nabla f(\mathbf{x}_0)^T (\mathbf{x} - \mathbf{x}_0) \quad \forall \mathbf{x}, \mathbf{x}_0 \in \text{dom } f$
(if and only if)

2) Second-order condition: $\text{dang pho' bien hanh}$:

Hessian: $\nabla^2 f(\mathbf{x}) \succeq 0$

Optimization

$$\mathbf{x}^* = \arg \min_{\mathbf{x}} f_0(\mathbf{x}) \text{ subject to } \begin{cases} f_i(\mathbf{x}) \leq 0, i=1, 2, \dots, m \\ h_j(\mathbf{x}) = 0, j=1, 2, \dots, p \end{cases}$$

$$D = \bigcap_{i=0}^m \text{dom } f_i \cap \bigcap_{j=1}^p \text{dom } h_j$$

\Rightarrow set of all \mathbf{x} satisfying all constraints $\xrightarrow{\text{(in)}}$ feasible set
infeasible point $\not\in$ feasible point

⊗ Convex optimization problem:

$$\mathbf{x}^* = \arg \min_{\mathbf{x}} f_0(\mathbf{x}) \text{ subject to } \begin{cases} f_i(\mathbf{x}) \leq 0, i=1, 2, \dots, m \\ a_j^T \mathbf{x} - b_j = 0, j=1, \dots, p \end{cases}$$

f_0 is convex func, f_i is convex func, h_j is affine func

$f_i(\mathbf{x}) \leq 0 \Rightarrow$ 0 sublevel set of f_i $\forall \mathbf{x}: h_j(\mathbf{x}) = 0 \Rightarrow$ hyperplane

\Rightarrow We optimize a convex function in a convex set domain

Linear Programming (Guy hoạch tuyến tính) f_0, f_i are affine funcs

cvxopt.solvers.lp $\mathbf{x} = \arg \min_{\mathbf{x}} c^T \mathbf{x} + d \text{ subject to } \begin{cases} G \mathbf{x} \leq h \\ A \mathbf{x} = b \end{cases}$

A Standard form LP: $\mathbf{x} = \arg \min_{\mathbf{x}} c^T \mathbf{x} \text{ subject to } \begin{cases} A \mathbf{x} = b \\ \mathbf{x} \geq 0 \end{cases}$

Quadratic Programming

`cvxopt.solvers.qp`

$$x = \arg \min_x \frac{1}{2} x^T P x + q^T x + r \text{ subject to } \begin{array}{l} Gx \leq h \\ Ax = b \end{array}$$

P là ma trận đối称 dương $P \geq 0$

→ Khi $P=0 \Rightarrow QP$ là LP

Geometric Programming

`cvxopt.solvers.gp`

→ Function $f: \mathbb{R}^n \rightarrow \mathbb{R}$ with $\text{dom } f = \mathbb{R}_{++}^n$ (all element > 0)

$$f(x) = c x_1^{a_1} x_2^{a_2} \dots x_n^{a_n} \quad c > 0 \quad a_i \in \mathbb{R}$$

f is a monomial function

$$\rightarrow f(x) = \sum_{k=1}^K c_k x_1^{a_{1k}} x_2^{a_{2k}} \dots x_n^{a_{nk}} \quad c_k > 0$$

f is a posynomial function

$$x = \arg \min_x f_0(x) \text{ subject to } \begin{array}{l} f_i(x) \leq 1, i = 1, 2, \dots, m \\ h_j(x) = 1, j = 1, 2, \dots, p \end{array}$$

where f_0, f_i are posynomials and h_j are monomials

⇒ geometric programming ($x > 0$ is hidden)

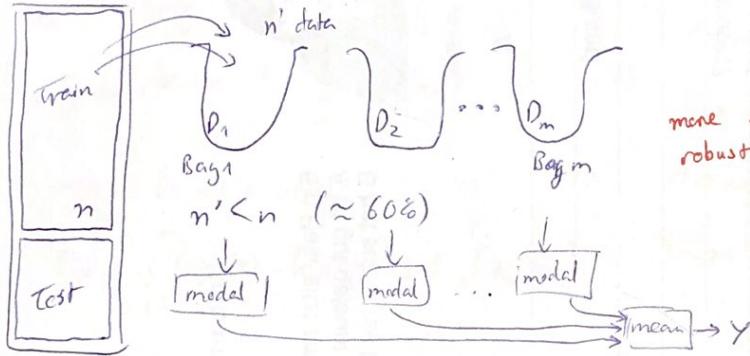
$$\text{Set } \begin{cases} x_i = e^{y_i} \Rightarrow f_0(x) = \exp(a^T y + b) \\ y_i = \log(x_i) \end{cases}$$

Ensemble of models

(also as additive model)

Bagging

Bootstrap aggregating
Data



random with replacement:

Bag i: can have multiple times data x_i

Simple, easy to implement, commonly used



Boosting

- 1, Ada boost (exponential error)
- 2, Gradient boosting
- 3, Xg boost (extreme gradient boosting)
- 4, Gentle Boost (cross entropy error)

Steps for Ada boost:

- 1, 1st weak classifier
- 2, Calculate error $J \Rightarrow \epsilon$
(in case of normalization w
 $\Rightarrow \sum w=1 \Rightarrow \epsilon = J$)
- 3, Calculate amount of say α from ϵ
- 4, Update w with α , (normalize w)
- 5, 2nd weak classifier with new weight w
or sample on new w

$$-g(x_i) = -\frac{\partial L(y_i, F(x_i))}{\partial F(x_i)}$$

\Rightarrow Regression model: h_i for $(x_i, -g(x_i))$

$$\Rightarrow F := F + g h, \quad g=1$$

Ada boost

$$H(x) = \text{sign} \left(\sum_{m=1}^M \alpha_m h_m(x) \right)$$

$$J_m = \sum_{n=1}^N w_n^{(m)} I(h_m(x) \neq t_n)$$

$$E_m = \frac{\sum_{n=1}^N w_n^{(m)} I(h_m(x) \neq t_n)}{\sum_{n=1}^N w_n^{(m)}}$$

$$\alpha_m = \ln \left(\frac{1-E_m}{E_m} \right)$$

$$w_n^{(m+1)} = w_n^{(m)} \cdot \exp \{ \alpha_m I(h_m(x) \neq t_n) \}$$

Gradient boosting (Regression)

(x_i, y_i) & first model $F_1(x)$
 \Rightarrow fit $F_2(x)$ to $(x_i, \underbrace{y_i - F_1(x_i)}_{\text{Residuals}})$

$$J = \sum L(y_i, F(x_i))$$

$$\Rightarrow F_{i+1} = F_i + h_i$$

$$L = \frac{(y_i - F_i)^2}{2} \Rightarrow -g = y_i - F_i(x_i)$$

$$L = |y - F| \quad \text{absolute loss}$$

$$L = \begin{cases} \frac{\gamma_2}{\delta} (y - F)^2 & |y - F| \leq \delta \\ \frac{\gamma_2}{\delta} (y - F - \delta)^2 & |y - F| > \delta \end{cases} \quad \text{Huber loss}$$

Bayesian Model Averaging:

$$p(x) = \sum_{h=1}^H p(x|h) p(h)$$

H different models $h = 1, \dots, H$
with prior probability $p(h)$

Error Reduction

Ex: bagging

$$y_{\text{com}}(x) = \frac{1}{M} \sum_{m=1}^M y_m(x)$$

$$\Rightarrow \text{Average error of individual models} \\ E_{\text{AV}} = \frac{1}{M} \sum_{m=1}^M \mathbb{E}_x [\varepsilon_m(x)^2]$$

$$y(x) = h(x) + \varepsilon(x)$$

$$\text{Average error of committee}$$

$$\mathbb{E}_x = \left[\{y_m(x) - h(x)\}^2 \right] = \mathbb{E}_x [\varepsilon_m(x)^2]$$

$$\begin{aligned} E_{\text{com}} &= \mathbb{E}_x \left[\left\{ \frac{1}{M} \sum_{m=1}^M y_m(x) - h(x) \right\}^2 \right] \\ &= \mathbb{E}_x \left[\left\{ \frac{1}{M} \sum_{m=1}^M \varepsilon_m(x) \right\}^2 \right] \end{aligned}$$

\Rightarrow if errors have 0 mean: $\mathbb{E}_x [\varepsilon_m(x)] = 0$

errors are uncorrelated: $\mathbb{E}_x [\varepsilon_m(x) \varepsilon_j(x)] = 0 \Rightarrow E_{\text{com}} = \frac{1}{M} E_{\text{AV}}$

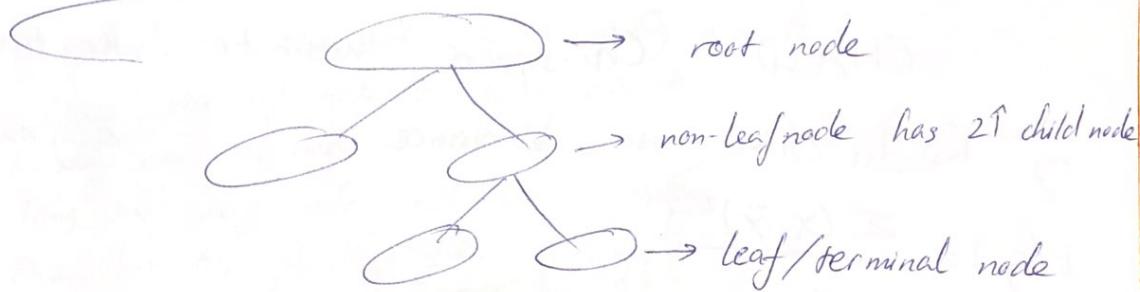
(?) \rightarrow

unrealistic

However
In general $E_{\text{com}} < E_{\text{AV}}$

- ④ The weak classifiers have to be a bit unstable algo like decision trees, neural net, linear discriminant func
not nearest neighbor, SVM, linear regression

Decision Trees



All non-leaf nodes have 2 child nodes \Rightarrow binary decision tree.

Iterative Dichotomiser 3 (ID3)
only for categorical attribute (discrete)

Classification & Regression Tree (CART)
for both categorical & continuous

6 question

Bin / multi valued
when node \rightarrow leaf
deal impure nodes?
how select query
pruned?
missing attributes?

Impurity Measures:

$$\rightarrow \text{Misclassification: } i(s_j) = 1 - \max_k p(G_k | s_j)$$

\rightarrow Information gain: C class với N_c là số phần tử mỗi class

$$H(S) = - \sum_{c=1}^C \frac{N_c}{N} \log\left(\frac{N_c}{N}\right) \quad \text{entropy tại node đầu}$$

Chọn thuộc tính x : \Rightarrow K child node s_1, \dots, s_K với m_k phần tử
entropy sum with weight $H(x, S) = \sum_{k=1}^K \frac{m_k}{N} H(s_k) \Rightarrow G(x, S) = H(S) - H(x, S)$

$$x^* = \arg \max_x G(x, S) = \arg \min_x H(x, S) \quad \text{Information gain}$$

Reduction in entropy = gain in information

$$\rightarrow \text{Gain Index: } i(s_j) = \sum_{k=1}^K p(G_k | s_j) p(C_k | s_j) \quad \text{variance impurity at node } s_j$$

$$= \frac{1}{2} \left(1 - \sum_k p^2(G_k | s_j) \right)$$

$$H(x, s_j) = \sum_{k=1}^K \frac{m_k}{N} i(s_j) \Rightarrow x^* = \min H(x, s_j)$$

> +) Chi-square = $\sqrt{\frac{(actual - expected)^2}{expected}}$

arg max

CHAID

Chi-square Automatic Interaction Detector

→ Redu chien in Variance

$$i(s_i) = \frac{\sum (x - \bar{x})^2}{n}$$

HỘ CHIẾU
HORIORITY

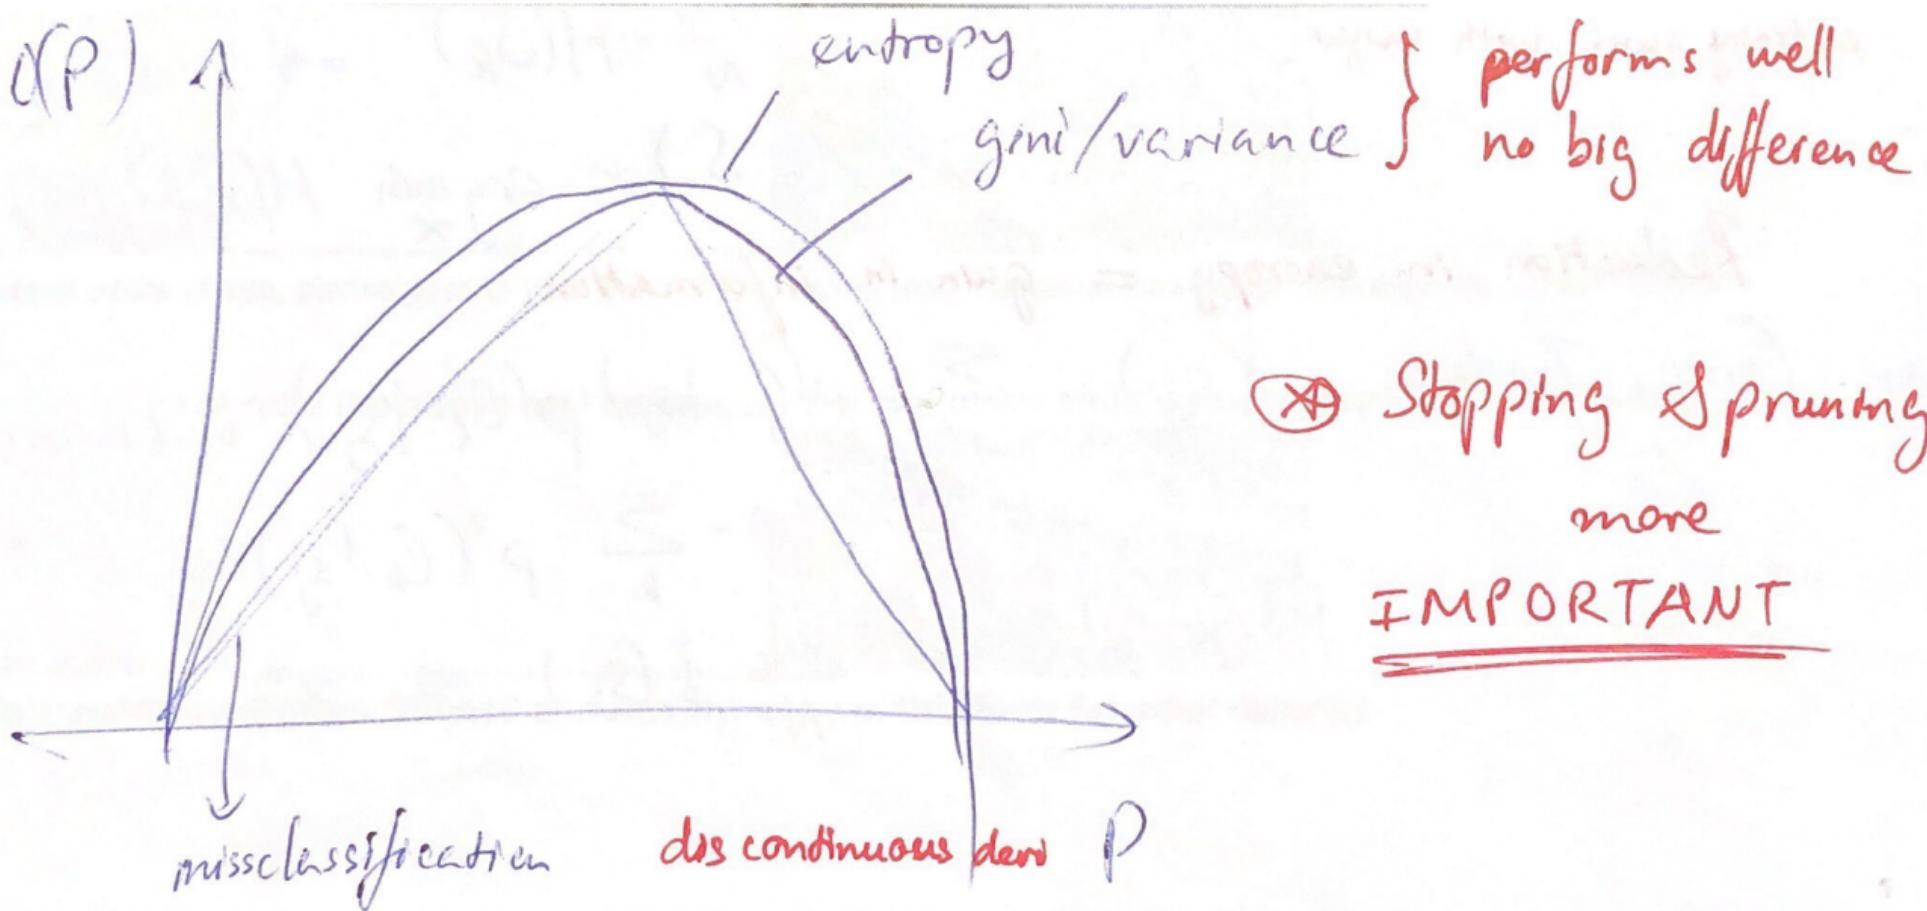
ẤT NHẬP CÀNH
partment
JÖNG
sector

hồng
Division

Thân

g hộ chiếu
bearer

Điểm



Stopping Conditions: (prevent over-fitting) = Prepruning

- $H(S) = 0$: entropy = 0 \Rightarrow All items $\in 1$ class
- Số phần tử trong node < 1 ngưỡng nào đó
Class cho leaf node = class đa số
- Khoảng cách từ node \rightarrow root = c \Rightarrow hạn chế chiều sâu
- Tổng số leaf node vượt quá 1 ngưỡng
- Phân chia tiếp theo giảm entropy quá nhiều

\Rightarrow Pruning: 3) Validation Set: Tùy node nào tăng Precision cho VS
(Reduced error pruning method)

Post-pruning 4) Regularized loss function:

$$L = \sum_{k=1}^K \frac{|S_k|}{|S|} H(S_k) + \lambda K$$

Minimum error

Sensitive to no. classes
least accurate in practice

Pessimistic

most crude & quickest
no need validation set
but + caution needed

Error complexity: $R(t) = r(t) p(t) + \alpha N t$

Critical value: the value we choose to decide query at each nodes
 \Rightarrow depends on how the tree is created

Reduced error: ...

More Stable & accurate

Computational Complexity:

N data points

D dimensionality

\Rightarrow

Storage $O(N)$

Test runtime $O(\log N)$

Training time $O(DN^2 \log N)$

max $O(DN \log N)$ times each node
max N nodes

the worst case

1 data each leaf
 $\Rightarrow k \text{ step} \Rightarrow 2^k = N$
 $\Rightarrow K \in O(\log N)$

Summary

Simple

Interpretable results

Resistance to overfitting

Memory consumption \Rightarrow suitable for
Noisy weak classifiers problems with
not generalized well
Sensitive outliers
expensive learning step

Random Forest

Handle missing values
while maintain accuracy
Won't overfit

Not good with regression
Have little control to modify

Step: Sample from training set

choose $m < M$ (input features). At each node, select m random data from M , to decide query attribute to split the node
Grow tree to largest, no pruning

Predict data: majority vote \Rightarrow classification
average \Rightarrow regression

④ choose randomly K attributes.
Training time: $O(KN^2 \log N)$ $K \ll D$ ($K = \sqrt{Ns}$)

Typical: $\begin{cases} K=10 & \text{root node} \\ K=100d & \text{level-d node} \end{cases}$

Ada boost

- ④ Ensembles of Classifiers: K classifiers, independent, error prob < 0.5
- ④ Suitable for: unstable algorithms: $\begin{cases} \text{decision trees} \\ \text{neural networks} \end{cases}$
- ④ good with stable methods: $\begin{cases} \text{nearest neighbor} \\ \text{SVMs} \\ \text{linear regression} \end{cases}$

$$\text{④ AdaBoost} \quad F(x) = \text{sign}\left(\sum_{m=1}^M \theta_m f_m(x)\right)$$

Initial weight for each data: $w_i(x_i, y_i) = \frac{1}{n}$

Classification error: $\epsilon_m = E_{w_m}[1_{y \neq f_m(x)}]$

\Rightarrow Update rule: $\theta_m = \frac{1}{2} \ln\left(\frac{1-\epsilon_m}{\epsilon_m}\right)$

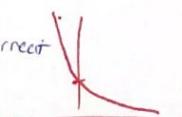
$$w_{m+1}(x_i, y_i) = \frac{w_m(x_i, y_i) \exp[-\theta_m y_i f_m(x_i)]}{Z_m}$$

Z_m is normalization factor

$$\text{Bayesian Model Averaging} \quad p(x) = \sum_{h=1}^H p(x|h) p(h)$$

Exponential error used in AdaBoost

Fast convergence
No penalty for too correct
less robust to outliers



Weighted error function $J_m = \sum_{n=1}^N w_n^{(m)} I(h_m(x) \neq t_n)$

$$H(x) = \text{sign} \left(\sum_{m=1}^M \alpha_m h_m(x) \right)$$

$$E_m = \sum_{n=1}^N w_n^{(m)} I(h_m(x) \neq t_n)$$

- Bargeldauszahlungen mit der Debitkarte [Deutsche Bank Card⁴] an Geldautomaten ausländischer Kooperationspartner⁶ in Fremdwährung
- Kontoauszüge im Online-Banking oder am Bankingterminal

Final

Classification

In begrenzter Anzahl

- Ausgabe einer Debitkarte [Deutsche Bank Card⁴]

Über diese Anzahl hinausgehende Dienste werden getrennt in Rechnung gestellt.

Zahlungen (ohne Karten)

Überweisung

[SEPA-Überweisung in EUR innerhalb des EWR¹]

amount of say

$$\sum_{n=1}^N w_n^{(m)}$$

Normalize

$$E = \sum_{n=1}^N w_n^{(m)} \exp\left(-\frac{1}{2} t_n \alpha_m h_m(x_n)\right)$$

$$\Rightarrow w_n^{(m+1)} = w_n^{(m)} \exp\left[\alpha_m I(h_m(x_n) \neq t_n)\right]$$

Je beleghafte Überweisung bzw. Überweisung über einen Mitarbeiter im telefonischen Kundenservice

$\alpha_m = \ln \left(\frac{1 - E_m}{E_m} \right)$	1,50 EUR
	4,00 EUR
	zzgl.
	1,50 EUR

Je formlos erteilte Überweisung zzgl. Entgelt für beleghafte Überweisung

Gradient Descent

ML

$$\theta_{t+1} = \theta_t - \eta \nabla_{\theta} f(\theta_t)$$

check derivative: $f'(x) \approx \frac{f(x+\varepsilon) - f(x-\varepsilon)}{2\varepsilon}$

+ Momentum:

\Rightarrow Nesterov accelerated gradient (NAG)

- Init: $v_{dw_0} = 0, v_{db_0} = 0$ $v_t = \gamma v_{t-1} + \eta \nabla_{\theta} f(\theta - \gamma v_{t-1})$
- Update: Calculate dW, db

$$\Rightarrow \begin{cases} v_{dw} = \beta v_{dw} + (1-\beta) dW \\ v_{db} = \beta v_{db} + (1-\beta) db \end{cases} \quad \begin{cases} w = w - \alpha v_{dw} \\ b = b - \alpha v_{db} \end{cases}$$
 $\beta_1 = 0,9$

+ RMSprop (Root mean squared prop)

- Init $s_{dw_0} = 0, s_{db_0} = 0$
- Calculate dW, db
- Update: $\begin{cases} s_{dw} = \beta s_{dw} + (1-\beta) dW^2 \\ s_{db} = \beta s_{db} + (1-\beta) db^2 \end{cases} \Rightarrow \begin{cases} w = w - \alpha \frac{dW}{\sqrt{s_{dw}} + \varepsilon} \\ b = b - \alpha \frac{db}{\sqrt{s_{db}} + \varepsilon} \end{cases}$
 $\beta_2 = 0,999$
 $\varepsilon = 10^{-7}$

+ Adam (Adaptive moment estimation)

- Init: $v_{dw}, s_{dw}, v_{db}, s_{db} = 0$
- Calculate dW, db
- $\begin{cases} v_{dw} = \beta_1 v_{dw} + (1-\beta_1) dW \\ v_{db} = \beta_1 v_{db} + (1-\beta_1) db \\ s_{dw} = \beta_2 s_{dw} + (1-\beta_2) dW^2 \\ s_{db} = \beta_2 s_{db} + (1-\beta_2) db^2 \end{cases} \Rightarrow \begin{cases} v_{dw}^{\text{corrected}} = \frac{v_{dw}}{1-\beta_1^t}, v_{db}^{\text{cor.}} = \frac{v_{db}}{1-\beta_1^t} \\ s_{dw}^{\text{cor.}} = \frac{s_{dw}}{1-\beta_2^t}, s_{db}^{\text{cor.}} = \frac{s_{db}}{1-\beta_2^t} \end{cases}$
 $\beta_1 = 0,9$
 $\beta_2 = 0,999$
 $\varepsilon = 10^{-7}$
- $w := w - \alpha \frac{v_{dw}^{\text{cor.}}}{\sqrt{s_{dw}^{\text{cor.}}} + \varepsilon}; b := b - \alpha \frac{v_{db}^{\text{cor.}}}{\sqrt{s_{db}^{\text{cor.}}} + \varepsilon}$

Neural Network

tips & tricks

- Shuffling: ✓
- Data augmentation: reshape, rescale, crops, zooming
change color (color PCA)

- = Normalizing the inputs: ✓

convergence is fastest if:
+ mean each input variable = 0
+ scale \Rightarrow same covariance

Mean Cancellation \Rightarrow K-L expansion \Rightarrow covariance equalization
(if possible)

- Leaky ReLU better than Relu, ELU ✓

- Initialize Weights:

Xavier Glorot: $W \sim U\left(0, \sqrt{\frac{6}{n_{in} + n_{out}}}\right)$ ✓

- Batch Normalization

Normalize after each layer (?)

\Rightarrow Learn this moving average

- Drop out \rightarrow Note: when applying, must multiply activation with the prob that it is set to 0

Neural Network

(General stuffs)

Forward pass

$$y^{(0)} = x$$

$$z^{(k)} = W^{(k)} y^{(k-1)} \quad k=1, \dots, l$$

$$y^{(k)} = g_k(z^{(k)})$$

$$y = y^{(l)}$$

$$E = L(t, y) + \lambda \Omega(W)$$

Backward Pass

$$h \leftarrow \frac{\partial E}{\partial y} = \frac{\partial}{\partial y} L(t, y) + \lambda \frac{\partial}{\partial y} \Omega$$

for $k = l \rightarrow 1$

$$h \leftarrow \frac{\partial E}{\partial z^{(k)}} = h \odot g'(y^{(k)})$$

$$\frac{\partial E}{\partial w^{(k)}} = h y^{(k-1)\top} + \lambda \frac{\partial \Omega}{\partial w^{(k)}}$$

$$h \leftarrow \frac{\partial E}{\partial y^{(k-1)}} = W^{(k)\top} h$$

Gradient Descent

$$w_{kj}^{(\sigma+1)} = w_{kj}^{(\sigma)} - \eta \frac{\partial E(w)}{\partial w_{kj}} \Big|_{w^{(\sigma)}}$$

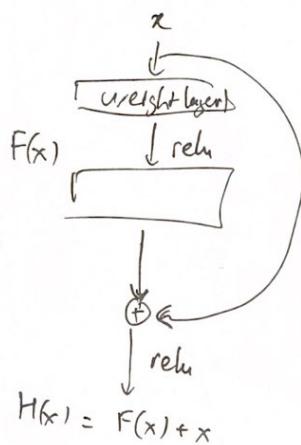
Momentum method

BPTT: Back propagation through time

$$\frac{\partial E_t}{\partial w_{ij}} = \sum_{1 \leq k \leq t} \left(\frac{\partial E_t}{\partial h_t} \frac{\partial h_t}{\partial h_k} \frac{\partial h_k}{\partial w_{ij}} \right) ; \quad E = \sum_{1 \leq t \leq T} E_t$$

$$\frac{\partial h_t}{\partial h_k} = \prod_{t \geq i \geq k} \frac{\partial h_i}{\partial h_{i-1}}$$

Residual



Word Embeddings

The idea of represent a word as a vector, but not a one-hot coding. In which similar words have similar vector values.

→ The trigram (n-gram method):

high amount of n-tuples of words \Rightarrow predict relative probabilities

Problems: scalability, partial observability

→ Word embedding:

$$x_{v \times 1} \rightarrow w_{v \times d} \rightarrow h_{d \times 1}$$

1 of K encoding

→ word 2vec :

CBOW:
syntactic
(grammar)

only care which words occur
don't care about order of occurrence

SKIP Gram:
semantic
(meaning)

less weight to the more distance word

→ Hierarchical softmax : organize word in binary search tree

CSiME

JY

3) Search Algorithm

④ Naive Bayes.

$$P(Y, f_1 \dots f_n)$$

$$\text{From data} \Rightarrow 1, P(y_j, f_1, \dots f_n) = P(y_j) \pi_i P(f_i | y_j)$$

$$2, P(f_1 \dots f_n) = \sum_j P(y_j) \pi_i P(f_i | y_j)$$

$$3, P(Y | f_1, \dots f_n) = \frac{P(Y, f_1 \dots f_n)}{P(f_1 \dots f_n)}$$

⑤ Hinge loss:

$$s_j = f(x_i, w),$$

$$L_i = \sum_{j \neq y_i} \max(0, s_j - s_{y_i} + 1)$$

⑥ Regularization :

$$L_1 : R(w) = \sum_k \sum_l |w_{k,l}|$$

$$L_2 : R(w) = \sum_k \sum_l w_{k,l}^2$$

$$\text{Elastic Net } R(w) = \sum_k \sum_l \beta w_{k,l}^2 + |w_{k,l}|$$

⑦ Gradient descent

$$\hat{x} = x - \xi \nabla_x f(x)$$

$$\text{⑧ Laplace smoothing: } P_{\text{Lap}, k}(x) = \frac{\frac{\text{freq of class } x}{d(x) + k}}{\frac{\text{total no. of data}}{N+k} \underbrace{|X|}_{\text{No. of Class}}}$$

Graphs

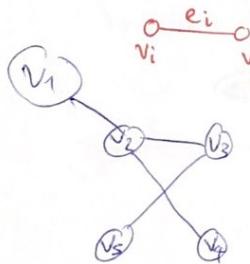
Definition: Undirected Graph vs Directed Graph

$$G = (V, E)$$

nodes & edges

$$V = \{v_i\}$$

$$E = \{e_i\} = \{\{v_i, v_j\}\}$$



Adjacency Matrix

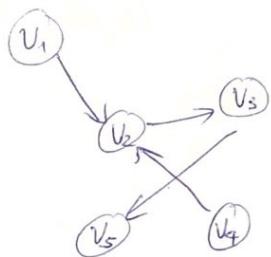
$$\begin{bmatrix} 0 & 1 & 0 & 0 & 0 \\ 1 & 0 & 1 & 1 & 0 \\ 0 & 1 & 0 & 0 & 1 \\ 0 & 1 & 0 & 0 & 0 \\ 0 & 0 & 1 & 0 & 0 \end{bmatrix}$$

symmetric

$$D = (V, A)$$

nodes & arcs

$$\begin{array}{c} \text{to} \\ \text{from} \end{array} \begin{bmatrix} 0 & 1 & 0 & 0 & 0 \\ 0 & 0 & 1 & 0 & 0 \\ 0 & 0 & 0 & 0 & 1 \\ 0 & 1 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 \end{bmatrix}$$



$A^k: (i, j) \Rightarrow$ how many ways from node $i \rightarrow j$ with path length k

Incidence Matrix

$$\begin{array}{c} e_1 \ e_2 \ e_3 \ e_4 \\ \hline v_1 & 1 & 0 & 0 & 0 \\ v_2 & 1 & 1 & 0 & 1 \\ v_3 & 0 & 1 & 1 & 0 \\ v_4 & 0 & 0 & 0 & 1 \\ v_5 & 0 & 0 & 1 & 0 \end{array}$$

$$\begin{array}{c} a_1 \ a_2 \ a_3 \ a_4 \\ \hline v_1 & 1 & 0 & 0 & 0 \\ v_2 & -1 & 1 & 0 & -1 \\ v_3 & 0 & -1 & 1 & 0 \\ v_4 & 0 & 0 & 0 & 1 \\ v_5 & 0 & 0 & -1 & 0 \end{array}$$

1: out (go from)
-1: in (go to)

Trees & Forest

→ Tree: undirected graph: connected & acyclic (not forming loop)
directed graph: exists a unique walk from root to all leaf

→ Forest: trees

Search Algorithms

Properties:

completeness, optimality, time complexity, space complexity
 guaranteed solution opt solution is (running time) (memory required)
 will be found if exists always found

Uninformed search

Depth-first search



go down until you can't

Breadth-first search



1 layer at a time

Minimum Spanning Tree (MST)

- + Subgraph of an undirected graph G
- + If edges have weights \Rightarrow MST can be found
- + Weight of tree := \sum weight of edges
- + MST: tree with minimal weight

MST heuristic: create MST for an undirected graph

Prim's Algorithm

- + Init at random node
- + In all available edges choose the minimum
- + Till all vertices are in the tree

expand from previous node

Kruskal's Algorithm

- + Add the smallest edge, as long as no circle is created, till all vertices \in tree

Dijkstra's Algorithm

Only positive weights

+ Init arbitrary

+ Follow expand closest distance at 1 vertex

Have table

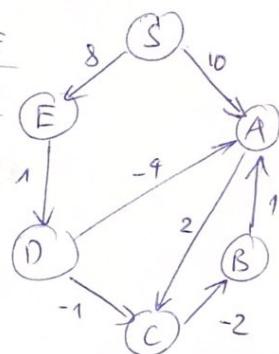
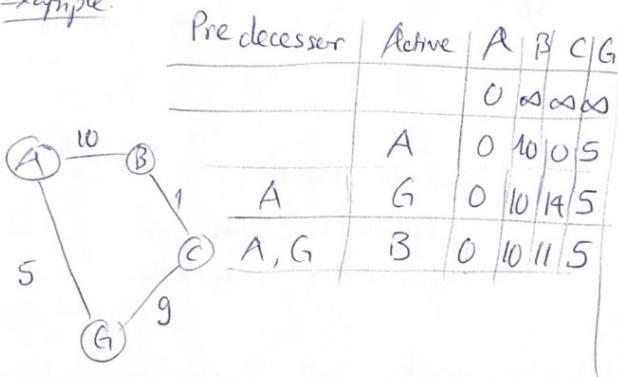
+ Each step with

+,

+ At all nodes, but in order (left \Rightarrow right)

Till converge

Example:



Use when we require to find the shortest path

$$\oplus \quad A^*: f(x) = g(x) + h(x)$$

value of

state x

cost to

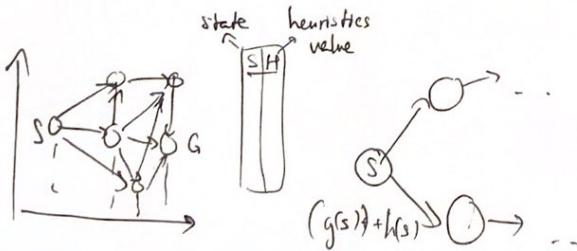
get to x

the heuristic

from x to the final goal

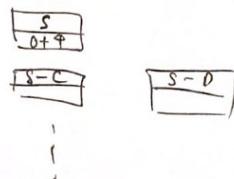
- Expand node with smallest $f(x)$, until every path is explored

A^* as search tree



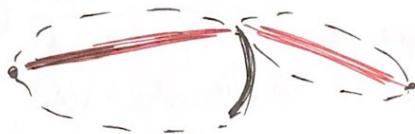
Closed list = { S, ... }

A^* with priority queue



60

- Statically weighted A^* : $h_w(v) = h_{a(v)} \cdot \varepsilon$, $\varepsilon > 1$
- Dynamically weighted A^* : $f(v) = g(v) + (1 + \varepsilon w(v))h(v)$
 $w(v) = \begin{cases} 1 - \frac{d(v)}{N} & \text{if } d(v) \leq N \\ 0 & \text{else} \end{cases}$, N is the expected path steps needed

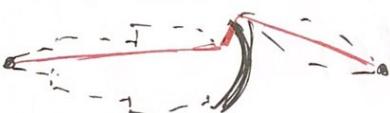


A^* : optimal solution

computational-expensive &
high memory consumption



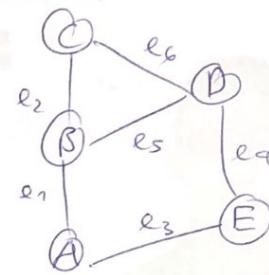
Statically weighted A^* : faster to a solution
loses solution quality



Dynamically weighted A^* : trade-off between
speed & quality

T1/ a) Adjacency matrix

$$A = \begin{bmatrix} A & B & C & D & E \\ A & 0 & 1 & 0 & 0 & 1 \\ B & 1 & 0 & 1 & 1 & 0 \\ C & 0 & 1 & 0 & 1 & 0 \\ D & 0 & 1 & 1 & 0 & 1 \\ E & 1 & 0 & 0 & 1 & 0 \end{bmatrix}$$



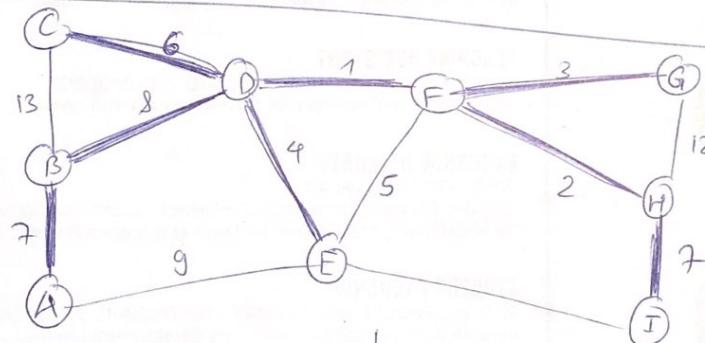
b) Number of paths from E → C with length 3

$$\Rightarrow A^3 = \begin{bmatrix} 0 & 5 & 2 & 14 \\ 5 & 2 & 4 & 61 \\ 2 & 4 & 2 & 42 \\ 1 & 6 & 4 & 25 \\ 4 & 1 & 2 & 50 \end{bmatrix} \Rightarrow 2 \text{ possible path } \begin{array}{l} EAABC \\ EDABC \end{array}$$

c) $G = (V, E)$ where $V = \{A, B, C, D, E\}$

$$E = \{e_1 \dots e_6\}$$

T2/



Prim (start node D)

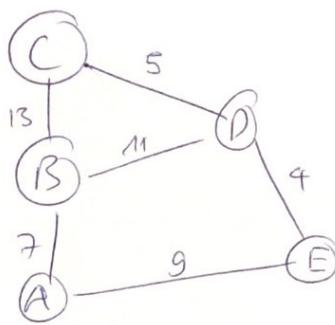
Step	Nodes	Edges
1	D	-
2	D, F	e _{DF}
3	D, F, H	e _{DF} , e _{FH}
4	D, F, H, G	e _{DF} , e _{FH} , e _{FG}
5	D, F, H, G, E	e _{DF} , e _{FH} , e _{FG} , e _{FE}
6	D, F, H, G, E, C	e _{DF} , e _{FH} , e _{FG} , e _{FE} , e _{DE} , e _{DC}
7	D, F, H, G, E, C, I	e _{DF} , e _{FH} , e _{FG} , e _{FE} , e _{DE} , e _{DC} , e _{CI} , e _{IH}
8	D, F, H, G, E, C, I, B	- - - - - e _{DB}
9	...	A - - - - - e _{BA}

Kruskal

Step	Nodes	Edges
1	D, F	e _{DF}
2	D, F, H	e _{DF} , e _{FH}
3	D, F, H, G	e _{DF} , e _{FH} , e _{FG}
4	D, F, H, G, E	e _{DF} , e _{FH} , e _{FG} , e _{FE}
5	D, F, H, G, E, C	e _{DF} , e _{FH} , e _{FG} , e _{FE} , e _{DE} , e _{PC}
6	D, F, H, G, E, C, I, A, B	e _{DF} , e _{FH} , e _{FG} , e _{FE} , e _{DE} , e _{PC} , e _{AB} , e _{HI}
7	- - - - -	- - - - - e _{BD}

T3) Dijkstra from A

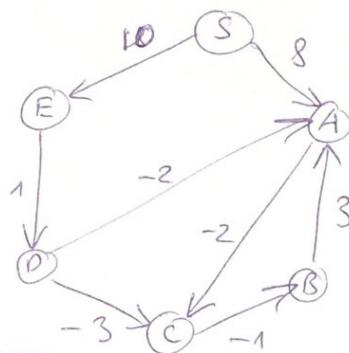
Predecessor	Active	A	B	C	D	E
		0	∞	∞	∞	∞
A	A	0	7	∞	∞	9
A	B	0	7	20	18	9
AB	E	0	7	20	13	9
ABE	D	0	7	18	13	9
ABED	C	0	7	18	13	9



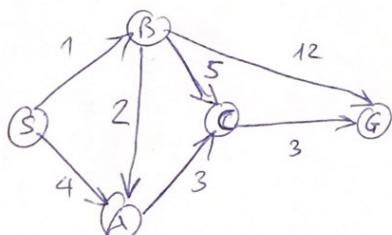
T4) Bellman and Ford algorithm

	S	A	B	C	D	E
Init	0	∞	∞	∞	∞	∞
Loop	0	8	5	6	11	10

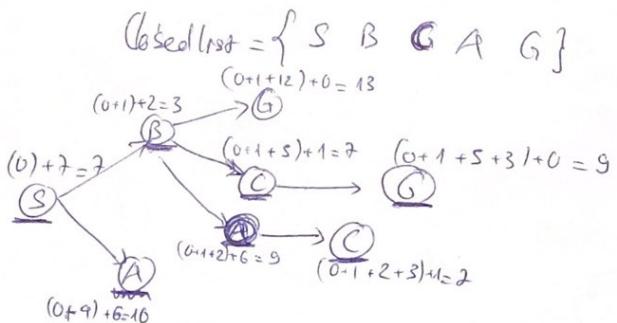
0 8 5 6 11 10



T5) A* algorithm with closed list



State	H
S	2
A	6
B	2
C	1
G	0

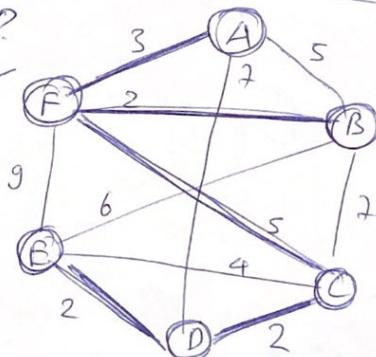


T6) $W_{MST} \leq 2 W_{opt}$, $W_{MST} = 14$?

Euler Circle: FAFBECDE~~E~~CF

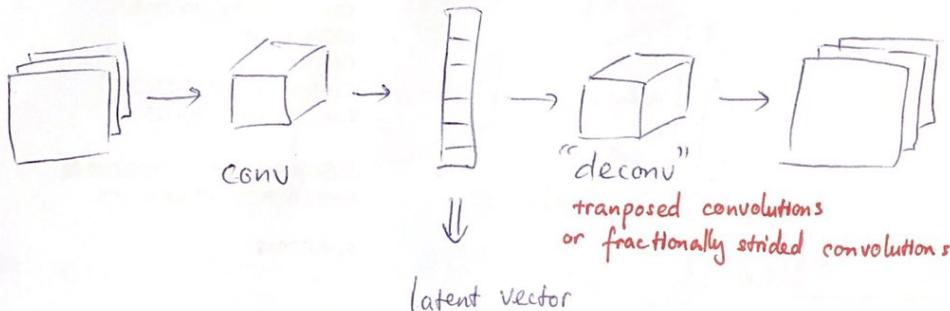
$$\Rightarrow FA \setminus B \setminus CD \setminus EF = 27$$

$$\Rightarrow \text{Weight of the ST} = \frac{27}{2} \approx 14$$



Variational Auto encoders

- Normal: Auto encoder



Normal Auto Encoder

- Learns hidden representation of input (\rightarrow learn latent vector)
- Can not generate new data
- Reconstruction loss

Variational Auto Encoder

- Learns to generate new data

what is optimized

Reconstruction loss + latent loss

Images are more blurry than GAN
(Generative Adversarial Nets)

The more dimensions in the latent vector

sharper, clearer outputs
needs more data

Application:

- + Denoising Autoencoder: add noise to init img, try loss with no noise
- + Neural inpainting: crop out a part, remove watermarks...

Normal AE

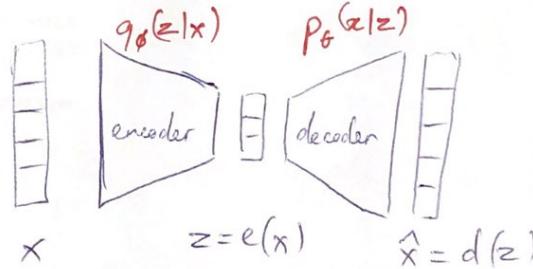
Input \Rightarrow fixed vector

VAE

Input \Rightarrow distribution

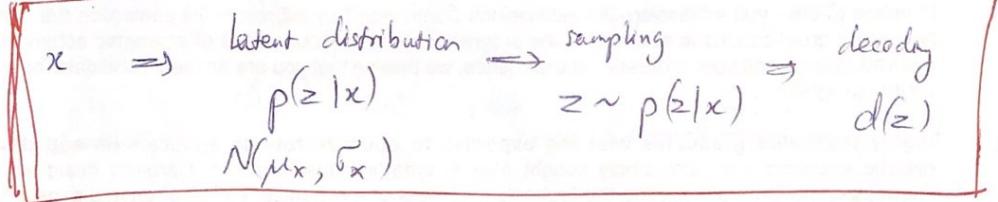
mean vector
standard deviation

- VAE loss:



Autoencoder loss = $\|x - \hat{x}\|^2 = \|x - d(e(x))\|^2$

VAE:



$$\text{loss} = \|x - d(z)\|^2 + \text{KL}[N(\mu_x, \sigma_x^2), N(0, 1)]$$

$$L(\theta, \phi; x, z) = \mathbb{E}_{q_\phi(z|x)} \log p_\theta(x|z) - D_{\text{KL}}(q_\phi(z|x) \parallel p(z))$$

(Explanation behind, last page)

+ Reparametrized trick: $z = \mu + \sigma \cdot \epsilon$ ($\epsilon \sim \text{Normal}(0, 1)$)
 \Rightarrow to ensure gradients flow

+ Disentangled VAE: add β to D_{KL} loss

β too small \Rightarrow overfitting

β too big \Rightarrow loose high definition details ... hurt performance

Kullback - Leibler Divergence

Kind of a distance metric between 2 distributions

Entropy: $H = -\sum_{i=1}^N p(x_i) \log p(x_i)$

$$D_{KL}(p||q) = \sum_{i=1}^N p(x_i) (\log p(x_i) - \log q(x_i)) = \sum_{i=1}^N p(x_i) \log \frac{p(x_i)}{q(x_i)}$$
$$= E [\log p(x) - \log q(x)]$$

⇒ How many bits of info we expect to lose

⇒ A function that we can optimize

$$\Rightarrow H(p, q) = H(p) + D_{KL}(p||q)$$

Cross entropy = entropy + KL Divergence

	State	Action model
classical planning	Observable	Deterministic, exact
Markov Decision Processes	Observable	Stochastic (Non-deterministic)
Partially observable Markov Decision Processes	Partially observable	Stochastic

KL Divergence between 2 normal distribution: $\mathcal{N}(\mu_1, \sigma_1^2) \times \mathcal{N}(\mu_2, \sigma_2^2)$

$$KL(p, q) = \log \frac{\sigma_2}{\sigma_1} + \frac{\sigma_1^2 + (\mu_1 - \mu_2)^2}{2\sigma_2^2} - \frac{1}{2}$$

CODE: μ, σ = encoder (\hat{x})

$z = \mu + \sigma * \text{random-normal}(0, 1)$

$y = \text{decoder}(z)$

recon-loss = $x \cdot \log(y) + (1-x) \log(1-y)$

KL-loss = $\frac{1}{2} [\text{sqn}(\mu) + \text{sqn}(\sigma) - \log(\text{sqn}(\sigma) + 1e^{-8}) - 1]$

$$\begin{aligned}\mu_2 &= 0 \\ \sigma_2 &= 1\end{aligned}$$

www.academy.rwth-aachen.de | masters@academy.rwth-aachen.de

ELBO = recon-loss + KL-loss

loss = -ELBO

The math behind VAE

VAE: learn model parameters to maximize likelihood of training data:

$$p_{\theta}(x) = \int p_{\theta}(z) p_{\theta}(x|z) dz$$

are chosen
(simple Gaussian)

Decoder neural network

$$\begin{aligned} \log p_{\theta}(x^{(i)}) &= E_{z \sim q_{\phi}(z|x^{(i)})} [\log p_{\theta}(x^{(i)})] \\ &= E_z \left[\log \frac{p_{\theta}(x^{(i)}|z) p_{\theta}(z)}{p_{\theta}(z|x^{(i)})} \cdot \frac{q_{\phi}(z|x^{(i)})}{q_{\phi}(z|x^{(i)})} \right] \\ &= E_z [\log p_{\theta}(x^{(i)}|z)] - E_z \left[\log \frac{q_{\phi}(z|x^{(i)})}{p_{\theta}(z)} \right] + E_z \left[\log \frac{q_{\phi}(z|x^{(i)})}{p_{\theta}(z|x^{(i)})} \right] \\ &= \underline{E_z [\log p_{\theta}(x^{(i)}|z)]} - D_{KL}(q_{\phi}(z|x^{(i)}) || p_{\theta}(z)) + D_{KL}(q_{\phi}(z|x^{(i)}) || p_{\theta}(z|x^{(i)})) \end{aligned}$$

Tractable lower bound

\Rightarrow can take gradient and optimized



intractable, can't compute

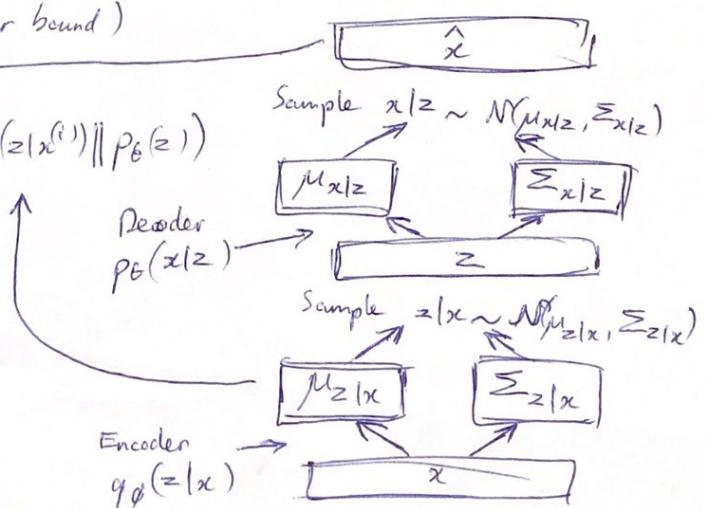
But, we know it ≥ 0

Intuition: the higher the probability, the closer the output to input (similar to maximum likelihood) || Make approximation posterior distribution close to prior

$$\log p_{\theta}(x^{(i)}) \geq L(x^{(i)}, \theta, \phi) \Rightarrow \theta^*, \phi^* = \underset{\theta, \phi}{\operatorname{arg\ max}} \sum_{i=1}^N L(x^{(i)}, \theta, \phi)$$

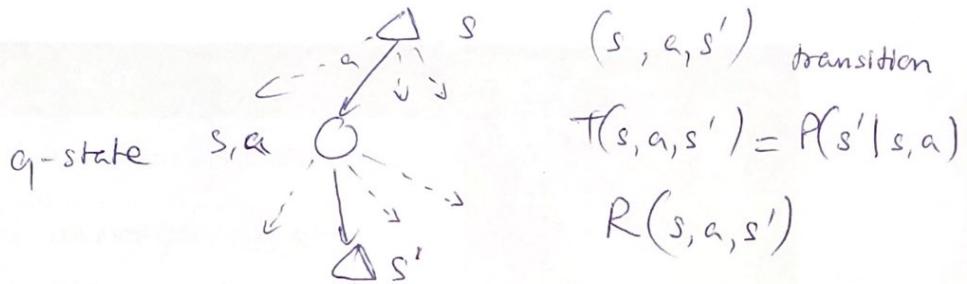
Variational lower bound "ELBO"
(evidence lower bound)

$$\begin{aligned} &E_z [\log p_{\theta}(x^{(i)}|z)] - D_{KL}(q_{\phi}(z|x^{(i)}) || p_{\theta}(z)) \\ &= L(x^{(i)}, \theta, \phi) \end{aligned}$$



MDP

MDP state Projects on expectimax-like search tree



Finite vs infinite Horizon

To solve infinite utilities:

finite horizon

discounting

absorbing state (like the firehose over heating)

MDP:

- Policy: choice of action (at each state)
- Utility: sum of (discounted) rewards

$V^*(s)$: value of a state: expected utility starting in s and acting optimally

$Q^*(s, a)$: expected utility starting in s and having taken a , then act optimally

$\pi^*(s)$: optimal action/policy from state s

$$V^*(s) = \max_a Q^*(s, a)$$

$$Q^*(s, a) = \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V^*(s')]$$

$$\Rightarrow V^*(s) = \max_a \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V^*(s')]$$

$$v(s) = \mathbb{E}[G_t | S_t = s] = \mathbb{E}[R_{t+1} + \gamma v(S_{t+1}) | S_t = s] = R_s + \gamma \sum_{s'} P_{ss'} v(s')$$

Bellman
equations

Lecture 8

- Markov process: random process without memories

is a $\boxed{\text{tuple } \langle S, P \rangle}$

state transition matrix

$P = s []$

$$P_{ss'} = \mathbb{P} [S_{t+1} = s' | S_t = s]$$

probabilities from s to s'

- Markov reward process: $\boxed{\text{tuple } \langle S, R, P, \gamma \rangle}$

state, transition matrix, reward func., discount factor $\in [0, 1]$

$$R_s = \mathbb{E} [R_{t+1} | S_t = s]$$

- Utility of state sequences G_t

$$G_t = U([s_t, s_{t+1}, \dots, s_\infty]) = \sum_{k=0}^{\infty} \gamma^k R(s_{t+k}) \in \frac{R_{\max}}{1-\gamma} (\gamma < 1)$$

$\gamma \rightarrow 0 \Rightarrow$ short-sighted evaluation

$\gamma \rightarrow 1 \Rightarrow$ far-sighted

$\gamma = 1 \Rightarrow$ additive utility

The Stupor

5 axioms of rational decisions

order $(A \succ B) \vee (B \succ A) \vee (A \sim B)$

transitivity $(A \succ B) \wedge (B \succ C) \Rightarrow (A \succ C)$

continuity $A \succ B \succ C \Rightarrow \exists p [p, A; 1-p, C] \sim B$

substitution: $(A \sim B) \Rightarrow [p, A; 1-p, C] \sim [p, B; 1-p, C]$

monotony: $A \succ B \Rightarrow (p \geq q \Leftrightarrow [p, A; 1-p, B] \succ [q, A; 1-q, B])$

- Value iteration: $V_{k+1}(s) \leftarrow \max_a \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V_k(s')]$

Repeat until convergence: unique (fixed) optimal value

But, policy may converge long before values do

- When the policy is fixed, Bellman equation is a linear equation:

$$v = R + \gamma Pv \quad \begin{bmatrix} v(1) \\ v(2) \\ \vdots \\ v(n) \end{bmatrix} = \begin{bmatrix} R_1 \\ \vdots \\ R_n \end{bmatrix} + \gamma \begin{bmatrix} P_{11} & \dots & P_{1n} \\ \vdots & \ddots & \vdots \\ P_{nn} & \dots & P_{nn} \end{bmatrix} \begin{bmatrix} v_1 \\ \vdots \\ v_n \end{bmatrix}$$

$\Rightarrow v = (I - \gamma P)^{-1} R$
(convergence)

Dynamic Programming | Value iteration

Synchronous

for all $s \in S$

$$v_{\text{new}}(s) \leftarrow \max_a (R_s^a + \gamma \sum_{s'} P_{ss'}^a v_{\text{old}}(s'))$$

$$v_{\text{old}} \leftarrow v_{\text{new}}$$

calculate all expected value $v^*(s)$

Update all at once

for all $s \in S$:

$$v(s) \leftarrow \max_a (R_s^a + \gamma \sum_{s'} P_{ss'}^a v(s'))$$

calculate each $v^*(s)$

then update immediately

Value iteration

- 1, Init: all expected utility = 0, except terminal state
- 2, Start at random state
- 3, Each state, try all action
- 4, ...
- 5, Draw out π^* for the next iteration
- 6, Continue until convergence

POMDP

Partially observable MDP

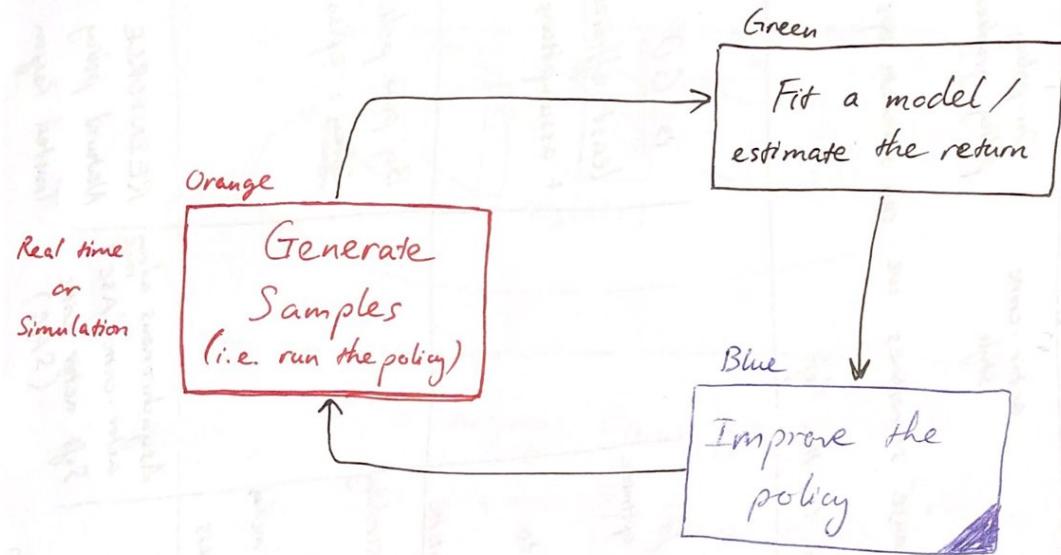
tuple $\langle S, A, \mathcal{O}, P, R, Z, \gamma \rangle$

state, action, observation, transition matrix, reward, observation function, discount

$$Z_{s', o}^a = P[O_{t+1} = o | S_{t+1} = s', A_t = a]$$

$$\Rightarrow \begin{cases} v^*(b) \leftarrow \max_a (R_b^a + \gamma \sum_{s'} P_{bb'}^a v(b')) \\ \pi^*(b) = \arg \max_a (R_b^a + \gamma \sum_{s'} P_{bb'}^a v(b')) \end{cases}$$

Reinforcement - Learning



Structure of RL algorithms

Some definition:

- + $Q^\pi(s_t, a_t) = \sum_{t'=t}^T E_{\pi_\theta}[r(s_{t'}, a_{t'}) | s_t, a_t]$: the expectation of total reward from your current time step t till the end, under your policy π_θ
- + $V^\pi(s_t) = \sum_{t'=t}^T E_{\pi_\theta}[r(s_{t'}, a_{t'}) | s_t] = E_{a_t \sim \pi(a_t | s_t)}[Q^\pi(s_t, a_t)]$: the expectation of total reward from s_t

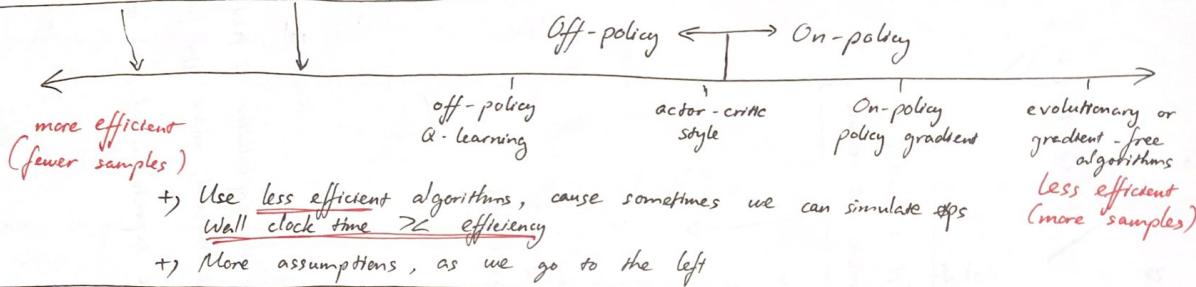
⊗ Why are there so many methods/algorithms?

There are trade-off: sampling efficiency & stability and ease of use & assumptions: stochastic or deterministic continuous or discrete episodic or infinite horizon

Model based		Value func fitting methods	Actor - critic algorithms	Policy Gradient
shallow RL	Deep RL			

Sample efficiency

(How many samples do we need to get good policy?)



Stability & ease of use

- Does it converge?
- If yes, to what?
- Does it ALWAYS converge?

RL: often not gradient descent

model will converge

BUT: better model

NOT GUARANTEE

better policy

minimize error of fit
at worst, doesn't optimize anything

can not prove to converge

\Rightarrow more like heuristic

is GD
least efficient
+ assumptions

Assumption

By some: episodic learning

Generally: full observability

By some continuous methods:
continuity / smoothness

By pure policy gradient methods,
Often: episodic learning

Example

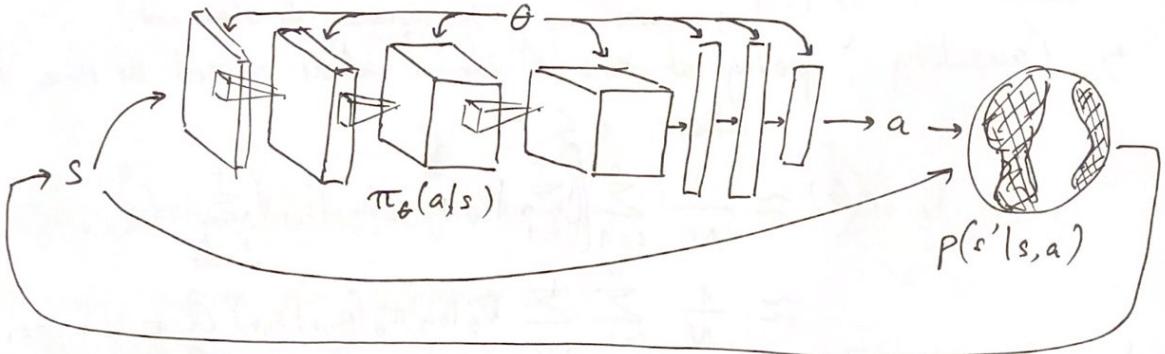
Dyna
Guided policy search

Q - learning
DQN
Temporal difference
Fitted value iteration

Asynchronous advantage actor-critic (A3C)
Soft actor critic (SAC)

REINFORCE
Natural policy gradient
Trusted Region policy optimization

Policy Gradient



Goal:

$$\theta^* = \arg \max_{\theta} E_{(s, a) \sim p_\theta(s, a)} [r(s, a)]$$

$$= \arg \max_{\theta} \sum_{t=1}^T E_{(s_t, a_t) \sim p_\theta(s_t, a_t)} [r(s_t, a_t)] \quad \begin{cases} \text{infinite,} \\ \text{finite} \end{cases} \quad \begin{cases} \text{horizon case} \end{cases}$$

\Rightarrow maximize: $J(\theta) = E_{\tau \sim p_\theta(\tau)} [r(\tau)] = \int \pi_\theta(\tau) r(\tau) d\tau$

$$\approx \frac{1}{N} \sum_i \sum_t r(s_{i,t}, a_{i,t})$$

sum over samples \nearrow sum over time steps \nwarrow

The Gradient:

$$\nabla_\theta J(\theta) = \int \nabla_\theta \pi_\theta(\tau) r(\tau) d\tau = \int \pi_\theta(\tau) \nabla_\theta \log \pi_\theta(\tau) r(\tau) d\tau$$

$$= E_{\tau \sim \pi_\theta(\tau)} \left[\left(\sum_{t=1}^T \nabla_\theta \log \pi_\theta(a_t | s_t) \right) \left(\sum_{t=1}^T r(s_t, a_t) \right) \right]$$

$$\nabla_\theta J(\theta) \approx \frac{1}{N} \sum_{i=1}^N \left[\left(\sum_{t=1}^T \nabla_\theta \log \pi_\theta(a_{i,t} | s_{i,t}) \right) \left(\sum_{t=1}^T r(s_{i,t}, a_{i,t}) \right) \right]$$

then $\theta \leftarrow \theta + \alpha \nabla_\theta J(\theta)$

REMARKS: somewhat like MLE, makes good stuff happen more, bad stuff less

④ Partial observability:

$$\nabla_\theta J(\theta) \approx \frac{1}{N} \sum_{i=1}^N \left[\left(\sum_{t=1}^T \nabla_\theta \log \pi_\theta(a_{i,t} | o_{i,t}) \right) \left(\sum_{t=1}^T r(s_{i,t}, a_{i,t}) \right) \right]$$

$$o_{i,t} \rightarrow \text{network} \rightarrow a_{i,t}$$

$$\pi_\theta(a_{i,t} | o_{i,t})$$

② Problem of Policy Gradient: HIGH VARIANCE with $r(\epsilon)$

Different samples \Rightarrow different gradient estimate

For a small finite No. Samples \Rightarrow gradient noisy

(At the beginning, policy θ is not good \Rightarrow random action \Rightarrow the not-so-good action-result accumulate \Rightarrow high variance in the end)

\Rightarrow Causality: policy at time t' cannot affect reward at time t
when $t \leq t'$

$$\begin{aligned}\Rightarrow \nabla_{\theta} J(\theta) &\approx \frac{1}{N} \sum_{i=1}^N \left(\sum_{t=1}^T \nabla_{\theta} \log \pi_{\theta}(a_{i,t} | s_{i,t}) \left(\sum_{t'=t}^T r(s_{i,t'}, a_{i,t'}) \right) \right) \\ &\approx \frac{1}{N} \sum_{i=1}^N \sum_{t=1}^T \nabla_{\theta} \log \pi_{\theta}(a_{i,t} | s_{i,t}) \hat{Q}_{i,t}\end{aligned}$$

reward to-go \times

\Rightarrow Baselines: $b = \frac{1}{N} \sum_{i=1}^N r(\epsilon)$

$$\Rightarrow \nabla_{\theta} J(\theta) \approx \frac{1}{N} \sum_{i=1}^N \nabla_{\theta} \log \pi_{\theta}(\epsilon) [r(\epsilon) - b] = Q(s_t, a_t) - V(s_t)$$

REINFORCE algorithm:

1, Sample $\{\epsilon^i\}$ from $\pi_{\theta}(a_t | s_t)$

$$2, \nabla_{\theta} J(\theta) \approx \frac{1}{N} \sum_{i=1}^N \nabla_{\theta} \log \pi_{\theta}(\epsilon) \left(\sum_{t'=t}^T r(s_{i,t'}, a_{i,t'}) - b \right) \quad b = \frac{1}{N} \sum_{i=1}^N \sum_{t'=t}^T r(s_{i,t'}, a_{i,t'})$$

3, $\theta \leftarrow \theta + \alpha \nabla_{\theta} J(\theta)$

Pseudo code:

logits = policy.predictions(states)

negative_likelihoods = tf.nn.softmax_cross_entropy(labels=actions, logits)

weighted_negative_likelihoods = tf.multiply(negative_likelihoods, q-values)

loss = tf.reduce_mean(weighted_negative_likelihoods)

gradients = loss.gradients(loss, variables)

with causality & baselines

\Rightarrow Can derive off-policy policy gradient with IMPORTANT SAMPLING

$$\bullet E_{x \sim p(x)}[f(x)] = E_{x \sim q(x)} \left[\frac{p(x)}{q(x)} f(x) \right] \Rightarrow \nabla_{\theta} J(\theta) = E_{\epsilon \sim \pi_{\theta_{old}}(\epsilon)} \left[\frac{\pi_{\theta}(\epsilon)}{\pi_{\theta_{old}}(\epsilon)} \nabla_{\theta} \dots \right]$$

Actor - Critic

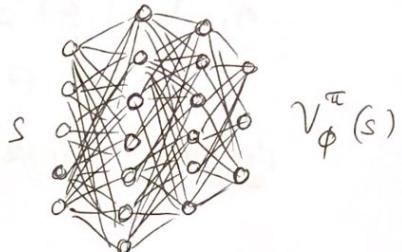
the policy - value function (a.k.a. policy evaluation)

An extra network:

$$Q^\pi(s_t, a_t) = \sum_{t'=t}^T E_{\pi_\theta}[r(s'_t, a'_t) | s_t, a_t]$$

$$V^\pi(s_t) = E_{a_t \sim \pi_\phi(a_t | s_t)}[Q^\pi(s_t, a_t)]$$

$$\Rightarrow A^\pi(s_t, a_t) = Q^\pi(s_t, a_t) - V^\pi(s_t)$$



- We can rewrite the gradient as:

$$\nabla_\theta J(\theta) \approx \frac{1}{N} \sum_{i=1}^N \sum_{t=1}^T \nabla_\theta \log \pi_\theta(a_{i,t} | s_{i,t}) A^\pi(s_{i,t} | a_{i,t})$$

As: $Q^\pi(s_t, a_t) = r(s_t, a_t) + E_{s_{t+1} \sim P(s_{t+1} | s_t, a_t)} [V^\pi(s_{t+1})]$

estimate as $\approx r(s_t, a_t) + V^\pi(s_{t+1})$ (in stead of the average, take value of 1 sample)

Then we have: $A^\pi(s_t, a_t) \approx r(s_t, a_t) + V^\pi(s_{t+1}) - V^\pi(s_t)$

= What to fit to:

$$y_{i,t} = \sum_{t'=t}^T E_{\pi_\theta}[r(s'_{t'}, a'_{t'} | s_{i,t})] \approx r(s_{i,t}, a_{i,t}) + V^\pi(s_{i,t+1})$$

$$\approx r(s_{i,t}, a_{i,t}) + \hat{V}_\phi^\pi(s_{i,t+1})$$

Bootstrap estimate

$$\mathcal{L}(\phi) = \frac{1}{2} \sum_i \| \hat{V}_\phi^\pi(s_i) - y_i \|^2$$

$$y_{i,t} = \sum_{t'=t}^T r(s_{i,t'}, a_{i,t'}) \quad \text{Monte-Carlo}$$

(do as before, estimate by running it in a real world
⇒ a single sample estimate)

Batch Actor-Critic algorithm:

- 1, Sample $\{s_i, a_i\}$ from $\pi_\theta(a|s)$
- 2, Fit $V_\phi^\pi(s)$ to sampled reward sum $\underbrace{\text{Monte Carlo}}_{\text{Bootstrap estimate}} \dots + \mathcal{L}(\phi) = \frac{1}{2} \dots$
- 3, Evaluate $\hat{A}^\pi(s_i, a_i) = r(s_i, a_i) + \hat{V}_\phi^\pi(s_{i'}) - \hat{V}_\phi^\pi(s_i)$
- 4, $\nabla_\theta J(\theta) \approx \sum \nabla_\theta \log \pi_\theta(a_i | s_i) \hat{A}^\pi(s_i, a_i)$
- 5, $\theta \leftarrow \theta + \alpha \nabla_\theta J(\theta)$

With discount factor:

$$3, \hat{A}^\pi(s_i, a_i) = r(s_i, a_i) + \gamma \hat{V}_\phi^\pi(s_{i'}) - \hat{V}_\phi^\pi(s_i)$$

On line Actor-Critic Algorithm:

- 1, Take action $a \sim \pi_\theta(a|s)$ get (s, a, s', r)
- 2, Update \hat{V}_ϕ^π using target value $r + \gamma \hat{V}_\phi^\pi(s')$
- 3, Evaluate $\hat{A}^\pi(s, a) = r(s, a) + \gamma \hat{V}_\phi^\pi(s') - \hat{V}_\phi^\pi(s)$
- 4, $\nabla_\theta J(\theta) \approx \nabla_\theta \log \pi_\theta(a|s) \hat{A}^\pi(s, a)$
- 5, $\theta \leftarrow \theta + \alpha \nabla_\theta J(\theta)$

Problem: single batch \Rightarrow need parallel actor-critic
(synchronous / asynchronous)

+ Architecture design: 2 networks vs shared network design
simple & stable vs two different gradients
no shared features which need to tuned, balanced

+ Critic as state-dependent baselines

$$\nabla_\theta J(\theta) \approx \frac{1}{N} \sum_{i=1}^N \sum_{t=1}^T \nabla_\theta \log \pi_\theta(a_{i,t}|s_{i,t}) \left(\left(\sum_{t'=t}^T \gamma^{t'-t} r(s_{i,t'}, a_{i,t'}) \right) - \hat{V}_\phi^\pi(s_{i,t}) \right)$$

+ Control variates: action-dependent baselines

Set ~~b~~ $b = \hat{V}_\phi^\pi(s_{i,t}) \Rightarrow \begin{cases} \text{no bias} \\ \text{lower variance} \end{cases}$

+ Eligibility traces & n-step returns:

Actor-critic: $\hat{A}_c^\pi(s_t, a_t) = r(s_t, a_t) + \gamma \hat{V}_\phi^\pi(s_{t+1}) - \hat{V}_\phi^\pi(s_t)$ low variance but biased

Monte Carlo: $\hat{A}_{MC}^\pi(s_t, a_t) = \sum_{t'=t}^{\infty} \gamma^{t'-t} r(s_{t'}, a_{t'}) - \hat{V}_\phi^\pi(s_t)$ no bias but higher variance

$$\Rightarrow \hat{A}_n^\pi(s_t, a_t) = \sum_{t'=t}^{t+n} \gamma^{t'-t} r(s_{t'}, a_{t'}) - \hat{V}_\phi^\pi(s_t) + \gamma^n \hat{V}_\phi^\pi(s_{t+n})$$

+ Generalized advantage estimation

$$\hat{A}_{GAE}^\pi(s_t, a_t) = \sum_{n=1}^{\infty} \omega_n \hat{A}_n^\pi(s_t, a_t) \quad \omega_n \propto \gamma^{n-1}$$

Policy iteration

Can we omit policy gradient completely?

High level idea - Algorithm:

1) Evaluate $A^\pi(s, a)$

2) Set $\pi \leftarrow \pi'$

Dynamic Programming

+ Assume: we know: $p(s'|s, a)$

s & a are both discrete & small \Rightarrow can store $V(s)$ in table

$$\pi'(a_t | s_t) = \begin{cases} 1 & \text{if } a_t = \arg \max_{a_t} A^\pi(s_t, a_t) \\ 0 & \text{otherwise} \end{cases}$$

π' : simply choose the current best action

Policy iteration with Dynamic Programming

$$1, V^\pi(s) \leftarrow r(s, \pi(s)) + \gamma \mathbb{E}_{s' \sim p(s'|s, \pi(s))}[V^\pi(s')]$$

$$2, \text{Set } \pi \leftarrow \pi'$$

Simpler

Value iteration Algorithm

$$\text{Since: } \arg \max_{a_t} A^\pi(s_t, a_t) = \arg \max_{a_t} Q^\pi(s_t, a_t)$$

$$1, \text{Set } Q^\pi(s, a) \leftarrow r(s, a) + \gamma \mathbb{E}[V^\pi(s')]$$

$$2, \text{Set } V^\pi(s) \leftarrow \max_a Q^\pi(s, a)$$

REMARKS: Still have a big table

when facing the Curse of Dimension (of s)

\Rightarrow Use neural network to evaluate $V(\cdot)$, $Q(\cdot, \cdot)$

Fitted Value Iteration

$$1, y_i \leftarrow \max_{a_i} r(s_i, a_i) + \gamma \mathbb{E}[V_\phi(s'_i)]$$

$$2, \phi \leftarrow \arg \min_{\phi} \frac{1}{2} \sum_i \|V_\phi(s_i) - y_i\|^2$$



Fully Fitted Q Iteration

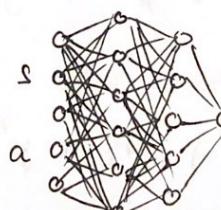
$$\mathbb{E}[V(s'_i)] \approx \max_{a'_i} Q_\phi(s'_i, a'_i)$$

1, Collect: $\{(s_i, a_i, s'_i, r_i)\}$ using some policy

2, Set $y_i \leftarrow r(s_i, a_i) + \gamma \max_{a'_i} Q_\phi(s'_i, a'_i)$

3, $\phi \leftarrow \arg \min_{\phi} \frac{1}{2} \sum_i \|Q_\phi(s_i, a_i) - y_i\|^2$

- off-policy



$Q_\phi(s, a)$ - not really converge

BUT: Still need to know transition dynamic

~~⊗ Online Q iteration algorithm:~~

- 1) Take some action a_i & get (s_i, a_i, s'_i, r_i)
- 2) $y_i = r(s_i, a_i) + \gamma \max_{a'} Q_\phi(s'_i, a')$
- 3, $\phi \leftarrow \phi - \alpha \frac{dQ_\phi}{d\phi}(s_i, a_i)(Q_\phi(s_i, a_i) - y_i)$

1 sample off policy
1 gradient step

+ How to choose action in step 1: Exploration vs Exploitation

- ⇒ Epsilon greedy
- Boltzmann exploration (very large or continuous action-space)

Problems with Online Q-learning:

- Sequential states are strongly correlated ⇒ Replay buffer
- Target value is always changing ⇒ Target network

"Classic" Deep Q-learning (DQN) with replay buffer & target network

- 1, Save target network parameters $\phi' \leftarrow \phi$
- 2, Collect dataset $\{(s_i, a_i, s'_i, r_i)\}$ using some policy, add to buffer \mathcal{B}
- $N \times$
 - 3, Sample a batch (s_i, a_i, s'_i, r_i) from \mathcal{B}
 - 4, $\phi \leftarrow \phi - \alpha \sum_i \frac{dQ_\phi}{d\phi}(s_i, a_i)(Q_\phi(s_i, a_i) - [r(s_i, a_i) + \gamma \max_{a'} Q_\phi(s'_i, a')])$

+ Alternative: 5, Update $\phi' \leftarrow \sigma \phi' + (1-\sigma) \phi$ $\sigma = 0,999$

+ Double Q-learning: helps a lot, solve over-estimate prob, no downside ⇒ should always use
 $y = r + \gamma \mathbb{E}_{a'} Q_{\phi'}(s', \arg \max_a Q_{\phi'}(s', a'))$ instead of $\max_{a'} Q_{\phi'}(s', a')$

+ Multi-Step returns: helps a lot, have DOWNSIDE ⇒ frequently use

$$y_{i,t} = \sum_{t'=t}^{t+N-1} \gamma^{t'-t} r_{i,t'} + \gamma^N \max_{a_{i,t+N}} Q_{\phi'}(s_{i,t+N}, a_{i,t+N})$$

ONLY CORRECT when learning on-policy

Q learning with continuous action-space

Tips for Q -Learning:

- Large replay buffers help improve stability
- Start with high exploration \Rightarrow gradually reduce
- Bellman error gradients can be big \Rightarrow clip gradients / use Huber loss

Model-based RL

Optimal control

- + Deterministic case vs Stochastic
- + Open-loop case vs Closed-loop

Maximize objective through sequence of actions:

$$a_1, \dots, a_T = \underset{a_1 \dots a_T}{\operatorname{argmax}} J(a_1 \dots a_T) \quad A = \underset{A}{\operatorname{argmax}} J(A)$$

Optimal control: discrete space

Random Shooting Method

- + Guess & Check:
 - 1, Pick $A_1 \dots A_N$ from some distribution (e.g. uniform)
 - 2, choose A_i based on $\operatorname{argmax}_i J(A_i)$

Cross-Entropy Method (CEM)

- 1, Sample A_1, \dots, A_N from $P(A)$
- 2, Evaluate $J(A_1) \dots J(A_N)$
- 3, Pick elites A_{i_1}, \dots, A_{i_M} with highest value $M < N$
- 4, Refit $P(A)$ to the elites $A_{i_1} \dots A_{i_M}$

Monte Carlo tree search (MCTS)

- + Discrete case:
 - 1, Find a leaf s_t using Tree Policy(s_t)
 - 2, Evaluate the leaf using Default Policy(s_t)
 - 3, Update all values in tree between s_t and s_f , take best action from s_t

$$\text{UCT Tree Policy}(s_t) : \text{Score}(s_t) = \frac{\alpha(s_t)}{N(s_t)} + 2C \sqrt{\frac{2 \ln N(s_t=1)}{N(s_t)}}$$

All of these is what we do IF we already know the model

Derivatives are hard to come by, but SOMETIMES you CAN, through Physics equation

$$\min_{u_1, \dots, u_T} \sum_{t=1}^T c(x_t, u_t), \quad x_t = f(x_{t-1}, u_{t-1})$$

$$= \min_{u_1, \dots, u_T} c(x_1, u_1) + c(f(x_1, u_1), u_2) + \dots + c(f(f(\dots)), u_T)$$

Special case: LQR - Linear Quadratic Regulator

Assume: $f(\cdot)$ has special structure:

$$f(x_t, u_t) = F_t \begin{bmatrix} x_t \\ u_t \end{bmatrix} + f_t \quad \text{linear dynamics}$$

$$c(x_t, u_t) = \frac{1}{2} \begin{bmatrix} x_t \\ u_t \end{bmatrix}^T C_t \begin{bmatrix} x_t \\ u_t \end{bmatrix} + \begin{bmatrix} x_t \\ u_t \end{bmatrix}^T c_t \quad \text{quadratic cost}$$

$$C_T = \begin{bmatrix} C_{x_T, x_T} & C_{x_T, u_T} \\ C_{u_T, x_T} & C_{u_T, u_T} \end{bmatrix}, \quad c_T = \begin{bmatrix} c_{x_T} \\ c_{u_T} \end{bmatrix}$$

Solve for u_T only:

$$Q(x_T, u_T) = \text{const} + \frac{1}{2} \begin{bmatrix} x_T \\ u_T \end{bmatrix}^T C_T \begin{bmatrix} x_T \\ u_T \end{bmatrix} + \begin{bmatrix} x_T \\ u_T \end{bmatrix}^T c_T \quad (2)$$

$$\text{Set } \nabla_{u_T} Q(x_T, u_T) = 0 \Rightarrow \dots \quad u_T = -C_{u_T, u_T}^{-1} (C_{u_T, x_T} \cdot x_T + c_{u_T})$$

$$= K_T x_T + k_T \quad (1)$$

Replace (1) to (2) $\Rightarrow V(x_T)$

Minimize objective on last action u_T , based on x_T
 u_{T-1} affects objective through linear dynamics func to x_T

Linear case: LQR (continue)

Backward recursion

for $t = T \rightarrow 1$:

$$Q_t = C_t + F_t^T V_{t+1} F_t$$

$$q_t = c_t + F_t^T V_{t+1} f_t + F_t^T v_{t+1}$$

$$Q(x_t, u_t) = \text{const} + \frac{1}{2} \begin{bmatrix} x_t \\ u_t \end{bmatrix}^T Q_t \begin{bmatrix} x_t \\ u_t \end{bmatrix} + \begin{bmatrix} x_t \\ u_t \end{bmatrix}^T q_t$$

$$u_t \leftarrow \underset{u_t}{\operatorname{argmin}} \quad Q(x_t, u_t) = K_t x_t + k_t$$

$$K_t = -Q_{u_t, u_t}^{-1} Q_{u_t, x_t}$$

$$k_t = -Q_{u_t, u_t}^{-1} q_{u_t}$$

$$V_t = Q_{x_t, x_t} + Q_{x_t, u_t} K_t + K_t^T Q_{u_t, x_t} + K_t^T Q_{u_t, u_t} K_t$$

$$v_t = q_{x_t} + Q_{x_t, u_t} k_t + K_t^T Q_{u_t} + K_t^T Q_{u_t, u_t} k_t$$

$$J(x_t) = \text{const} + \frac{1}{2} x_t^T V_t x_t + x_t^T v_t$$

Forward recursion:

for $t = 1 \rightarrow T$:

$$u_t = K_t x_t + k_t$$

$$x_{t+1} = f(x_t, u_t)$$

Non-linear case: DDP (Differential Dynamic Programming)
or Iterative LQR

Every iteration: linearize local dynamic at loca. with Taylor expansion

$$f(x_t, u_t) \approx f(\hat{x}_t, \hat{u}_t) + \nabla_{x_t, u_t} f(\hat{x}_t, \hat{u}_t) \begin{bmatrix} x_t - \hat{x}_t \\ u_t - \hat{u}_t \end{bmatrix}$$

$$c(x_t, u_t) \approx c(\hat{x}_t, \hat{u}_t) + \nabla_{x_t, u_t} c(\hat{x}_t, \hat{u}_t) \begin{bmatrix} x_t - \hat{x}_t \\ u_t - \hat{u}_t \end{bmatrix} + \frac{1}{2} \begin{bmatrix} x_t - \hat{x}_t \\ u_t - \hat{u}_t \end{bmatrix}^T \nabla_{x_t, u_t}^2 c(\hat{x}_t, \hat{u}_t) \begin{bmatrix} x_t - \hat{x}_t \\ u_t - \hat{u}_t \end{bmatrix}$$

$$F_t = \nabla_{x_t, u_t} f(\hat{x}_t, \hat{u}_t)$$

$$c_t = \nabla_{x_t, u_t} c(\hat{x}_t, \hat{u}_t)$$

$$C_t = \nabla_{x_t, u_t}^2 c(\hat{x}_t, \hat{u}_t)$$

Run LQR backward pass on state $\delta x_t = x_t - \hat{x}_t$, $\delta u_t = u_t - \hat{u}_t$

forward pass with real nonlinear dynamics and

$$u_t = K_t(x_t - \hat{x}_t) + k_t + \hat{u}_t$$

Update \hat{x}_t and \hat{u}_t

How to learn model

⊗ Model-based RL v.0.5:

- 1, Run based policy $\pi_0(a_t|s_t) \Rightarrow$ collect $\mathcal{D} = \{(s, a, s')_i\}$
- 2, Learn dynamic model $f(s, a)$ to minimize $\sum_i \|f(s_i, a_i) - s'_i\|^2$
- 3, Plan through $f(s, a)$ to choose actions
- Good based policy should be taken good care
- Particularly effective if can hand-engineer a dynamics representation with our knowledge of physics ... \Rightarrow need to fit only a few parameters

+ Problems: might go beyond to new state distribution different to the one that collected the data

⊗ Model-based RL v.1.0: (start to work with neural network)

- 1, Run base policy $\pi_0(a_t|s_t) \Rightarrow$ collect $\mathcal{D} = \{(s, a, s')_i\}$
- 2, Learn model $f(s, a)$ to minimize $\sum_i \|f(s_i, a_i) - s'_i\|^2$
- 3, Plan through $f(s, a) \Rightarrow$ choose actions
- 4, Execute these actions & add resulting data $\{(s, a, s')_j\}$ to \mathcal{D}

+ Problems: sometimes we make mistakes
 \Rightarrow learn faster if we correct the mistake more often

⊗ Model-based RL v.1.5:

- 1, Run base policy $\pi_0(a_t|s_t) \Rightarrow$ collect $\mathcal{D} = \{(s, a, s')_i\}$
- 2, Learn model $f(s, a)$ to minimize $\sum_i \|f(s_i, a_i) - s'_i\|^2$
- 3, Plan through $f(s, a) \Rightarrow$ choose actions
- 4, Execute the first planned action, observe resulting state s'
- 5, Append (s, a, s') to \mathcal{D}

⊗ Problems: overfitting early, especially with high-dimensional data

⇒ Introduce: uncertainty estimation

3, Take action with high expected reward

⊗ This goes against exploration

Depends on problems: use $\begin{cases} \text{expected} \\ \text{optimistic} \\ \text{pessimistic} \end{cases}$ value / planning

+ How to get uncertainty-aware model

- Not use output entropy: aleatoric / statistical uncertainty
vs epistemic / model uncertainty

- Bayesian neural network: $p(\theta | D) = \prod_i p(\theta_i | D)$

$$p(\theta_i | D) = \mathcal{N}(\mu_i, \sigma_i)$$

complicate

- Bootstrap ensembles: $p(\theta | D) \approx \frac{1}{N} \sum_i \delta(\theta_i)$

Sampled with replacement from D , usually less than 10 models

⊗ How to plan with uncertainty?

In step 3, we do optimal control to choose action

In optimal control, for example LQR:

- Before: $J(a_1, \dots, a_H) = \sum_{t=1}^H r(s_t, a_t)$ where $s_{t+1} = f(s_t, a_t)$

- Now: $J(a_1, \dots, a_H) = \frac{1}{N} \sum_i^N \sum_{t=1}^H r(s_{t,i}, a_{t,i})$ where $f_i(s_{t,i}, a_{t,i}) = \underbrace{\text{distribution}}_{\text{over deterministic model}}$

1, Sample $\theta \sim p(\theta | D)$

2, At each time step t , sample $s_{t+1} \sim p(s_{t+1} | s_t, a_t, \theta)$

3, Calculate $R = \sum_t r(s_t, a_t)$

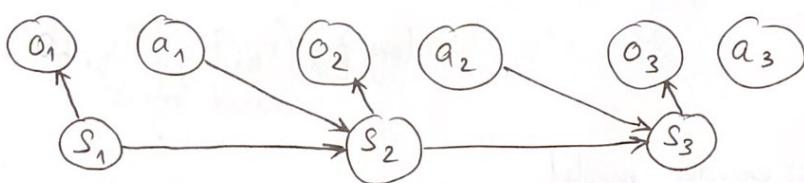
4, Repeat step 1 → 3 and accumulate average reward

→ (Generalized for stochastic model)



to deal with complex observations (images ~)

Model-based RL with state/latent space models



$p(o_t | s_t)$ observation model

$p(s_{t+1} | s_t, a_t)$ dynamics model

$p(r_t | s_t, a_t)$ reward model

+) State space (latent space) models:

Input is given observations, not states

+) Before, all model-based RL algorithms.

Fully observable model

$$\max_{\phi} \frac{1}{N} \sum_{i=1}^N \sum_{t=1}^T \log p_{\phi}(s_{t+1,i} | s_{t,i}, a_{t,i})$$

Now: latent space model

$$\max_{\phi} \frac{1}{N} \sum_{i=1}^N \sum_{t=1}^T E \left[\log p_{\phi}(s_{t+1,i} | s_{t,i}, a_{t,i}) + \log p_{\phi}(o_{t,i} | s_{t,i}) \right]$$

expectation with regards to $(s_t, s_{t+1}) \sim p(s_t, s_{t+1} | o_{1:T}, a_{1:T})$

very complicated

⇒ Can do it by learn approximate posterior $q_{\psi}(s_t | o_{1:t}, a_{1:t})$ encoder

+) Full smoothing posterior: $q_{\psi}(s_t, s_{t+1} | o_{1:T}, a_{1:T})$

look at all the observations & action in trajectory

↑ most accurate

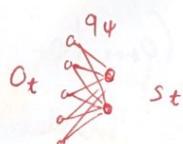
↓ most complicated
(big RNN..)

+) Single-step encoder: $q_{\psi}(s_t | o_t)$

look at only 1 observation

↑ simplest (CNN..)

↓ least accurate



⊗ Simple special case: $q(s_t | o_t)$ is deterministic

$$q_{\psi}(s_t | o_t) = \delta(s_t = g_{\psi}(o_t)) \Rightarrow s_t = g_{\psi}(o_t)$$

⇒ Problems: to maximize likelihood:

$$\max_{\phi, \psi} \frac{1}{N} \sum_{i=1}^N \sum_{t=1}^T \log p_{\phi}(g_{\psi}(o_{t+1,i}) | g_{\psi}(o_{t,i}), a_{t,i}) + \log p_{\phi}(o_{t,i} | g_{\psi}(o_{t,i}))$$

The whole recipe: [model-based RL with latent space]

$$\max_{\phi, \psi} \frac{1}{N} \sum_{i=1}^N \sum_{t=1}^T \log p_\phi(g_\psi(o_{t+1|i}) | g_\psi(o_{t,i}), a_{t,i}) + \log p_\phi(o_{t,i} | g_\psi(o_{t,i}))$$

latent space dynamics image reconstruction

$$+ \log p_\phi(r_{t,i} | g_\psi(o_{t,i}))$$

reward model

This is deterministic encoder model

⊗ Model-based RL with latent space model v. 1.5:

- 1, Run base policy $\pi_0(a_t | o_t)$ to collect $\mathcal{D} = \{(o, a, o')\}_i\}$
 - 2, Learn $p_\phi(s_{t+1}|s_t, a_t)$, $p_\phi(r_t, s_t)$, $p(o_t|s_t)$, $g_\psi(o_t)$ Learn model
 - 3, Plan through the model to choose actions (any, like MCTS, LQR, random shooting..)
 - 4, Execute the first planned action, observing result o' (NPC)
 - 5, Append (o, a, o') to dataset \mathcal{D}
- every N steps

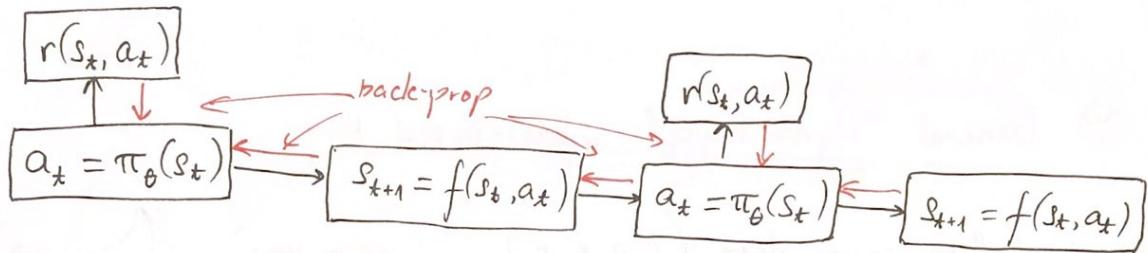
⊗ Also can learn directly in observation space:

learn $p(o_{t+1}|o_t, a_t)$ (take in image \rightarrow split out image)

Gigantic model: RNN ..

⊗ Open-loop case: observe ^{only} initial state \mathbf{v} \Rightarrow formulate a complete plan
 \Rightarrow commit to that plan

Closed-loop case: get state feed-back \Rightarrow learn policy



⊗ +, Model-based RL v. 2.0:

- 1, Run base policy $\pi_0(a_t | s_t)$ \Rightarrow collect $\mathcal{D} = \{(s_i, a_i, s'_i)\}_i$
- 2, Learn dynamics model $f(s, a)$ to minimize $\sum_i \|f(s_i, a_i) - s'_i\|^2$
- ⊗ 3, Back-propagate through $f(s, a)$ into the policy to optimize $\pi_\theta(a_t | s_t)$
- 4, Run $\pi_\theta(a_t | s_t)$, appending the visited tuples (s, a, s') to \mathcal{D}

⊗ Problem: (like RNN) exploding / vanishing gradients

⇒ Solutions:

- +, Use derivative-free (model-free) algorithms with the model used to generate synthetic samples
Example: Policy gradients has high variance, which can be reduced with lots of data, which can be generated by learned model
- + Use simpler policies than neural nets
 - LQR with learned models (LQR-FLM Fitted local model)
 - Train local policies to solve simple tasks
 - Combine into global policies via supervised learning

1) Use "model-free" RL algorithm

Problems: when generate longer & longer trajectories

If your model is inaccurate (which always is), the longer we roll-out the model, the more these errors compound

⇒ Not very nice for Policy Gradients, -.

④ General "Dyna-style" model-based RL:

1, Collect some data (s, a, s', r) (1-million steps)

2, Learn model $\hat{p}(s'|s, a)$ ($\&$ optionally $\hat{r}(s, a)$)

3, Repeat K times:

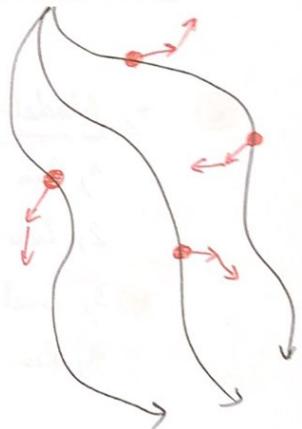
4, Sample $s \sim \beta$ from buffer

5, Choose action a (from β , π , or random)

6, Simulate $s' \sim \hat{p}(s'|s, a)$ (and $r = \hat{r}(s, a)$)

7, Train on (s, a, s', r) with model-free RL

8, (Optional) take N more model-based steps



④ Note: ↗ Step 5: Classic Dyna choose action from buffer

↑ only requires short roll-outs (as few as 1 step)
still see diverse states

⇒ Solve above problem

↪ Use simpler policies than neural nets

🚫 In order to use LQR, we need $\frac{df}{dx_t}$ & $\frac{df}{du_t}$

⇒ Idea: fit $\frac{df}{dx_t} \propto \frac{df}{du_t}$ around current trajectory/policy

If continuous system sufficiently smooth & initial state distribution quite light ⇒ do linear regression at every time step

⌚ What controller to execute? how to choose u_t

+ Version 0.5: $p(u_t | x_t) = \delta(u_t = \hat{u}_t)$ doesn't correct deviations or drift

+ Version 1.0: $p(u_t | x_t) = \delta_{u_t = K_t(x_t - \hat{x}_t) + k_t + \hat{u}_t}$

Better, but a little too good

We need data to be a little bit cluster, but not too much still need to be varied

⇒ + Version 2.0: $p(u_t | x_t) = N(K_t(x_t - \hat{x}_t) + k_t + \hat{u}_t, \Sigma_t)$
Set $\Sigma_t = Q^{-1}_{u_t, u_t}$

⌚ How to fit the dynamics? $p(x_{t+1} | x_t, u_t)$

+ Version 1.0: At each time step using linear regression

$$p(x_{t+1} | x_t, u_t) = N(A_t x_t + B_t u_t + c, N_t); A_t \approx \frac{df}{dx_t}; B_t \approx \frac{df}{du_t}$$

Problems: linear regression requires number of samples that scale with dimensional states

⇒ + Version 2.0: fit $p(x_{t+1} | x_t, u_t)$ using Bayesian linear regression
Use your favorite global model as prior

⇒ Can get away with fewer samples

⌚ How to stay close to old controller?

Keep KL-Divergence small (between old & new trajectory)

We want to stay close around local region of trajectories where we have linearized to appear

By add that element to loss function

+) Guided policy search:

From LQR, we only get local policy
⇒ need to combine it to global policy

- 1) Have many

Guided policy search algorithm:

- 1) Optimize each local policy $\pi_{LQR,i}(u_t | x_t)$ on initial state $x_{0,i}$ w.r.t. $\tilde{c}_{k,i}(x_t, u_t)$ to keep $\pi_{LQR,i}$ close to π_θ
- 2) Use sample from (1) to train $\pi_\theta(u_t | x_t)$ to mimic each $\pi_{LQR,i}(u_t | x_t)$
- 3) Update cost function $\tilde{c}_{k+1,i}(x_t, u_t) = c(x_t, u_t) + \lambda_{k+1,i} \log \pi_\theta(u_t | x_t)$

State estimation

+ Bayes Filters:

$$\text{bel}(x_t) = p(x_t | z_{1:t}, u_{1:t}) \quad \text{belief over a state}$$

$$\overline{\text{bel}}(x_t) = p(x_t | z_{1:t-1}, u_{1:t}) \quad \text{a posterior (before adapt to } z_t \text{)}$$

⇒ Calculating $\text{bel}(x_t)$ from $\overline{\text{bel}}(x_t)$: correction/measurement update

Algorithm Bayes Filter:

2 for all x_t do:

$$3 \quad \overline{\text{bel}}(x_t) = \int p(x_t | u_t, x_{t-1}) \text{bel}(x_{t-1}) dx \quad \Rightarrow \text{prediction step}$$

$$4 \quad \text{bel}(x_t) = n p(z_t | x_t) \overline{\text{bel}}(x_t) \quad \Rightarrow \text{update step}$$

⇒ Can only be implemented: for very simple estimation problems
finite state space

④ Important Assumption: Markov property (each state is a complete summary of the past)

⇒ Problem: can not be implemented on digital computer

+ The Kalman Filter (KF)

- For continuous state space, not discrete or hybrid
- Assumption: posterior are Gaussians & Markov property
 - $p(x_t | u_t, x_{t-1})$ must be linear func (linear system dynamics)
 - $p(z_t | x_t)$ also linear
 - initial belief(x_0) must be Gaussian

$$x_t = A_t x_{t-1} + B_t u_t + \epsilon_t$$

$$p(x_t | u_t, x_{t-1}) = \det(2\pi R_t)^{-\frac{1}{2}} \exp \left\{ -\frac{1}{2} (x_t - A_t x_{t-1} - B_t u_t)^T R_t^{-1} (x_t - A_t x_{t-1} - B_t u_t) \right\}$$

$$z_t = C_t x_t + \delta_t (= y) (= y)$$

$$p(z_t | x_t) = \det(2\pi Q_t)^{-\frac{1}{2}} \exp \left\{ -\frac{1}{2} (z_t - C_t x_t)^T Q_t^{-1} (z_t - C_t x_t) \right\}$$

$$\text{bel}(x_0) = p(x_0) = \det(2\pi \Sigma_0)^{-\frac{1}{2}} \exp \left\{ -\frac{1}{2} (x_0 - \mu_0)^T \Sigma_0^{-1} (x_0 - \mu_0) \right\}$$

Algorithm Kalman Filter ($\mu_{t-1}, \Sigma_{t-1}, u_t, z_t$)

$$2 \quad \bar{\mu}_t = A_t \mu_{t-1} + B_t u_t \quad \left. \begin{array}{l} \text{incorporate } u_t - \text{prediction step} \\ O(n^2) \end{array} \right\}$$

$$3 \quad \bar{\Sigma}_t = A_t \Sigma_{t-1} A_t^T + R_t \quad O(n^2)$$

$$4 \quad K_t = \bar{\Sigma}_t C_t^T (C_t \bar{\Sigma}_t C_t^T + Q_t)^{-1} \quad \left. \begin{array}{l} \text{Kalman gain} \\ \text{incorporate } z_t - \text{correction step} \end{array} \right\}$$

$$5 \quad \mu_t = \bar{\mu}_t + K_t (z_t - C_t \bar{\mu}_t) \quad O(n^{2,3})$$

$$6 \quad \Sigma_t = (I - K_t C_t) \bar{\Sigma}_t$$

$$7 \quad \text{return } \mu_t, \Sigma_t \Rightarrow \text{belief at time } t$$

\Rightarrow quite computationally expensive, everything are Gaussian

+), Extended Kalman Filter (EKF)

overcome linearity assumption by only approximate by Gaussians

$$x_t = g(u_t, x_{t-1}) + \epsilon_t$$

$$z_t = h(x_t) + \delta_t$$

Jacobian of state
 $n \times n$

$$\begin{aligned} g(u_t, x_{t-1}) &\approx g(u_t, \mu_{t-1}) + g'(u_t, \mu_{t-1})(x_{t-1} - \mu_{t-1}) \\ &= g(u_t, \mu_{t-1}) + G_t (x_{t-1} - \mu_{t-1}) \end{aligned}$$

$$p(x_t | u_t, x_{t-1}) \approx \det(2\pi R_t)^{-\frac{1}{2}} \exp \left\{ -\frac{1}{2} [x_t - g(u_t, \mu_{t-1}) - G_t (x_{t-1} - \mu_{t-1})]^T R_t^{-1} [x_t - g(u_t, \mu_{t-1}) - G_t (x_{t-1} - \mu_{t-1})] \right\}$$

$$\begin{aligned} h(x_t) &\approx h(\bar{\mu}_t) + h'(\bar{\mu}_t)(x_t - \bar{\mu}_t) \\ &= h(\bar{\mu}_t) + H_t (x_t - \bar{\mu}_t) \end{aligned}$$

$$p(z_t | x_t) \approx \det(2\pi Q_t)^{-\frac{1}{2}} \exp \left\{ -\frac{1}{2} [z_t - h(x_t)]^T Q_t^{-1} [z_t - h(x_t)] \right\}$$

Algorithm Extended Kalman Filter ($\mu_{t-1}, \Sigma_{t-1}, u_t, z_t$)

$$\bar{\mu}_t = g(u_t, \mu_{t-1})$$

$$\bar{\Sigma}_t = G_t \Sigma_{t-1} G_t^T + R_t$$

$$K_t = \bar{\Sigma}_t H_t^T (H_t \bar{\Sigma}_t H_t^T + Q_t)^{-1}$$

$$\mu_t = \bar{\mu}_t + K_t (z_t - h(\bar{\mu}_t))$$

$$\Sigma_t = (I - K_t H_t) \bar{\Sigma}_t$$

return μ_t, Σ_t

+ Can extend EKF \Rightarrow multi-hypothesis (extended) Kalman filter (MHEKF)

+ EKF's performance depends on degree of nonlinearities & uncertainty

+ Unscented KF & moments matching KF are better

⊗ Gaussian moment representation $\mu \& \Sigma$

canonical representation $\xi \& \Omega$

Information/precision matrix $\Omega = \Sigma^{-1}; \Sigma = \Omega^{-1}$

Information vector $\xi = \Sigma^{-1}\mu; \mu = \Omega^{-1}\xi$

$$\begin{aligned} p(x) &= \det(2\pi\Sigma)^{-\frac{1}{2}} \exp\left\{-\frac{1}{2}(x-\mu)^T \Sigma^{-1}(x-\mu)\right\} \\ &= n \exp\left\{-\frac{1}{2} x^T \Omega x + x^T \xi\right\} \end{aligned}$$

Algorithm Information filter (IF) ($\xi_{t-1}, \Omega_{t-1}, u_t, z_t$)

$$2 \quad \bar{\Omega}_t = (A_t \Omega_{t-1} A_t^T + R_t)^{-1} \quad \left. \right\} O(n^{2,8})$$

$$3 \quad \bar{\xi}_t = \bar{\Omega}_t (A_t \Omega_{t-1} \xi_{t-1} + B_t u_t) \quad \left. \right\} O(n^2)$$

$$4 \quad \Omega_t = C_t^T Q_t^{-1} C_t + \bar{\Omega}_t$$

$$5 \quad \xi_t = C_t^T Q_t^{-1} z_t + \bar{\xi}_t$$

Algorithm Extended Information filter

$$2 \quad \mu_{t-1} = \Sigma_{t-1}^{-1} \xi_{t-1}$$

$$3 \quad \bar{\Sigma}_t = (G_t \Sigma_{t-1}^{-1} G_t^T + R_t)^{-1}$$

$$4 \quad \bar{\xi}_t = \bar{\Sigma}_t g(u_t, \mu_{t-1})$$

$$5 \quad \bar{\mu}_t = g(u_t, \mu_{t-1})$$

$$6 \quad \Omega_t = \bar{\Sigma}_t + H_t^T Q_t^{-1} H_t$$

$$7 \quad \xi_t = \bar{\xi}_t + H_t^T Q_t^{-1} [z_t - h(\bar{\mu}_t) - H_t \bar{\mu}_t]$$

8 return ξ_t, Ω_t

$$\begin{aligned} x_t &= g(u_t, x_{t-1}) + \epsilon_t \\ z_t &= h(x_t) + \delta_t \\ G_t &= g'(u_t, x_{t-1}) \\ H_t &= h'(\bar{\mu}_t) \end{aligned}$$

- Global uncertainty: set $\Omega = 0$ better than set $|\Sigma| = \infty$

- EIF tends to be numerically more stable than KF

- EIF is better for multi-robot problems

- For high dimensional state, EKF is better computational than EIF

④ Major advantage of Gaussian filters is computational

disadvantage is assumption on unimodal distribution

Measurements

Maps: $m = \{m_1, \dots, m_N\}$

feature-based

m_n : properties of a feature
& location of feature

only the shape of the environment
at the specific locations

easy to adjust positions of objects
→ popular in the robotic mapping field

2 ways to represent a map

location-based

a specific location

volumetric: label for any location
in the world

occupancy map

+ Correct range with local measurement noise:

$$p_{\text{nois}}(z_t^k | x_t, m) = \begin{cases} n N(z_t^k; z_t^{k*}, \sigma_{\text{nois}}^2) & \text{if } 0 \leq z_t^k \leq z_{\max} \\ 0 & \text{otherwise} \end{cases}$$

+ Unexpected object:

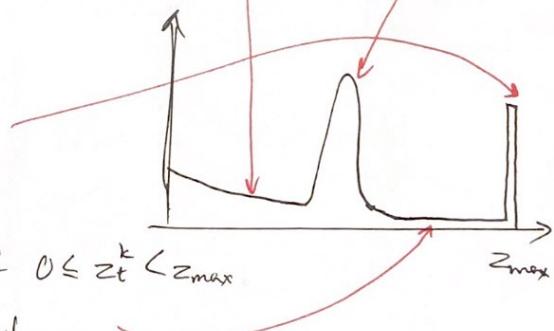
$$p_{\text{short}}(z_t^k | x_t, m) = \begin{cases} \eta \gamma_{\text{short}} e^{-\gamma_{\text{short}} z_t^k} & \text{if } 0 \leq z_t^k \leq z_t^{k*} \\ 0 & \text{otherwise} \end{cases}$$

+ Failures:

$$p_{\text{max}}(z_t^k | x_t, m) = I(z = z_{\max})$$

+ Random measurements

$$p_{\text{rand}}(z_t^k | x_t, m) = \begin{cases} \frac{1}{z_{\max}} & \text{if } 0 \leq z_t^k \leq z_{\max} \\ 0 & \text{otherwise} \end{cases}$$



$$p(z_t^k | x_{t,m}) = \begin{bmatrix} z_{\text{hot}} \\ z_{\text{short}} \\ z_{\text{max}} \\ z_{\text{rand}} \end{bmatrix}^T \begin{bmatrix} p_{\text{hot}}(z_t^k | x_{t,m}) \\ p_{\text{short}}(z_t^k | x_{t,m}) \\ p_{\text{max}}(z_t^k | x_{t,m}) \\ p_{\text{rand}}(z_t^k | x_{t,m}) \end{bmatrix}$$

Robot motion

Pose $\begin{bmatrix} x \\ y \\ \theta \end{bmatrix}$ } location
orientation

⊗ Motion model (Probabilistic kinematic model): $p(x_t | u_t, x_{t-1})$

Velocity commands

Odometry (distance traveled,
angle turned..)

- more accurate
- but post-the-fact (not for motion planning)

Probabilistic motion planning \leftarrow Use for \rightarrow estimation

Each has: closed form calculation & sampling algorithm

⊗ Velocity motion model: assumes we can control a robot through velocities:

$$u_t = \begin{bmatrix} v_t \\ \omega_t \end{bmatrix}$$

$$x_{t-1} = \begin{bmatrix} x \\ y \\ \theta \end{bmatrix}; x_t = \begin{bmatrix} x' \\ y' \\ \theta' \end{bmatrix}$$

Algorithm Motion model velocity (x_t, u_t, x_{t-1})

$$2 \quad \mu = \frac{1}{2} \frac{(x-x')\cos\theta + (y-y')\sin\theta}{(y-y')\cos\theta - (x-x')\sin\theta}$$

$$3 \quad x^* = \frac{x+x'}{2} + \mu(y-y')$$

$$4 \quad y^* = \frac{y+y'}{2} + \mu(x'-x)$$

$$5 \quad r^* = \sqrt{(x-x^*)^2 + (y-y^*)^2}$$

$$6 \quad \Delta\theta = \text{atan}2(y-y^*, x'-x^*) - \text{atan}2(y-y^*, x-x^*)$$

$$7 \quad \hat{v} = \frac{\Delta\theta}{\Delta t} r^*$$

$$8 \quad \hat{\omega} = \frac{\Delta\theta}{\Delta t}$$

$$9 \quad \hat{\theta} = \frac{\theta' - \theta}{\Delta t} - \hat{\omega}$$

$$10 \quad \text{return } \text{prob}(v-\hat{v}, \alpha_1|v| + \alpha_2|\omega|) \cdot \text{prob}(\omega - \hat{\omega}, \alpha_3|v| + \alpha_4|\omega|) \cdot \text{prob}(\hat{\theta}, \alpha_5|v| + \alpha_6|\omega|)$$

Invert the motion model

compared actual with the commanded

Algorithm Sample motion model velocity: $(\dot{x}_t, \dot{x}_{t+1})$

- 2 $\hat{v} = v + \text{sample}(\alpha_1|v| + \alpha_2|\omega|)$
- 3 $\hat{\omega} = \omega + \text{sample}(\alpha_3|v| + \alpha_4|\omega|)$
- 4 $\hat{\gamma} = \text{sample}(\alpha_5|v| + \alpha_6|\omega|)$
- 5 $\dot{x}' = x - \frac{\hat{v}}{\hat{\omega}} s\theta + \frac{\hat{v}}{\hat{\omega}} s(\theta + \hat{\omega}\Delta t)$
- 6 $y' = y + \frac{\hat{v}}{\hat{\omega}} c\theta - \frac{\hat{v}}{\hat{\omega}} c(\theta + \hat{\omega}\Delta t)$
- 7 $\theta' = \theta + \hat{\omega}\Delta t + \hat{\gamma}\Delta t$
- 8 return $x_t = [\dot{x}', y', \theta']^T$

Prob-normal-distribution (a, b): return $\frac{1}{\sqrt{2\pi b}} e^{-\frac{1}{2} \frac{a^2}{b}}$

Prob-triangular-distribution (a, b): return if $|a| > \sqrt{6b}$ return 0
else return $\frac{\sqrt{6b} - |a|}{6b}$

Sample-normal-distribution (b): return $\frac{b}{6} \sum_{i=1}^{12} \text{rand}(-1, 1)$

Sample-triangular-distribution (b): return $b \cdot \text{rand}(-1, 1) \cdot \text{rand}(-1, 1)$

⊗ Odometry motion model

only available after the robot has moved
 ⇒ only use for filter algorithm
 not for accurate motion planning & control

Algorithm motion mode odometry (x_t, u_t, x_{t-1})

$$2 \quad \delta_{\text{rot}1} = \text{atan} 2 (\bar{y}' - \bar{y}, \bar{x}' - \bar{x}) - \bar{\theta}$$

$$3 \quad \delta_{\text{trans}} = \sqrt{(\bar{x} - \bar{x}')^2 + (\bar{y} - \bar{y}')^2}$$

$$4 \quad \delta_{\text{rot}2} = \bar{\theta}' - \bar{\theta} - \delta_{\text{rot}1} \quad \begin{matrix} \text{hypothesized} \\ \text{final} \\ \text{pose} \end{matrix}$$

$$5 \quad \hat{\delta}_{\text{rot}1} = \text{atan} 2 (y' - y, x' - x) - \theta$$

$$6 \quad \hat{\delta}_{\text{trans}} = \sqrt{(x - x')^2 + (y - y')^2}$$

$$7 \quad \hat{\delta}_{\text{rot}2} = \theta' - \theta - \hat{\delta}_{\text{rot}1}$$

$$8 \quad p_1 = \text{prob} (\delta_{\text{rot}1} - \hat{\delta}_{\text{rot}1}, \alpha_1 \hat{\delta}_{\text{rot}1} + \alpha_2 \hat{\delta}_{\text{trans}})$$

$$9 \quad p_2 = \text{prob} (\delta_{\text{trans}} - \hat{\delta}_{\text{trans}}, \alpha_3 \hat{\delta}_{\text{trans}} + \alpha_4 (\hat{\delta}_{\text{rot}1} + \hat{\delta}_{\text{rot}2}))$$

$$10 \quad p_3 = \text{prob} (\delta_{\text{rot}2} - \hat{\delta}_{\text{rot}2}, \alpha_1 \hat{\delta}_{\text{rot}2} + \alpha_2 \hat{\delta}_{\text{trans}})$$

$$11 \quad \text{return } p_1 \cdot p_2 \cdot p_3 \quad (= p(x_t | u_t, x_{t-1}))$$

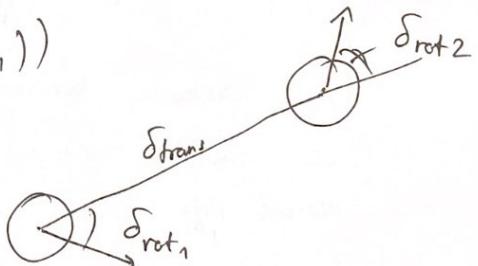
$$\bar{x}_{t-1} = [\bar{x} \quad \bar{y} \quad \bar{\theta}]^T$$

$$\bar{x}_t = [\bar{x}' \quad \bar{y}' \quad \bar{\theta}']^T$$

bar \Leftrightarrow measurements

hat \Leftrightarrow estimations

\Rightarrow inverse motion model



Algorithm Sample motion model odometry (u_t, x_{t-1})

$$2 \quad \delta_{\text{rot}1} = \text{atan2}(\bar{y}' - \bar{y}, \bar{x}' - \bar{x}) - \bar{\theta}$$

$$3 \quad \delta_{\text{trans}} = \sqrt{(\bar{x}' - \bar{x})^2 + (\bar{y}' - \bar{y})^2}$$

$$4 \quad \delta_{\text{rot}2} = \bar{\theta}' - \bar{\theta} - \delta_{\text{rot}1}$$

$$5 \quad \hat{\delta}_{\text{rot}1} = \delta_{\text{rot}1} - \text{sample}(\alpha_1 \delta_{\text{rot}1} + \alpha_2 \delta_{\text{trans}})$$

$$6 \quad \hat{\delta}_{\text{trans}} = \delta_{\text{trans}} - \text{sample}(\alpha_3 \delta_{\text{trans}} + \alpha_4 (\delta_{\text{rot}1} + \delta_{\text{rot}2}))$$

$$7 \quad \hat{\delta}_{\text{rot}2} = \delta_{\text{rot}2} - \text{sample}(\alpha_1 \delta_{\text{rot}2} + \alpha_2 \delta_{\text{trans}})$$

$$8 \quad x' = x + \hat{\delta}_{\text{trans}} \cos(\theta + \hat{\delta}_{\text{rot}1})$$

$$9 \quad y' = y + \hat{\delta}_{\text{trans}} \sin(\theta + \hat{\delta}_{\text{rot}1})$$

$$10 \quad \theta' = \theta + \hat{\delta}_{\text{rot}1} + \hat{\delta}_{\text{rot}2}$$

$$11 \quad \text{return } x_t = [x', y', \theta']^\top$$

Map-based motion model $p(x_t | u_t, x_{t-1}, m)$

- Occupancy maps: $p(x_t | m) = 0$ iff the robot collides

- If distance between $x_{t-1} \rightarrow x_t$ is small enough (< half robot's diameter)
we can estimate: $p(x_t | u_t, x_{t-1}, m) \approx \eta p(x_t | u_t, x_{t-1}) p(x_t | m)$
(discard info relating robot path to x_t)

Algorithm motion model with map (x_t, u_t, x_{t-1}, m)

$$\text{return } p(x_t | u_t, x_{t-1}).p(x_t | m)$$

Algorithm Sample motion model with maps (u_t, x_{t-1}, m)

do $x_t = \text{sample-motion-model}(u_t, x_{t-1})$

$$\pi = p(x_t | m)$$

until $\pi > 0$

return $\langle x_t, \pi \rangle$