

Robotics Notes

Huu Duc Nguyen M.Sc.

07 May 2022

Contents

| | |
|----------------------------------------------------------------------|-----------|
| Abbreviations | 1 |
| 1 Probabilistic Robotics | 2 |
| 1.1 State Estimation | 2 |
| 1.1.1 Bayes Filters | 2 |
| 1.1.2 The Kalman Filter (KF) | 2 |
| 1.1.3 Extended Kalman Filter (EKF) | 3 |
| 1.1.4 Information Filter (IF) | 4 |
| 1.1.5 Extended Information Filter (IF) | 5 |
| 1.2 Measurements | 5 |
| 1.2.1 Map Representation | 5 |
| 1.2.2 Measurement Noise | 6 |
| 1.3 Robot Motion | 6 |
| 1.3.1 Motion Model | 7 |
| 1.3.2 Velocity Motion Model | 7 |
| 1.3.3 Odometry Motion Model | 8 |
| 1.3.4 Map-based Motion Model | 9 |
| 2 Markov Decision Process | 11 |
| 2.1 Definitions | 11 |
| 2.1.1 Markov Chain | 11 |
| 2.1.2 Markov Decision Process (MDP) | 11 |
| 2.1.3 Partially Observable Markov Decision Process (POMDP) | 12 |
| 2.2 Bellman equations | 13 |
| 2.3 Partially Observable Markov Decision Process | 13 |

Abbreviations

| | |
|---------------|----------------------------------------------|
| prob. | probability |
| a.k.a. | also known as |
| func. | function |
| vs. | versus |
| KF | Kalman Filter |
| EKF | Extended Kalman Filter |
| IF | Information Filter |
| EIF | Extended Information Filter |
| MHEKF | Multi-Hypothesis Extended Kalman Filter |
| MDP | Markov Decision Process |
| POMDP | Partially Observable Markov Decision Process |

1 Probabilistic Robotics

Reference from the great book: [thrun2006probabilistic].

1.1 State Estimation

1.1.1 Bayes Filters

$$\begin{aligned} \text{bel}(x_t) &= p(x_t | z_{1:t}, u_{1:t}) && \text{belief over a state} \\ \overline{\text{bel}}(x_t) &= p(x_t | z_{1:t-1}, u_{1:t}) && \text{a posterior (before adapt to } z_t) \\ \Rightarrow \text{Calculating } \text{bel}(x_t) \text{ from } \overline{\text{bel}}(x_t) & && \text{correction/measurement update} \end{aligned}$$

Bayes Filter Algorithm:

- 2 for all x_t do:
- 3 $\overline{\text{bel}}(x_t) = \int p(x_t | u_t, x_{t-1}) \text{bel}(x_{t-1}) dx$ prediction step
- 4 $\text{bel}(x_t) = \eta p(z_t | x_t) \overline{\text{bel}}(x_t)$ update step

\Rightarrow Can only be implemented for very simple estimation problems, finite state space

Important assumption: Markov property (each state is a complete summary of the past)

Problem: can not be implemented on digital computers

The next subsections describe two Gaussian filters (Kalman Filter ([KF](#)) and Information Filter ([IF](#))) and two extensions of them (Extended Kalman Filter ([EKF](#)) and Extended Information Filter ([EIF](#))). The Gaussian filters have major advantage in computational cost, with the disadvantage of having assumption on uni-model distribution.

1.1.2 The Kalman Filter (KF)

- **Learning Resources:** www.bzarg.com
- For continuous space, not discrete or hybrid

- **Assumption:** posterior are Gaussians and Markov property

$p(x_t|u_t, x_{t-1})$ must be linear **func.** (linear system dynamics)

$p(z_t, x_t)$ also linear

$bel(x_0)$ initial belief must be Gaussian

$$x_t = A_t x_{t-1} + B_t u_t + \varepsilon_t$$

$$p(x_t|u_t, x_{t-1}) = \det(2\pi\Sigma_t)^{-\frac{1}{2}} \exp\left[-\frac{1}{2}(x_t - A_t x_{t-1} - B_t u_t)^T R_t^{-1} (x_t - A_t x_{t-1} - B_t u_t)\right]$$

$$z_t = C_t x_t + \delta_t \quad (=y)$$

$$p(z_t|x_t) = \det(2\pi Q_t)^{-\frac{1}{2}} \exp\left[-\frac{1}{2}(z_t - C_t x_t)^T Q_t^{-1} (z_t - C_t x_t)\right]$$

$$bel(x_0) = p(x_0) = \det(2\pi\Sigma_0)^{-\frac{1}{2}} \exp\left[-\frac{1}{2}(x_t - \mu_0)^T \Sigma_0^{-1} (x_t - \mu_0)\right]$$

Kalman Filter Algorithm: $(\mu_{t-1}, \Sigma_{t-1}, u_t, z_t)$

$$\begin{aligned} 2 \quad & \bar{\mu}_t = A_t \mu_{t-1} + B_t u_t \\ 3 \quad & \bar{\Sigma}_t = A_t \Sigma_{t-1} A_t^T + R_t \\ 4 \quad & K_t = \bar{\Sigma}_t C_t^T \left(C_t \bar{\Sigma}_t C_t^T + Q_t \right)^{-1} \quad (\text{Kalman gain}) \\ 5 \quad & \mu_t = \bar{\mu}_t + K_t (z_t - C_t \bar{\mu}_t) \\ 6 \quad & \Sigma_t = (I - K_t C_t) \bar{\Sigma}_t \\ 7 \quad & \text{return } \mu_t, \Sigma_t \end{aligned} \quad \left. \begin{array}{l} \text{incorporate } u_t - \text{prediction step} - \mathcal{O}(n^2) \\ \text{incorporate } z_t - \text{correction step} - \mathcal{O}(n^{2,8}) \\ \Rightarrow \text{belief at time } t \end{array} \right\}$$

⇒ quite computationally expensive, **everything are Gaussians**

1.1.3 Extended Kalman Filter (EKF)

Overcome the assumption on linearity by only approximate by Gaussians

$$x_t = g(u_t, x_{t-1}) + \varepsilon_t$$

$$z_t = h(x_t) + \delta_t$$

1 Probabilistic Robotics

$$g(u_t, x_{t-1}) \approx g(u_t, \mu_{t-1}) + g'(u_t, \mu_{t-1})(x_t - \mu_{t-1})$$

$$= g(u_t, \mu_{t-1}) + G_t(x_t - \mu_{t-1})$$

g' is the Jacobian of state ($n \times n$ matrix)

$$p(x_t|u_t, x_{t-1}) \approx \det(2\pi R_t)^{-\frac{1}{2}} \exp \left\{ -\frac{1}{2} [x_t - g(u_t, x_{t-1})]^T R_t^{-1} [x_t - g(u_t, x_{t-1})] \right\}$$

$$h(t) \approx h(\bar{\mu}_t) + h'(\mu_t)(x_t - \mu_{t-1})$$

$$= h(\bar{\mu}_t) + H_t(x_t - \mu_{t-1})$$

$$p(z_t|x_t) = \det(2\pi Q_t)^{-\frac{1}{2}} \exp \left\{ -\frac{1}{2} [z_t - h(x_t)]^T Q_t^{-1} [z_t - h(x_t)] \right\}$$

Extended Kalman Filter Algorithm: $(\mu_{t-1}, \Sigma_{t-1}, u_t, z_t)$

- 2 $\bar{\mu}_t = g(u_t, \mu_{t-1})$
- 3 $\bar{\Sigma}_t = G_t \Sigma_{t-1} G_t^T + R_t$
- 4 $K_t = \bar{\Sigma}_t H_t^T \left(H_t \bar{\Sigma}_t H_t^T + Q_t \right)^{-1}$
- 5 $\mu_t = \bar{\mu}_t + K_t [z_t - h(\bar{\mu}_t)]$
- 6 $\Sigma_t = (I - K_t H_t) \bar{\Sigma}_t$
- 7 return μ_t, Σ_t

- Can extend EKF \Rightarrow Multi-Hypothesis Extended Kalman Filter (MHEKF)
- EKF's performance depends on degree of nonlinearities and uncertainty
- Unscented KF and moments matching KF are better

1.1.4 Information Filter (IF)

- Moment representation: μ & Σ
- Canonical representation: ξ & Ω
- Information precision matrix: $\Omega = \Sigma^{-1}; \quad \Sigma = \Omega^{-1}$
- Information vector: $\xi = \Sigma^{-1}\mu; \quad \mu = \Omega^{-1}\xi$

$$\begin{aligned} p(x) &= \det(2\pi\Sigma)^{-\frac{1}{2}} \exp \left\{ -\frac{1}{2}(x - \mu)^T \Sigma^{-1} (x - \mu) \right\} \\ &= \eta \exp \left\{ -\frac{1}{2} x^T \Omega x + x^T \xi \right\} \end{aligned}$$

Information Filter Algorithm: $(\xi_{t-1}, \Omega_{t-1}, u_t, z_t)$

$$\begin{aligned} 2 \quad & \bar{\Omega}_t = (A_t \Omega_{t-1}^{-1} A_t^T + R_t)^{-1} \\ 3 \quad & \bar{\xi}_t = \bar{\Omega}_t (A_t \Omega_{t-1}^{-1} \xi_{t-1} + B_t u_t) \\ 4 \quad & \Omega_t = C_t^T Q_t^{-1} C_t + \bar{\Omega}_t \\ 5 \quad & \xi_t = C_t^T Q_t^{-1} z_t + \bar{\xi}_t \\ 6 \quad & \text{return } \xi_t, \Omega_t \end{aligned} \quad \left. \begin{array}{l} \\ \\ \\ \end{array} \right\} \mathcal{O}(n^{2,8}) \quad \left. \begin{array}{l} \\ \\ \\ \end{array} \right\} \mathcal{O}(n^2)$$

1.1.5 Extended Information Filter (IF)

$$x_t = g(u_t, x_{t-1}) + \varepsilon_t$$

$$z_t = h(x_t) + \delta_t$$

$$G_t = g'(u_t, \mu_{t-1})$$

$$H_t = h'(\mu_t)$$

Extended Information Filter Algorithm: $(\xi_{t-1}, \Omega_{t-1}, u_t, z_t)$

$$\begin{aligned} 2 \quad & \mu_{t-1} = \Omega_{t-1}^{-1} \xi_{t-1} \\ 3 \quad & \bar{\Omega}_t = (G_t \Omega_{t-1}^{-1} G_t^T + R_t)^{-1} \\ 4 \quad & \bar{\xi}_t = \bar{\Omega}_t g(u_t, \mu_{t-1}) \\ 5 \quad & \bar{\mu}_t = g(u_t, \mu_{t-1}) \\ 6 \quad & \Omega_t = \bar{\Omega}_t + H_t^T Q_t^{-1} H_t \\ 7 \quad & \xi_t = \bar{\xi}_t + H_t^T Q_t^{-1} [z_t - h(\bar{\mu}_t) - H_t \bar{\mu}_t] \end{aligned}$$

- **Global uncertainty:** set $\Omega = 0$ is better than set $|\Sigma| = \infty$
- **IF** tends to be numerically more stable than **KF**
- **IF** is better for multi-robot problems
- For high dimensional state, **EKF** is computational better than **EIF**

1.2 Measurements

1.2.1 Map Representation

Maps: $m = \{m_1, m_2, \dots, m_N\}$

There are 2 ways to represent a map:

[TODO: Add images]

| feature-based | location-based |
|-----------------------------------------------------------------------------------------------------------------------------------------|-----------------------------------------------------------|
| m_n : properties of a feature and location of feature only the shape of the environment at the specific locations | a specific location |
| easy to adjust positions of objects ⇒ popular in the robotic mapping field | volumetric: label for any location in the world |
| | occupancy grid map |

1.2.2 Measurement Noise

The 4 types of measurement noise:

- Correct range with local measurement noise

With z_t^{k*} as the correct distance

$$p_{hit}(z_t^k | x_t, m) = \begin{cases} \eta \mathcal{N}(z_t^k | z_t^{k*}, \sigma_{hit}^2) & \text{if } 0 \leq z_t^k \leq z_{\max} \\ 0 & \text{otherwise} \end{cases} \quad (1.1)$$

- Unexpected object

$$p_{short}(z_t^k | x_t, m) = \begin{cases} \eta \lambda_{short} e^{-\lambda_{short} z_t^k} & \text{if } 0 \leq z_t^k \leq z_t^{k*} \\ 0 & \text{otherwise} \end{cases} \quad (1.2)$$

- Failures

$$p_{max}(z_t^k | x_t, m) = I(z = z_{\max}) \quad (1.3)$$

- Random measurements

$$p_{rand}(z_t^k | x_t, m) = \begin{cases} \frac{1}{z_{\max}} & \text{if } 0 \leq z_t^k < z_{\max} \\ 0 & \text{otherwise} \end{cases} \quad (1.4)$$

[TODO: Add image, plot]

$$p(z_t^k | x_t, m) = \begin{bmatrix} z_{hit} \\ z_{short} \\ z_{max} \\ z_{rand} \end{bmatrix}^T \cdot \begin{bmatrix} p_{hit}(z_t^k | x_t, m) \\ p_{short}(z_t^k | x_t, m) \\ p_{max}(z_t^k | x_t, m) \\ p_{rand}(z_t^k | x_t, m) \end{bmatrix} \quad (1.5)$$

1.3 Robot Motion

Pose: $[x, y, \theta]^T$ at location $[x, y]^T$ and orientation θ

1.3.1 Motion Model

Motion Model, also known as (a.k.a.) Probabilistic Kinematic Model: $p(x_t, u_t, x_{t-1})$

| Velocity commands | Odometry (distance traveled, angle turned, etc.) |
|---------------------------------------|------------------------------------------------------------------------------------------------|
| Use for Probabilistic motion planning | more accurate but post-the-fact (not for motion planning) Use for estimation |

Each has closed form calculation and sampling algorithm.

1.3.2 Velocity Motion Model

Assuming we can control a robot through velocities:

$$u_t = \begin{bmatrix} v_t \\ \omega_t \end{bmatrix}; \quad x_{t-1} = \begin{bmatrix} x \\ y \\ \theta \end{bmatrix}; \quad x_t = \begin{bmatrix} x' \\ y' \\ \theta' \end{bmatrix}$$

Motion Model Velocity Algorithm: (x_t, u_t, x_{t-1})

$$\left. \begin{array}{l}
 2 \quad \mu = \frac{1}{2} \frac{(x - x') \cos \theta + (y - y') \sin \theta}{(y - y') \cos \theta - (x - x') \sin \theta} \\
 3 \quad x^* = \frac{x+x'}{2} + \mu(y - y') \\
 4 \quad y^* = \frac{y+y'}{2} + \mu(x' - x) \\
 5 \quad r^* = \sqrt{(x - x^*)^2 + (y - y^*)^2} \\
 6 \quad \Delta\theta = \text{atan2}(y' - y^*, x' - x^*) - \text{atan2}(y - y^*, x - x^*) \\
 7 \quad \hat{v} = \frac{\Delta\theta}{\Delta t} r^* \\
 8 \quad \hat{\omega} = \frac{\Delta\theta}{\Delta t} \\
 9 \quad \hat{\gamma} = \frac{\theta' - \theta}{\Delta t} - \hat{\omega} \\
 10 \quad \text{return } p(v - \hat{v}, \alpha_1|v| + \alpha_2|\omega|) \cdot p(\omega - \hat{\omega}, \alpha_3|v| + \alpha_4|\omega|) \cdot p(\hat{\gamma}, \alpha_5|v| + \alpha_6|\omega|)
 \end{array} \right\} \begin{array}{l} \text{Invert the motion model} \\ \text{compared actual velocities with the desired} \end{array}$$

Sample Motion Model Velocity Algorithm: (u_t, x_{t-1})

```

2    $\hat{v} = v + \text{sample}(\alpha_1|v| + \alpha_2|\omega|)$ 
3    $\hat{\omega} = \omega + \text{sample}(\alpha_3|v| + \alpha_4|\omega|)$ 
4    $\hat{\gamma} = \text{sample}(\alpha_5|v| + \alpha_6|\omega|)$ 
5    $x' = x - \frac{\hat{v}}{\hat{\omega}} \sin \theta + \frac{\hat{v}}{\hat{\omega}} \sin(\theta + \hat{\omega}\Delta t)$ 
6    $y' = y + \frac{\hat{v}}{\hat{\omega}} \cos \theta - \frac{\hat{v}}{\hat{\omega}} \cos(\theta + \hat{\omega}\Delta t)$ 
7    $\theta' = \theta + \hat{\omega}\Delta t + \hat{\gamma}\Delta t$ 
8   return  $x_t = [x', y', \theta']^T$ 

```

- Probability normal distribution(a, b): return $\frac{1}{\sqrt{2\pi b}} e^{-\frac{a^2}{2b}}$
- Probability triangular distribution(a, b): return $\begin{cases} 0 & \text{if } |a| > \sqrt{6b} \\ \frac{\sqrt{6b}-|a|}{6b} & \text{else} \end{cases}$
- Sample normal distribution(b): return $\frac{b}{6} \sum_{i=1}^{12} \text{rand}(-1, 1)$
- Sample triangle distribution(b): return $b.\text{rand}(-1, 1).\text{rand}(-1, 1)$

1.3.3 Odometry Motion Model

Only available after the robot has moved

\Rightarrow only use for filter algorithm

not for accurate motion planning and control

Motion Model Odometry Algorithm: (x_t, u_t, x_{t-1})

```

2    $\delta_{rot1} = \text{atan2}(\bar{y}' - \bar{y}, \bar{x}' - \bar{x}) - \bar{\theta}$ 
3    $\delta_{trans} = \sqrt{(\bar{x} - \bar{x}')^2 + (\bar{y} - \bar{y}')^2}$ 
4    $\delta_{rot2} = \bar{\theta}' - \bar{\theta} - \delta_{rot1}$ 
5    $\hat{\delta}_{rot1} = \text{atan2}(y' - y, x' - x) - \theta$ 
6    $\hat{\delta}_{trans} = \sqrt{(x - x')^2 + (y - y')^2}$ 
7    $\hat{\delta}_{rot2} = \theta' - \theta - \hat{\delta}_{rot1}$ 
8    $p_1 = \text{prob}(\delta_{rot1} - \hat{\delta}_{rot1}, \alpha_1 \hat{\delta}_{rot1} + \alpha_2 \hat{\delta}_{trans})$ 
9    $p_2 = \text{prob}(\delta_{trans} - \hat{\delta}_{trans}, \alpha_3 \hat{\delta}_{trans} + \alpha_4 (\hat{\delta}_{rot1} + \hat{\delta}_{rot2}))$ 
10   $p_3 = \text{prob}(\delta_{rot2} - \hat{\delta}_{rot2}, \alpha_1 \hat{\delta}_{rot2} + \alpha_2 \hat{\delta}_{trans})$ 
11  return  $p_1 \cdot p_2 \cdot p_3$       ( $= p(x_t | u_t, x_{t-1})$ )

```

\Rightarrow inverse motion model

NOTE:

- Bar \Leftrightarrow measurements

$$\bar{x}_{t-1} = [\bar{x} \quad \bar{y} \quad \bar{\theta}]^T$$

$$\bar{x}_t = [\bar{x}' \quad \bar{y}' \quad \bar{\theta}']^T$$
- Hat \Leftrightarrow estimations
- No bar and hat \Leftrightarrow hypothesized final pose x, y

Sample Motion Model Odometry Algorithm: (u_t, x_{t-1})

```

2       $\delta_{rot1} = \text{atan2}(\bar{y}' - \bar{y}, \bar{x}' - \bar{x}) - \bar{\theta}$ 
3       $\delta_{trans} = \sqrt{(\bar{x} - \bar{x}')^2 + (\bar{y} - \bar{y}')^2}$ 
4       $\delta_{rot2} = \bar{\theta}' - \bar{\theta} - \delta_{rot1}$ 
5       $\hat{\delta}_{rot1} = \delta_{rot1} - \text{sample}(\alpha_1 \delta_{rot1} + \alpha_2 \delta_{trans})$ 
6       $\hat{\delta}_{trans} = \delta_{trans} - \text{sample}(\alpha_3 \delta_{trans} + \alpha_4 (\delta_{rot1} + \delta_{rot2}))$ 
7       $\hat{\delta}_{rot2} = \delta_{rot2} - \text{sample}(\alpha_1 \delta_{rot2} + \alpha_2 \delta_{trans})$ 
8       $x' = x + \hat{\delta}_{trans} \cos(\theta + \hat{\delta}_{rot1})$ 
9       $y' = y + \hat{\delta}_{trans} \sin(\theta + \hat{\delta}_{rot1})$ 
10      $\theta' = \theta + \hat{\delta}_{rot1} + \hat{\delta}_{rot2}$ 
11     return  $x_t = [x' \quad y' \quad \theta']^T$ 

```

1.3.4 Map-based Motion Model

Map-based Motion Model: $p(x_t|u_t, x_{t-1}, m)$

- Occupancy maps: $p(x_t|m) = 0 \Leftrightarrow$ the robot collides
- If the distance from $x_{t-1} \rightarrow x_t$ is small enough ($<$ half the robot's diameter), we can estimate the probability (prob.) $p(x_t|u_t, x_{t-1}, m) \approx \eta p(x_t|u_t, x_{t-1})p(x_t|m)$, which discards the info relating the robot's path to x_t

Motion Model with Map Algorithm: (x_t, u_t, x_{t-1}, m)

return $p(x_t|u_t, x_{t-1}) \cdot p(x_t|m)$

1 Probabilistic Robotics

Sample Motion Model with Map Algorithm: (u_t, x_{t-1}, m)

do :

$$x_t = \text{sample_motion_model}(u_t, x_{t-1})$$

$$\pi = p(x_t | m)$$

until $\pi > 0$

return $< x_t, \pi >$

2 Markov Decision Process

2.1 Definitions

2.1.1 Markov Chain


A Markov Chain \mathcal{M} is a tuple consisting of:

$$\mathcal{M} = \{\mathcal{S}, \mathcal{T}\}$$

\mathcal{S} – state space $s \in \mathcal{S}$ (discrete or continuous)

\mathcal{T} – transition operator $p(s_{t+1}|s_t)$ or $\vec{\mu}_{t+1} = \mathcal{T}\vec{\mu}_t$ (\mathcal{T} is a matrix)

Markov chain is a process without memories. In other words, the next state s_{t+1} depends only on the current state s_t , not the previous state s_{t-1} .



2.1.2 Markov Decision Process (MDP)

A Markov Decision Process (MDP) \mathcal{M} is a tuple consisting of:

$$\mathcal{M} = \{\mathcal{S}, \mathcal{A}, \mathcal{T}, r, \gamma\}$$


\mathcal{S} – state space $s \in \mathcal{S}$ (discrete or continuous)

\mathcal{A} – action space $a \in \mathcal{A}$ (discrete or continuous)

\mathcal{T} – transition operator $p(s_{t+1}|s_t)$ (now a tensor)

r – reward function $r : \mathcal{S} \times \mathcal{A} \rightarrow \mathbb{R}$, $r(s_t, a_t)$

γ – discount factor $\gamma \in [0, 1]$ (optional)



There are definitions relating to MDP:

2 Markov Decision Process

- Policy: choice of action (at each state): $\pi_\theta(\mathbf{a}_t|\mathbf{s}_t)$
- Utility: sum of (discounted) rewards

A MDP can also be considered as a Markov chain on the augmented state (\mathbf{s}, \mathbf{a}) , a.k.a. Q-state. Knowing the state transition $p(\mathbf{s}_{t+1}|\mathbf{s}_t, \mathbf{a}_t)$ and policy $\pi_\theta(\mathbf{a}_t|\mathbf{s}_t)$, the transition of these Q-states can be derived as follows:

$$p((\mathbf{s}_{t+1}, \mathbf{a}_{t+1})|(\mathbf{s}_t, \mathbf{a}_t)) = p(\mathbf{s}_{t+1}|\mathbf{s}_t, \mathbf{a}_t)\pi_\theta(\mathbf{a}_{t+1}|\mathbf{s}_{t+1}) \quad (2.1)$$

MDP state projects an expectimax-like search tree [TODO: add image]

2.1.3 Partially Observable Markov Decision Process (POMDP)

A Partially Observable Markov Decision Process (POMDP) \mathcal{M} is a tuple consisting of:

$$\mathcal{M} = \{\mathcal{S}, \mathcal{A}, \mathcal{O}, \mathcal{T}, \mathcal{E}, r, \gamma\}$$

\mathcal{S} – state space $s \in \mathcal{S}$ (discrete or continuous)

\mathcal{A} – action space $a \in \mathcal{A}$ (discrete or continuous)


\mathcal{O} – observation space $o \in \mathcal{O}$ (discrete or continuous)

\mathcal{T} – transition operator $p(s_{t+1}|s_t)$

\mathcal{E} – emission prob. $p(o_t|s_t)$

r – reward function $r : \mathcal{S} \times \mathcal{A} \rightarrow \mathbb{R}$

γ – discount factor $\gamma \in [0, 1]$ (optional)



Finite versus (vs.) infinite horizon: [TODO:]

To solve infinite utilities:

- Finite horizon
- Discounting
- Absorbing state (like the fire hole, overheating)

2.2 Bellman equations

- $T(s, a, s') = P(s'|s, a)$ the prob. of reaching state s' from taking action a at state s
- $R(s, a, s')$ the reward of making the transition
- $Q^*(s, a)$: expected utility starting in state s and having taken action a , then act optimally

$$Q^*(s, a) = \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V^*(s')] \quad (2.2)$$

It is the sum over possible next state s' , because there is uncertainty of which state s' will be reached, even with the same starting state s and action a .

- $V^*(s)$: value of a state - expected utility starting in state s and acting optimally

$$V^*(s) = \max_a Q^*(s, a) \quad (2.3)$$

$$\Rightarrow V^*(s) = \max_a \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V^*(s')] \quad (2.4)$$

- $\pi^*(s)$: optimal action / policy from state s

[TODO: Explain this??]

$$\mathbf{V}(\mathbf{s}) = \mathbb{E}[\mathbf{G}_t | \mathbf{S}_t = \mathbf{s}] = \mathbb{E} [\mathbf{R}_{t+1} + \gamma \mathbf{V}(\mathbf{S}_{t+1}) | \mathbf{S}_t = \mathbf{s}] = \mathbf{R}_s + \gamma \sum_{\mathbf{s}'} \mathbf{P}_{ss'} \mathbf{V}(\mathbf{s}') \quad (2.5)$$

2.3 Partially Observable Markov Decision Process

POMDP is defined with a tuple: $< S, A, O, P, R, Z, \gamma >$:

- S : state
- A : action
- O : **observation**
- P : transition matrix
- R : reward
- Z : **observation func.**
- γ : discount factor (optional)

$$Z_{s'o}^a = P [O_{t+1} = o | S_{t+1} = s', A_t = a] \quad (2.6)$$

$$\Rightarrow \begin{cases} V^*(b) \leftarrow \max_a \left[R_b^a + \gamma \sum_{s'} P_{bb'}^a V(b') \right] \\ \pi^*(b) = \arg \max_a \left[R_b^a + \gamma \sum_{s'} P_{bb'}^a V(b') \right] \end{cases} \quad (2.7)$$